UPPSALA UNIVERSITET

Machine learning – Block 1(b)

Måns Magnusson
Department of Statistics, Uppsala University

Autumn 2022

UPPSALA UNIVERSITET

This weeks lectures

- Regularization
- Model Selection and Assement
- Cross-Validation
- Evaluate classification models

UPPSALA UNIVERSITET

Section 1

Predictive Performance

UPPSALA UNIVERSITET

Previous Model Evaluation

- In the past, tools for assessing models, e.g.:
  - Residuals
  - Leverage, Cook's distance
  - p-values
  - R^2
  - AIC
- Model diagnoses and how well the model fits the data.
- In statistics: focus on estimation and attribution.

UPPSALA UNIVERSITET

Predictive Performance

- Supervised learning: focus on predictive performance:
- How well our model $\hat{f}$ trained on

$$\mathcal{T} = \{(y_i, x_i), i = 1, ..., n\}$$

work when predicting a new observation y_0 from the data generating process $P_{y,x}$.

$$\mathbb{E}(L(Y_0, X_0)|\mathcal{T}) = \int L(Y_0, \hat{f}(X_0))P_{y,x}d(Y_0, X_0)|\mathcal{T}$$

- ability to perform well on previously unobserved inputs is called generalization
- Models can be overly optimistic[1]:
  - explain training data well
  - poor generalizability
- a phenomenon known as overfitting.

[1] See e.g. Picard, R.R., Cook, R.D. (1984). Cross-validation of regression models. *Journal of the American Statistical Association*, **79(387)**, 575–583.

UPPSALA UNIVERSITET

Section 2

Measuring Performance

UPPSALA UNIVERSITET

Loss Functions

- To assess the performance we use the loss function for a new unseen observation y_0 and the prediction of that observation $\hat{f}(x_0)$

$$L(y_0, \hat{f}(x_0))$$

- This is quite general and we choose based L based on what we want the model performe well on.
- Examples:
  - Regression problems (squared loss/error):

$$L(y_0, \hat{f}(x_0)) = (y_0 - \hat{f}(x_0))^2$$

  - Classification (0-1 loss)

$$L(y_0, \hat{f}(x_0)) = I(y_0 = \hat{f}(x_0))$$

UPPSALA UNIVERSITET

Cross-Entropy Loss

- When we predict probabilities $\hat{f}(x_0) = \hat{p}$:

$$L(y_0, \hat{p}) = -(y_0 \log \hat{p}) + ((1 - y_0) \log (1 - \hat{p}))$$

Question: Do you recognize the (cross-entropy) loss function?

- Maximizing the likelihood is the same as minimizing the cross-entropy.
- Multi class cross-entroy over M classes

$$L(\mathbf{y}_0, \hat{\mathbf{p}}) = -\sum_{j=1}^{M} y_{0,j} \log \hat{p}_j$$

UPPSALA UNIVERSITET

The Confusion Matrix

- A common way to present performance in classification is the confusion matrix:

	Prediction	
Actual	Positive	Negative
Positive	True Positive (TP)	False Negative (FN)
Negative	False Positive (FP)	True Negative ((TN)

UPPSALA UNIVERSITET

The Confusion Matrix: Multi-class

	Prediction		
Actual	a	b	c
a	T_a	F_{ab}	F_{ac}
b	F_{ba}	T_b	F_{bc}
c	F_{ca}	F_{cb}	T_c

and

	Prediction		
Actual	TP	FP	FN
a	T_a	$F_{ba} + F_{ca}$	$F_{ab} + F_{ac}$
...	...	...	...

UPPSALA UNIVERSITET

Accuracy

$$\text{Accuracy} = \frac{(\text{TP}+\text{TN})}{(\text{TP}+\text{FP}+\text{FN}+\text{TN})}$$

or

$$\text{Accuracy} = \frac{T_a + T_b + T_c}{N}$$

Question: What is the problem with Accuracy?

UPPSALA UNIVERSITET

Precision

Of all the predicted positives, how many are actually positive?

$$\text{Precision} = \frac{(\text{TP})}{(\text{TP}+\text{FP})}$$

or

$$\text{Precision}_a = \frac{(T_a)}{T_a + F_{ba} + F_{ca})}$$

All predicted a: $T_a + F_{ba} + F_{ca}$

If we want one precision estimate for all classes:

1. Macro-average $(\text{Precision}_a, ..., \text{Precision}_c)$
2. Micro-average (use Table 2)

UPPSALA UNIVERSITET

Recall

Of all positives, how many are predicted correctly (recalled)?

$$\text{Recall} = \frac{(\text{TP})}{(\text{TP}+\text{FN})}$$

and

$$\text{Recall}_a = \frac{(T_a)}{T_a + F_{ab} + F_{ac})}$$

All true/actual a: $T_a + F_{ab} + F_{ac}$

If we want one precision estimate for all classes:

1. Macro-average $(\text{Recall}_a, ..., \text{Recall}_c)$
2. Micro-average (use Table 2)

UPPSALA UNIVERSITET

Sensitivity and specificity

$$\text{Sensitivity} = \text{Recall of positive class} = \frac{\text{TP}}{\text{TP+FN}}$$

and

$$\text{Specificity} = \text{Recall of negative class} = \frac{\text{TN}}{\text{TN+FP}}$$

UPPSALA UNIVERSITET

F1-score

Harmonic mean of Precision and Recall.

$$F_1 = \frac{2}{\text{Precision}^{-1} + \text{Recall}^{-1}} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Very common performance measurement in practice.

If we want one precision estimate for all classes:

1. Macro-average ($F_{1a}, ..., F_{1c}$)
2. Micro-average (use Table 2)

UPPSALA UNIVERSITET

Example

Say that we want to classify spam vs. ham.

	$\hat{f}(x) = 1$	$\hat{f}(x) = 0$
$y = 1$	515	91
$y = 0$	85	569

The cell counts yield us estimates of

1. Accuracy: $\frac{515+569}{515+91+85+569} \approx 0.86$
2. Precision: $\frac{515}{515+85} \approx 0.86$
3. Recall: $\frac{515}{515+91} \approx 0.85$
4. F_1: $\frac{2 \cdot 0.85 \cdot 0.86}{0.85+0.86} \approx 0.855$

In this example, we let $\hat{y}_i = 1$ whenever $\hat{\pi}_i > 0.5$.

What if we choose another cut-off level $\hat{\pi}_i > \alpha$ instead?

UPPSALA UNIVERSITET

Classification tables

$\alpha = 0.5$	$\hat{f}(x) = 1$	$\hat{f}(x) = 0$
$y = 0$	515	91
$y = 1$	85	569

Now let $\alpha = 0.3$ instead, so that we are more prone to say that $\hat{y} = 1$:

$\alpha = 0.3$	$\hat{f}(x) = 1$	$\hat{f}(x) = 0$
$y = 1$	462	144
$y = 0$	38	616

The cell counts yield us estimates of

1. Accuracy: $\frac{462+616}{462+38+144+616} \approx 0.86$
2. Precision: $\frac{462}{462+38} \approx 0.92$
3. Recall: $\frac{462}{462+144} \approx 0.76$
4. F_1: $\frac{2 \cdot 0.92 \cdot 0.76}{0.92+0.76} \approx 0.83$

The Precision has increased, but the Recall has decreased...

UPPSALA UNIVERSITET

A more problematic example

A highly unbalanced example. 1001 ham and 17 spam.

Our new classifier: Everything is ham!

	$\hat{f}(x) = 1$	$\hat{f}(x) = 0$
$y = 1$	1001	0
$y = 0$	17	0

The cell counts yield us estimates of

1. Accuracy: $\frac{1001}{1001+17} \approx 0.99$
2. Precision: $\frac{1001}{1001+0} \approx 1.0$
3. Recall: $\frac{1001}{1001+17} \approx 0.99$
4. F_1: $\frac{2 \cdot 1 \cdot 0.99}{0.99+1} \approx 0.99$
5. But Specificity is 0!

UPPSALA UNIVERSITET

Section 3

Test and training error

UPPSALA UNIVERSITET

Test Error

- The main error of interest - *generalization error*
- Conditional Test Error
  (Model performance for the actual training data):

$$\mathrm{Err}_{\mathcal{T}} = \mathbb{E}_{Y,X}(L(Y_0, \hat{f}(X_0)|\mathcal{T})$$

- Expected Test Error
  (Model performance over different training datasets):

$$\mathrm{Err} = \mathbb{E}_{\mathcal{T}}(\mathbb{E}_{Y,X}(L(Y_0, \hat{f}(X_0))))$$

- Sometimes refered to as generalization error.

UPPSALA UNIVERSITET

Training Error

- The loss function the algorithm try to minimize
- The Error in the training data:

$$\overline{\text{err}} = \frac{1}{N} \sum_{i=1}^{N} L(y_i, \hat{f}(x_i))$$

UPPSALA UNIVERSITET

Model complexity/capacity

- Model complexity/capacity: The flexibility of the model.
- Underfitting: Too low capacity of model
- Overfitting: Too high capacity of model
- Example: Polynomial regression with higher order terms

Figure: Model complexity (Goodfellow et al, 2017, Figure 5.2)

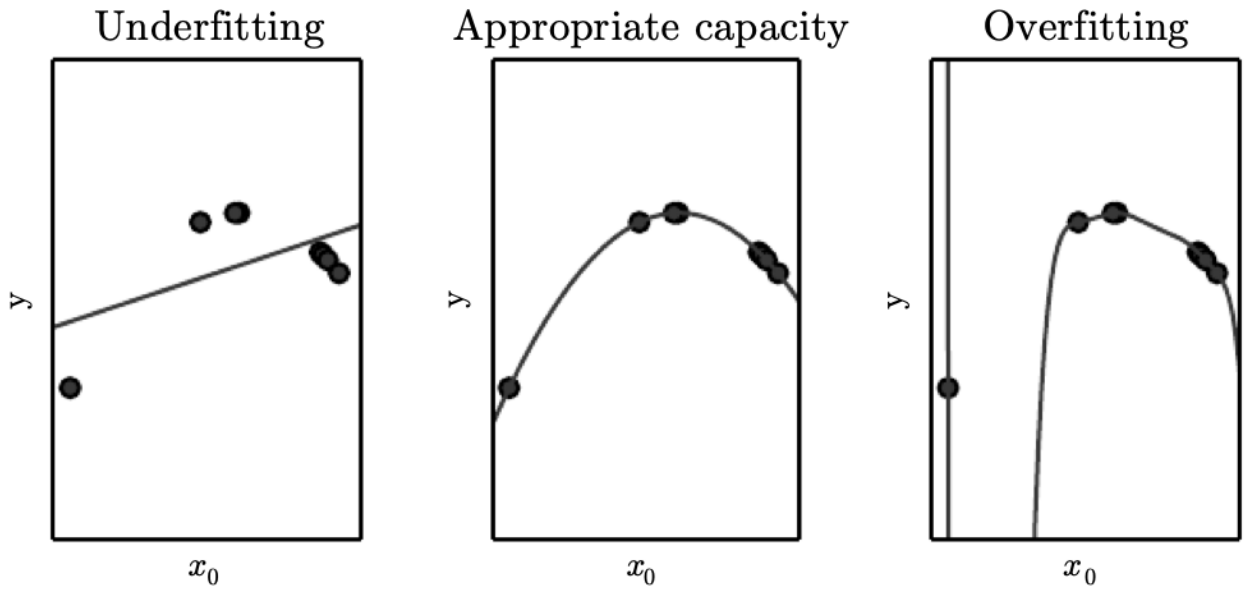

UPPSALA UNIVERSITET

Training, test, and complexity

Figure: Test, training, and model complexity (Goodfellow et al, 2017, Figure 5.3)

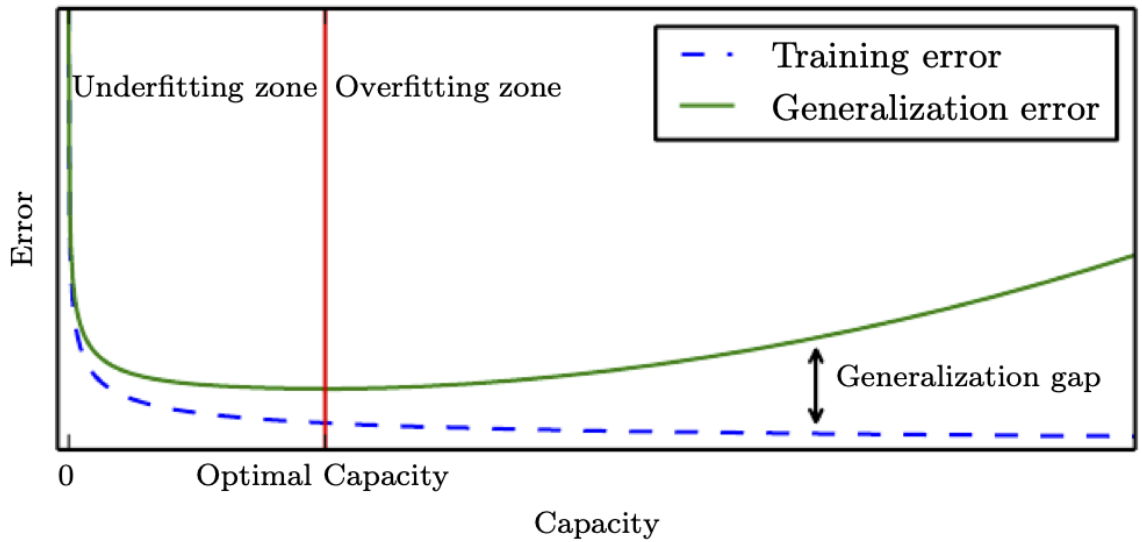

UPPSALA UNIVERSITET

How to estimate the Test Error (Model Assesment)

- We set aside a *test set* from the data
- Use as the last step to *estimate* the test error
- Should only be used ONCE
- Size of testset:
  - Common suggestion 10%
  - A statistical estimation problem (choice of samling size)

UPPSALA UNIVERSITET

Multiple Use of Test Set for Model Assesment

- What happens if we use the test set to pick the model?

UPPSALA UNIVERSITET

Section 5

Bias and Variance

UPPSALA UNIVERSITET

Bias and Variance

Assume we have the following data generating process:

$$Y = f(X) + \epsilon\,,$$

where $\mathbb{E}(\epsilon) = 0$ and $V(\epsilon) = \sigma_\epsilon$.

$$\begin{aligned}
\text{Err}(x_0) &= \mathbb{E}\{(Y - \hat{f}(x_0))^2 | X = x_0\} \\
&= \sigma_\epsilon^2 + \{\mathbb{E}(\hat{f}(x_0)) - f(x_0)\}^2 + \mathbb{E}\{\hat{f}(x_0) - \mathbb{E}(\hat{f}(x_0))\}^2 \\
&= \sigma_\epsilon^2 + \text{Bias}^2(\hat{f}(x_0)) + V(\hat{f}(x_0))
\end{aligned}$$

- *Bias*: How close can $\hat{f}$ get to the true model f
- *Variance*: The variability of the predictions from $\hat{f}$
- *Irreducible/Bayes error*: The minimum possible error

UPPSALA UNIVERSITET

Bias and Variance: Linear regression

In linear regression we have:

$$\hat{f}(x_i) = \hat{\beta} x_i$$

This give us the following error decomposition:

$$\frac{1}{N} \sum_{i}^{N} \text{Err}(x_i) = \sigma_\epsilon^2 + \frac{1}{N} \sum_{i}^{N} (f(x_i) - E(\hat{f}(x_i))^2 + \frac{p}{N} \sigma_\epsilon^2$$

UPPSALA UNIVERSITET

Bias and Variance

Figure: Test, training, and model complexity (Hastie et al, 2009, Figure 7)

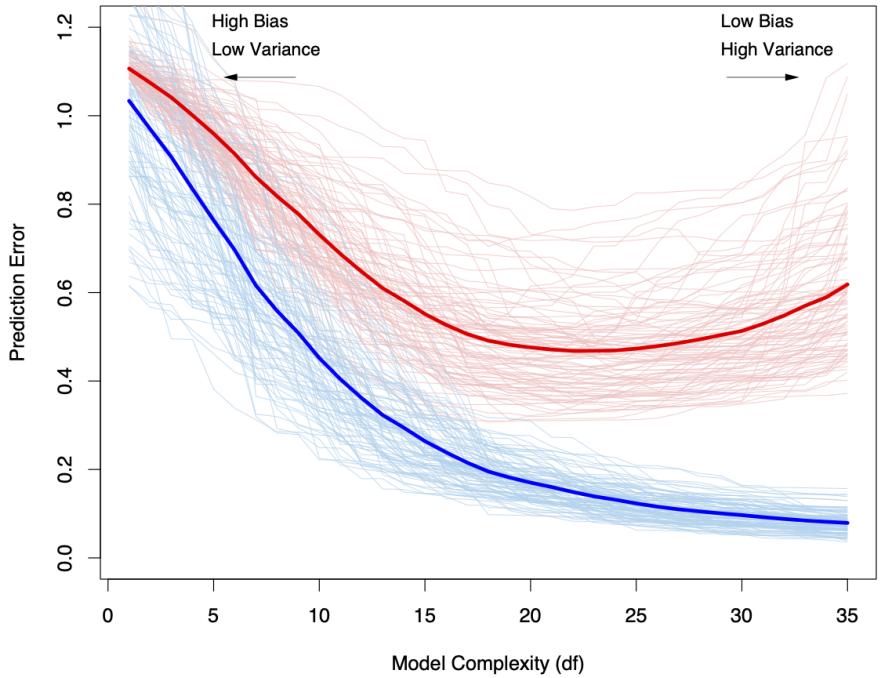

- High Bias: Underfit
- High Variance: Overfit
- High Irreducible error: No model is good

UPPSALA UNIVERSITET

Bias and Variance

Figure: Bias and variance (Goodfellow et al., 2017, Figure 5.6)

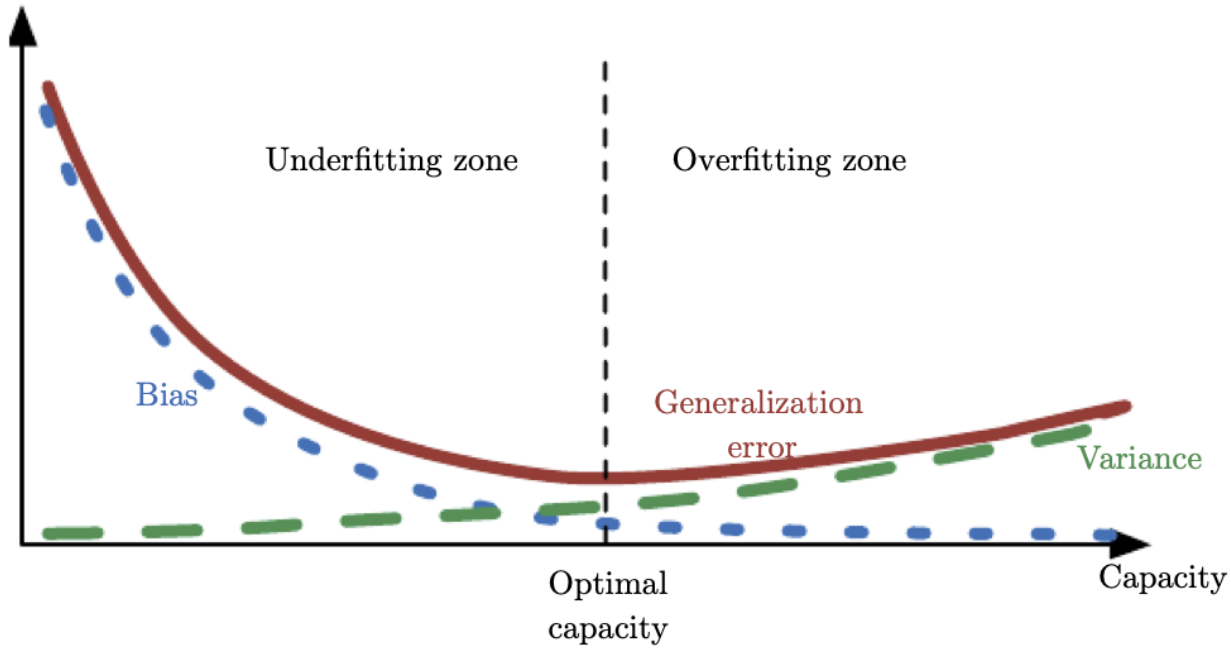

UPPSALA UNIVERSITET

Optimism

The in-sample test error:

$$\text{Err}_{\text{in}} = \frac{1}{N}\sum_{i=1}^{N} \mathbb{E}_{Y^0}\{L(Y_i^0, \hat{f}(x_i))|\mathcal{T}\},$$

where $Y_{0,i}$ is a new variable conditioned on x_i.

We have that

$$\mathbb{E}_{\mathbf{y}}(\text{Err}_{\text{in}}) = \mathbb{E}_{\mathbf{y}}(\overline{\text{err}}) + \underbrace{\frac{2}{N}\sum_{i=1}^{N} \text{Cov}(\hat{f}(x_i), y_i)}_{\text{optimism}},$$

where $\overline{\text{err}}$ is the training error.

Question: How could we create an optimistic classifier for the training data?

UPPSALA UNIVERSITET

Estimating Optimism

- Under certain conditions we can estimate this optimism.
- AIC is an example of this – asymptotic predictive performance.
- Find the optimism

UPPSALA UNIVERSITET

The double descent of large models

Figure: The double descent of large models (Nakkiran et al., 2019, Figure 1)

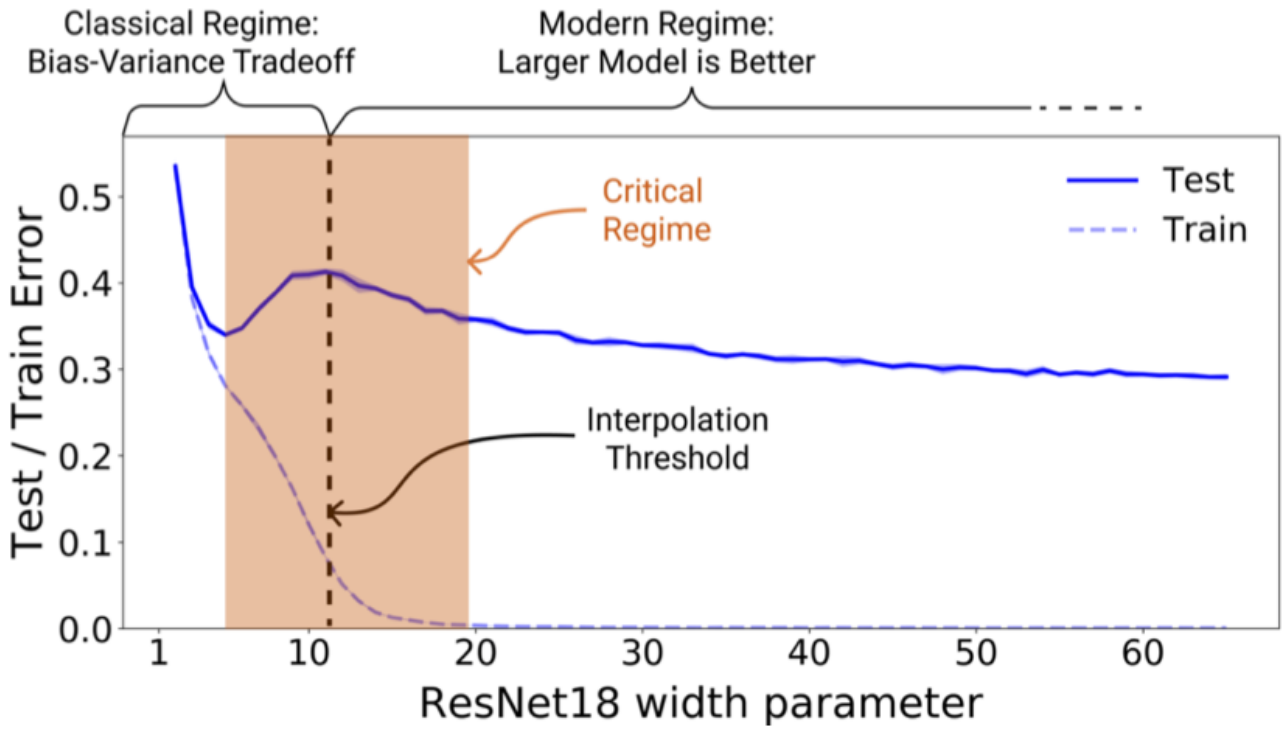

UPPSALA UNIVERSITET

Section 6

Cross-validation

UPPSALA UNIVERSITET

Cross-Valdidation

- We want to estimate Err for different models and to choose the best model where

$$\begin{aligned} \text{Err} &= \mathbb{E}_{\mathcal{T}}(\text{Err}_{\mathcal{T}}) \\ &= \mathbb{E}_{\mathcal{T}}(\mathbb{E}(L(Y_0, X_0)|\mathcal{T})) \\ &= \int (\int L(Y_0, \hat{f}(X_0)) P_{y,x} d(Y_0, X_0)|\mathcal{T}) d\mathcal{T} \end{aligned}$$

- Cross-Validation is probably the most popular approach to estimate Err and choose between models because it is
  1. Conceptually easy to understand
  2. Easy to implement
  3. No need for rules-of-thumbs to verify that it is applicable
  4. Equally useful for many different type of models
  5. Flexible for the use case at hand
- Common approach to learn hyper parameters (that is a model choice)

UPPSALA UNIVERSITET

Cross-Valdidation Algorithm

Figure: Cross-Validation (Hastie et al, 2009, p. 222, 242)

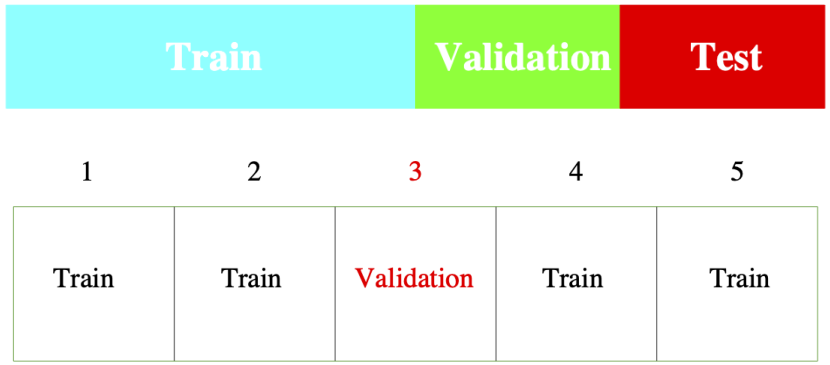

1. Split data in K folds
2. For each fold $k = 1, 2, ..., K$
   - 2.1 Use all samples except those in k to build $\hat{f}(x)$
   - 2.2 Use the model and predict the observations in fold k

$$\text{Err}_{CV}(\hat{f}) = \frac{1}{N} \sum_{i=1}^{N} L(y_i, \hat{f}_{-\kappa(i)}(x))$$

UPPSALA UNIVERSITET

The Bias of Cross-Valdidaton

- Cross-validation estimation of Err will be biased
- The training data size is smaller than the full data

Figure: Cross-Validation Bias (Hastie et al, 2009, Fig. 7.8)

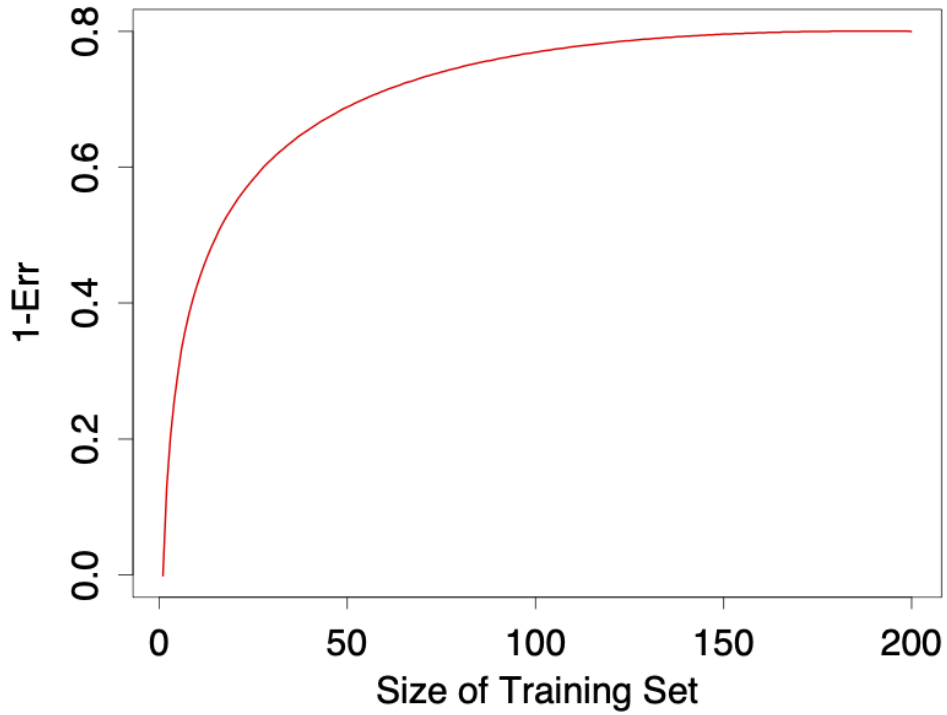

UPPSALA UNIVERSITET

K-fold Cross Validation

- Common K are: $K = \{2, 5, 10\}$
- Smaller K gives larger bias
- Larger K is computationally more costly
- $K = 10$ is a common approach

UPPSALA UNIVERSITET

Leave-One-Out Cross Validation

- When $K = N$
- Benefits
  - Almost unbiased estimate of Err
  - Sometimes we only need to train our model once
- Drawbacks
  - Higher Variance in estimate of Err
  - Can be more computationally very costly (naive implementation)
  - Can be unstable/less robust in some settings

UPPSALA UNIVERSITET

- Predictive Performance
- Measuring Performance
- Test and training error
- Estimating the test error
- Bias and Variance
- Cross-validation
- Regularisation

Leave-One-Out Cross Validation

Figure: Cross-Validation Bias (Hastie et al, 2009, Fig. 7.14)

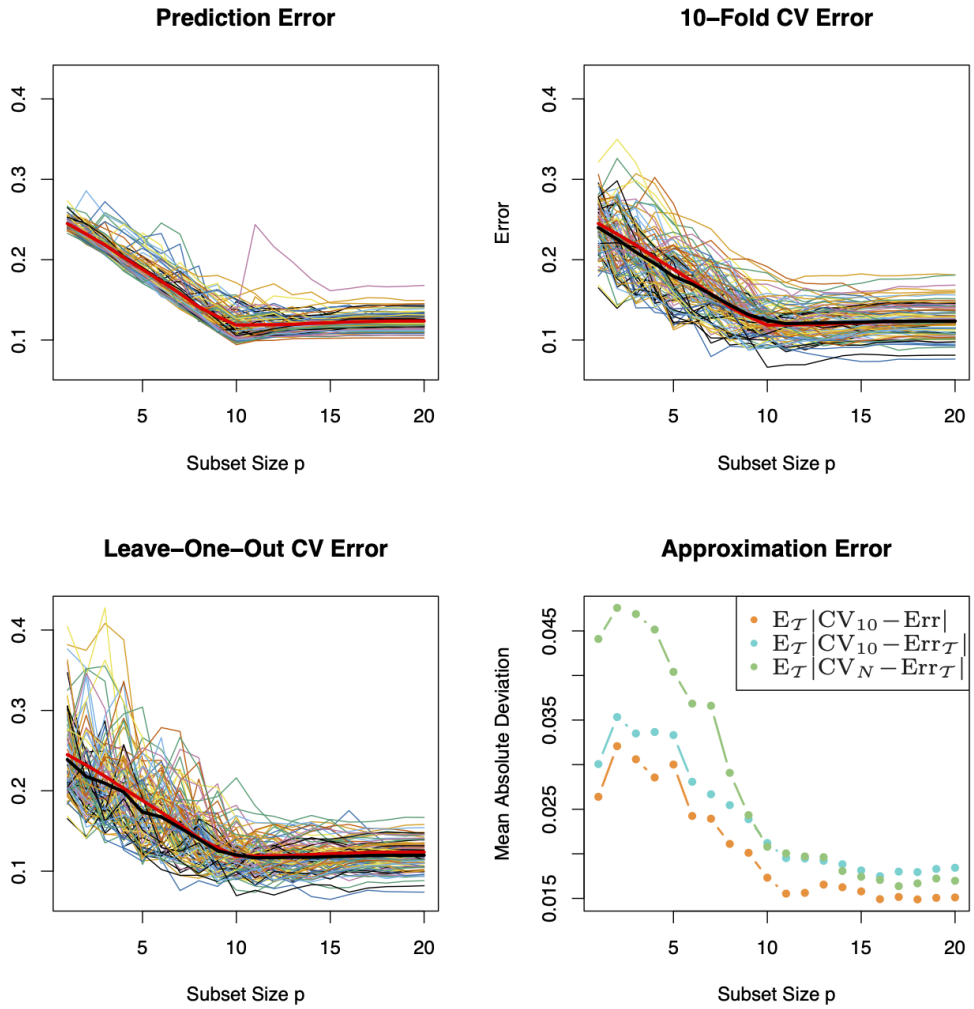

UPPSALA UNIVERSITET

The role of the data generating process

- we assume that testset and train set are different observations from the same data generating process

$$\mathbf{d} = \{(y_i, \mathbf{x}_i), i = 1, ..., n\} \sim P_{y,x}$$

- The (naive) assumption: independence
- Things that can go wrong:
  - temporal leak/concept drift
  - duplicated observations
  - non-randomized data

Example: Election prediction

- we want to predict the next election
- we know that there are "concept drift"
- Solution in Frölander and Uddhammar (2021) and Olsson and Ölfvingsson (2021)
  1. LOO-CV on the elections 1973-2014
  2. The elections 2018 as the final validation set

Figure: Predictive distr. (Olsson and Ölfvingsson, 2021, Fig. 6)

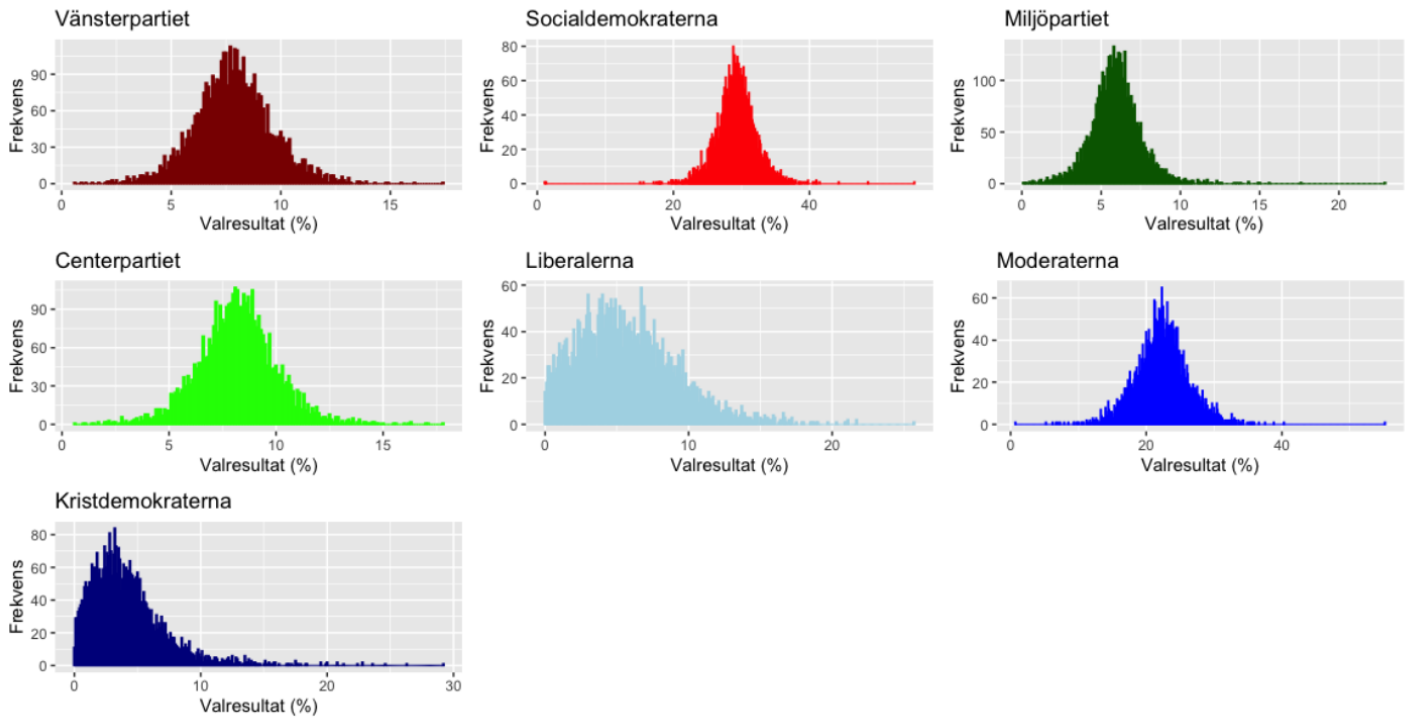

UPPSALA UNIVERSITET

Section 7

Regularisation

UPPSALA UNIVERSITET

Regression and GLM

- Linear regression and logistic regression are examples of generalised linear models, GLMs.
- Both use maximum likelihood estimation for fitting the model, where the likelihood function $L(\beta)$ is maximised.

UPPSALA UNIVERSITET

Regularised regression models

- In some situations, for instance when the predictors are highly collinear, when there are too many predictors or when there is complete separation in the data, maximum likelihood estimation is unstable.
  - Either the solution is not unique, or minuscule changes in the data can change the solution completely.
  - Such datasets are increasingly common in e.g. genomics, finance, astronomy and image analysis.
- In such cases, regularisation/shrinkage methods can be used instead.
- In a regularized GLM, it is not the likelihood $L(\beta)$ that is maximized, but a regularised function $L(\beta) \cdot p(\beta)$, where p is a penalty function that typically forces the resulting estimates to be closer to 0, which leads to a stable solution.

UPPSALA UNIVERSITET

Regularised regression models

Regularised linear regression models increase the bias of the estimates, but lowers their variance, thereby potentially decreasing the MSE.

UPPSALA UNIVERSITET

Connection to Bayesian estimation

In Bayesian estimation, a prior distribution $p(\beta)$ for the parameters β_i is chosen.

The estimates are then computed from the conditional distribution of the β_i given the data, called the posterior distribution.

Using Bayes' theorem, we find that

$$P(\beta|\mathbf{x}) \propto L(\beta) \cdot p(\beta),$$

i.e. that the posterior distribution is proportional to the likelihood times the prior.

A special type of Bayesian estimator is the maximum a posteriori (MAP) estimator, which is found by maximizing the above expression (i.e. finding the mode of the posterior).

This is equivalent to the estimates from a regularised frequentist model with penalty function $p(\beta)$!

UPPSALA UNIVERSITET

Inference and invariance

- Regularised regression models are not invariant under linear rescaling of the predictors.
  - If a predictor is multiplied by a scalar $a \neq 0$, this can change the entire model.
  - A model with measurements in inches might yield completely different results from a model with measurements in cm.
- For this reason, it is widely agreed that the predictors should be standardized to have mean 0 and variance 1 before a regularised model is fitted.
  - With this approach we choose a particular (natural?) scaling, among all possible scalings.
  - All predictors are on the same scale and are therefore treated equally by the penalty function.

UPPSALA UNIVERSITET

Inference and invariance

- Hypothesis tests are available (e.g. Lockhart et al. (2014), A significance test for the lasso, *Annals of Statistics*), but I advise against using them.
- Note that the hypothesis tests will be conditioned on the choice of scaling.
  - Because of this, regularised models are not appropriate for hypothesis testing – the p-values could change completely if we rescaled the data!
- Regularised regression models are however very useful for predictive modelling.

UPPSALA UNIVERSITET

L_q-penalties

The most popular penalty terms correspond to common L_q-norms. On a log-scale, the function to be maximized is

$$\ell(\beta) + \lambda \sum_{i=1}^{p} |\beta_i|^q,$$

where $\ell(\beta)$ is the loglikelihood of β and $\sum_{i=1}^{p} |\beta_i|^q$ is the L_q-norm, with $q \geq 0$.

This is equivalent to maximizing $\ell(\beta)$ under the constraint that $\sum_{i=1}^{p} |\beta_i|^q \leq \frac{1}{h(\lambda)}$, for some increasing positive function h.

- Relies on the sparsity assumption that most β are 0.

$\lambda \geq 0$ is a smoothing parameter:

- When $\lambda = 0$, we are back at the standard ML-estimate.
- The $\hat{\beta}$ are forced to be closer to 0 when λ increases.
- λ is usually chosen using cross-validation.

UPPSALA UNIVERSITET

Ridge regression

When the L_2 penalty is used, the regularised model is called ridge regression, for which we maximize

$$\ell(\beta) + \lambda \sum_{i=1}^{p} \beta_i^2.$$

- Invented and reinvented by several authors, from the 1940's onwards.
- In a linear model, the OLS estimate is $\hat{\beta} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$, whereas the ridge estimate is $\hat{\beta} = (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I})^{-1}\mathbf{X}^T\mathbf{y}$. The $\lambda\mathbf{I}$ is the 'ridge'.
- The β_i can become very small, but are never pushed all the way down to 0.
- In a Bayesian context, this corresponds to putting a standard normal prior on the β_i.

UPPSALA UNIVERSITET

Lasso

When the L_1 penalty is used, the regularised model is called the lasso (Least Absolute Shrinkage and Selection Operator), for which we maximize

$$\ell(\beta) + \lambda \sum_{i=1}^{p} |\beta_i|.$$

- Introduced by Robert Tibshirani in 1996.
- As λ increases, more and more β_i become 0.
  - Simultaneously performs estimation and variable selection!
- In a Bayesian context, this corresponds to putting a standard Laplace prior on the β_i.

UPPSALA UNIVERSITET

Examples in R

Functions for regularised generalized linear models (linear, logistic, Poisson, multinomial, and more) are available e.g. in the `glmnet` package for R.

The syntax used is somewhat different from that for `glm` and `lm`.

UPPSALA UNIVERSITET

Generalizations

Regularised models have been a hot research topics in the last 20 years. Some additional important models are:

- Elastic net: a compromise between ridge and lasso, in which

$$\ell(\beta) + \lambda_1 \sum_{i=1}^{p} |\beta_i| + \lambda_2 \sum_{i=1}^{p} \beta_i^2$$

  is maximized.
  - Introduced by Zou and Hastie in 2005.
  - Is better than the lasso at handling correlated predictors.
  - Has two smoothing parameters that we need to choose.
  - Available in the `glmnet` package.

- Group lasso: a version of the lasso in which variables can be grouped before fitting the model. The group lasso then selects groups of variables rather than individual variables.
  - Introduced by Yuan and Lin in 2006.
  - Useful e.g. when we have dummies for categorical variables (in contrast, the lasso may choose to only include the dummies for some of the categories).