



UPPSALA
UNIVERSITET

Machine learning – Block 3

Måns Magnusson
Department of Statistics, Uppsala University

Autumn 2022

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



UPPSALA
UNIVERSITET

Evaluation assignment 2

- Practicalities

- Introduction

- Feed-Forward Neural Networks

- Feed-Forward Neural Networks
- Hyper-parameters

- Regularization

- Optimization of Neural Networks

- Local minima
- Plateaus and Saddle Points
- Cliffs, exploding and vanishing gradients
- Parameter initialization

- Neural Networks in Practice

- t



UPPSALA
UNIVERSITET

This week's lecture

- **Practicalities**

- Introduction

- Feed-Forward Neural Networks

- Feed-Forward Neural Networks
- Hyper-parameters

- Regularization

- Optimization of Neural Networks

- Local minima
- Plateaus and Saddle Points
- Cliffs, exploding and vanishing gradients
- Parameter initialization

- Neural Networks in Practice

- Feed-Forward Neural Networks
- Regularization of Neural Networks
- Neural Network Optimization



UPPSALA
UNIVERSITET


- Practicalities
- **Introduction**
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Section 2

Introduction



Figure: ImageNet performance (Roessler, 2019)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

The Hype: Speech Recognition




Figure: Speech recognition performance (source: <https://eff.org/ai/metrics>)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

The Hype: Natural Language Processing

GLUE scores evolution over 2018-2019

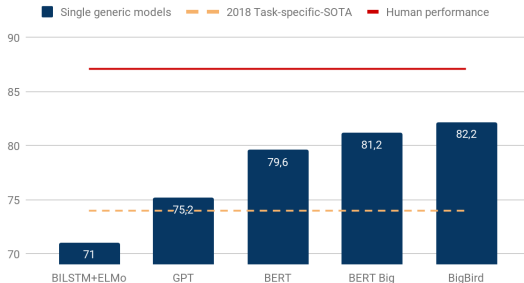


Figure: General Language Understanding (source: <https://www.programmersought.com/article/4251948498/>)

Work is very much ongoing:

<https://gluebenchmark.com/leaderboard>



UPPSALA
UNIVERSITET

The Hype

- Practicalities
- **Introduction**
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Although - Neural Networks is not a silver bullet



UPPSALA
UNIVERSITET

The Hype

- Practicalities
- **Introduction**
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Although - Neural Networks is not a silver bullet
- Remember the **Bayes error**



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Although - Neural Networks is not a silver bullet
- Remember the **Bayes error**
- Some times a linear regression (or Random Forest) is enough



UPPSALA
UNIVERSITET

- Practicalities
- Introduction
- **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Section 3

Feed-Forward Neural Networks



- Practicalities
- Introduction
- **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

The Feed-Forward Network




Figure 6.2: An example of a feedforward network, drawn in two different styles. Specifically, this is the feedforward network we use to solve the XOR example. It has a single hidden layer containing two units. (*Left*) In this style, we draw every unit as a node in the graph. This style is very explicit and unambiguous but for networks larger than this example it can consume too much space. (*Right*) In this style, we draw a node in the graph for each entire vector representing a layer's activations. This style is much more compact.

Figure: A simple feed-forward network (Goodfellow et al, 2017)

Important concepts:

Layers, neurons, input, output, weights, bias, architecture



Different Architectures for Different Purposes

- Practicalities
- Introduction
- **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Different networks for different purposes
 - **Convolutional Neural Networks:** Computer Vision



Different Architectures for Different Purposes

- Practicalities
- Introduction
- **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Different networks for different purposes
 - **Convolutional Neural Networks**: Computer Vision
 - **Recurrent Neural Networks**: Speech Audio (?)



Different Architectures for Different Purposes

- Practicalities
 - Introduction
 - **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Different networks for different purposes
 - **Convolutional Neural Networks**: Computer Vision
 - **Recurrent Neural Networks**: Speech Audio (?)
 - **Transformers/Attention**: Textual data
 - The Neural Network Zoo: <https://www.asimovinstitute.org/neural-network-zoo/>



UPPSALA
UNIVERSITET

Areas of Use: All fields

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Supervised learning
- Unsupervised learning
- Reinforcement learning



UPPSALA
UNIVERSITET

Why and when neural nets?

- Practicalities
 - Introduction
 - **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Learning feature representations
 - Needs a lot of data to learn complex representations
 - Good for sensor data (high-dimensional)



UPPSALA
UNIVERSITET

Why and when neural nets?

- Practicalities
- Introduction
- **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Learning feature representations
- Needs a lot of data to learn complex representations
- Good for sensor data (high-dimensional)
- When should we not use neural networks?



- Practicalities
- Introduction
- **Feed-Forward Neural Networks**
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Learning Representations

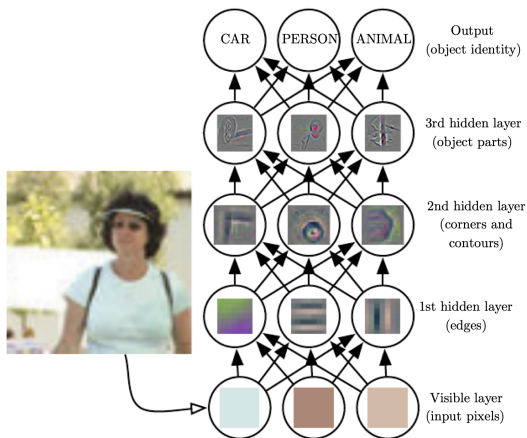


Figure: Learning representations can be crucial (Goodfellow et al, 2017, Fig. 1.2)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

The Feed-Forward Network




Figure 6.2: An example of a feedforward network, drawn in two different styles. Specifically, this is the feedforward network we use to solve the XOR example. It has a single hidden layer containing two units. (Left) In this style, we draw every unit as a node in the graph. This style is very explicit and unambiguous but for networks larger than this example it can consume too much space. (Right) In this style, we draw a node in the graph for each entire vector representing a layer's activations. This style is much more compact.

Figure: A simple feed-forward network (Goodfellow et al, 2017, Fig. 6.2)

In mathematical notation:

$$y_i = \mathbf{w}^T g(\mathbf{W}^T \mathbf{x}_i + \mathbf{b}_1) + \mathbf{b}_2$$



$$y_i = \mathbf{w}^T g(\mathbf{W}^T \mathbf{x}_i + \mathbf{b}_1) + \mathbf{b}_2$$

$$W = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}, w = \begin{pmatrix} 1 \\ -2 \end{pmatrix}, b_1 = \begin{pmatrix} 1 \\ -1 \end{pmatrix}, b_2 = (0)$$

$$g(z) = \text{ReLU}(z) = \max(0, z), x_i = \begin{pmatrix} 0 \\ 0 \end{pmatrix},$$

$$y_i = \begin{pmatrix} 1 \\ -2 \end{pmatrix}^T g\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 1 \\ -1 \end{pmatrix}\right) + (0)$$

$$y_i = \begin{pmatrix} 1 \\ -2 \end{pmatrix}^T \begin{pmatrix} 1 \\ 0 \end{pmatrix} + (0) = 1$$

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



The Feed-Forward Network

A feed-forward network for one observation (x_i).

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

$$\underbrace{\mathbf{h}_1}_{1 \times k_1} = g_1 \left(\underbrace{\mathbf{x}^T}_{1 \times p} \underbrace{\mathbf{W}_1}_{p \times k_1} + \underbrace{\mathbf{b}_1}_{1 \times k_1} \right)$$

\vdots

$$\underbrace{\mathbf{h}_l}_{1 \times k_l} = g_l \left(\underbrace{\mathbf{h}_{l-1}^T}_{1 \times k_{l-1}} \underbrace{\mathbf{W}_l}_{k_{l-1} \times k_l} + \underbrace{\mathbf{b}_l}_{1 \times k_l} \right)$$

\vdots

$$\underbrace{\hat{\mathbf{y}}}_{1 \times m} = g_L \left(\underbrace{\mathbf{h}_{L-1}^T}_{1 \times k_{L-1}} \underbrace{\mathbf{W}_L}_{k_{L-1} \times m} + \underbrace{\mathbf{b}_L}_{1 \times m} \right)$$

$$\hat{y} = f_L(f_{L-1}(\dots f_1(x)\dots))$$



Activation functions (g_l)

- Sometimes use notation σ as in $\sigma(Wh + b)$

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



Activation functions (g_l)

- Sometimes use notation σ as in $\sigma(Wh + b)$
- Historically $g(z)$ has been the sigmoid or or hyperbolic tangent (\tanh)

$$g_{\text{sigmoid}}(z) = \frac{e^z}{e^z + 1} = \frac{1}{1 + e^{-z}}$$

$$g_{\text{tanh}}(z) = \frac{\sinh z}{\cosh z} = \frac{e^{2z} - 1}{e^{2z} + 1}$$

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



Activation functions (g_l)

- Sometimes use notation σ as in $\sigma(Wh + b)$
- Historically $g(z)$ has been the sigmoid or or hyperbolic tangent (\tanh)

$$g_{\text{sigmoid}}(z) = \frac{e^z}{e^z + 1} = \frac{1}{1 + e^{-z}}$$

$$g_{\text{tanh}}(z) = \frac{\sinh z}{\cosh z} = \frac{e^{2z} - 1}{e^{2z} + 1}$$

- Now, usually variants of Rectified linear unit (ReLU)

$$g_{\text{ReLU}}(z) = \max(0, z)$$

- Easier to estimate with SGD
- Easier for deep models
- Last activation is the output function g_L , usually a softmax (if classification)

$$f(z_i) = \frac{e^{z_i}}{\sum_{j=1}^J e^{z_j}}$$

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



Activation functions (g_l)




Figure: Rectified Linear Unit (Goodfellow et al, 2017, Fig. 6.3)

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



Universal Approximation Theorem

"A feed-forward neural network with a linear output layer and at least one hidden layer with any 'squashing' activation function can approximate any Borel measurable function from one finite-dimensional space to another with any desired non-zero amount of error, provided that the network is given enough hidden units." (Goodfellow et al. 2017, p. 198)

- Also holds for ReLU
- No guarantee we can learn the network
- No guarantee that it will generalize
- No indication of how large the network need to be

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



UPPSALA
UNIVERSITET

Hyper-parameters in feed-forward networks

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- The number of layers
- The number of neurons
- Activation functions
- The type of layers (CNN, MaxPooling, Multi-head attention)



UPPSALA
UNIVERSITET

How to choose parameters

- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Trial and error on validation sets
 - Art rather than science
 - Specialized approaches (Bayesian Optimization)



How to choose parameters

- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Trial and error on validation sets
 - Art rather than science
 - Specialized approaches (Bayesian Optimization)
 - Grid search (combinatorial explosion)
 - Really bad with many parameters with less effects
 - If we have 5 irrelevant parameters we try 3 values for: 125 training per relevant run
 - Instead use...



How to choose parameters

- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Trial and error on validation sets
 - Art rather than science
 - Specialized approaches (Bayesian Optimization)
 - Grid search (combinatorial explosion)
 - Really bad with many parameters with less effects
 - If we have 5 irrelevant parameters we try 3 values for: 125 training per relevant run
 - Instead use...
 - Random search



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Grid search vs. Random Search

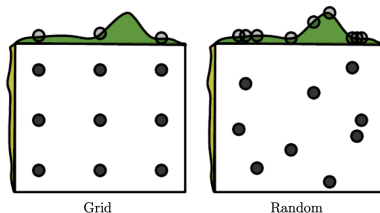


Figure: Grid search and random search (Goodfellow et al, 2017, Fig. 11.2)



UPPSALA
UNIVERSITET

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Section 4

Regularization



UPPSALA
UNIVERSITET

Regularization of Neural Networks

- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Reduce training error but improve test/validation error



UPPSALA
UNIVERSITET

Regularization of Neural Networks

- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Reduce training error but improve test/validation error
 - Neural Networks are extremely flexible / high model capacity



UPPSALA
UNIVERSITET

Regularization of Neural Networks

- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Reduce training error but improve test/validation error
 - Neural Networks are extremely flexible / high model capacity
 - Regularization is crucial for good generalizability of NN



UPPSALA
UNIVERSITET

Regularization of Neural Networks

- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Reduce training error but improve test/validation error
 - Neural Networks are extremely flexible / high model capacity
 - Regularization is crucial for good generalizability of NN



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- **Regularization**
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Let

$$\tilde{J}(W, b) = J(W, b) + \alpha\Omega(W),$$

where $J(W, b)$ is the cost function and $\alpha\Omega(W)$ is the penalty for the weight matrices.

- α is the strength of the penalty.



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Let

$$\Omega_1(W) = \sum_i \sum_j |w|_{i,j},$$

and

$$\Omega_2(W) = \sum_i \sum_j w_{i,j}^2,$$

be the L_1 and L_2 regularization respectively.

- We can then get the cost function

$$\tilde{J}(W, b) = J(W, b) + \sum_l \alpha_l \Omega_2(W_l),$$



- Lets define the cost function as

$$\begin{aligned}\tilde{J}(w) &= J(w) + \alpha \Omega_2(w) \\ &= J(w) + \alpha w^T w\end{aligned}$$

- Then the gradient update becomes

$$\nabla_w \tilde{J}(w) = \nabla_w J(w) + 2\alpha w$$

- To update our weights with gradient descent

$$\begin{aligned}w &\leftarrow w - \epsilon(\nabla_w J(w) + 2\alpha w) \\ w &\leftarrow (1 - 2\alpha\epsilon)w - \epsilon\nabla_w J(w)\end{aligned}$$

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

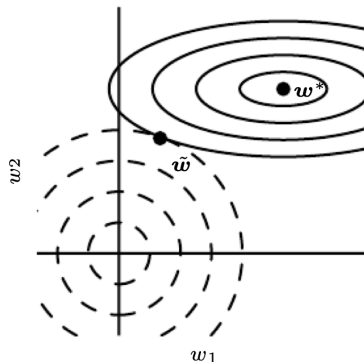


Figure: L_2 regularization (Goodfellow et al, 2017, Fig. 7.1)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Early Stopping

- Stop optimization early based on validation error
- Rerun to that number of epochs (hyperparameter)
- Can be shown to be equivalent (under strict assumptions) to L_2 regularization

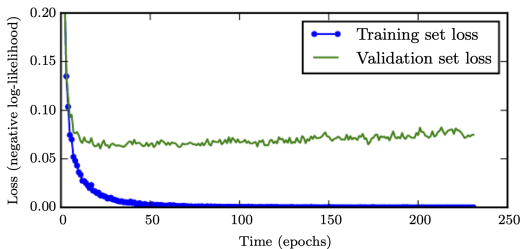


Figure: Early Stopping (Goodfellow et al, 2017, Fig. 7.3)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Early Stopping

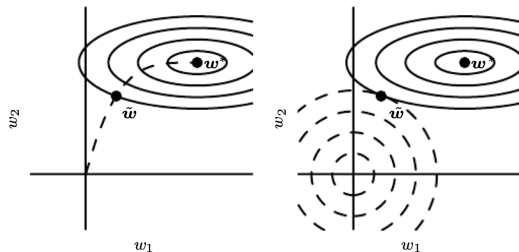


Figure: Early Stopping (Goodfellow et al, 2017, Fig. 7.4)



- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- In each iteration:
 - Sample an indicator I_i for each node i
 - Set the value h_i to 0 with probability p
 - The dropout probability is typically 0.8 for input nodes and 0.5 for hidden nodes
 - Forces the network to
 - not rely on individual nodes
 - spread out the weights over more nodes
 - Can be seen as an ensemble method



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Dropout




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

Figure: Dropout (Srivastava et al, 2014)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- In CNN: Dataset augmentation
- Get more data...



UPPSALA
UNIVERSITET

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Section 5

Optimization of Neural Networks



UPPSALA
UNIVERSITET

Neural Network Learning

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Usually, a lot of data and many parameters ($\theta = (W, b)$)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Usually, a lot of data and many parameters ($\theta = (W, b)$)
- We usually minimize our training cost function

$$J(\theta) = \sum_i^N L(\text{NN}(x_i), y_i) + \Omega(\theta),$$

where L is the observation level loss, $\text{NN}()$ is our neural network and Ω is the regularization term.



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Usually, a lot of data and many parameters ($\theta = (W, b)$)
- We usually minimize our training cost function

$$J(\theta) = \sum_i^N L(\text{NN}(x_i), y_i) + \Omega(\theta),$$

where L is the observation level loss, $\text{NN}()$ is our neural network and Ω is the regularization term.

- **Learning Target:** Find $\hat{\theta}$ that minimize the **generalization** error



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Usually, a lot of data and many parameters ($\theta = (W, b)$)
- We usually minimize our training cost function

$$J(\theta) = \sum_i^N L(\text{NN}(x_i), y_i) + \Omega(\theta),$$

where L is the observation level loss, $\text{NN}()$ is our neural network and Ω is the regularization term.

- **Learning Target:** Find $\hat{\theta}$ that minimize the **generalization** error



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Stochastic Gradient Descent methods

$$\theta_t = \theta_{t-1} - \eta_t \hat{\nabla} J(\theta_{t-1})$$

- In practice, better optimizers are used:
 - Adam
 - RMSprop



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Stochastic Gradient Descent methods

$$\theta_t = \theta_{t-1} - \eta_t \hat{\nabla} J(\theta_{t-1})$$

- In practice, better optimizers are used:
 - Adam
 - RMSprop
- To compute gradients ($\hat{\nabla} J(\theta_{t-1})$):
 - **Backpropagation** algorithm (chain-rule for derivatives)



- Practicalities
 - Introduction
 - Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
 - Regularization
 - Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
 - Neural Networks in Practice
- Neural Networks cost function $J(\theta)$ are (usually) **not a convex** function
 - We can have local minima
 - **When will this be a problem?**



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Neural Networks cost function $J(\theta)$ are (usually) **not a convex** function
- We can have local minima
- **When will this be a problem?**
- A problem if local optima has a **high cost**



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Neural Networks cost function $J(\theta)$ are (usually) **not a convex** function
- We can have local minima
- **When will this be a problem?**
- A problem if local optima has a **high cost**
- **Good thing:** This is probably not the situation!



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Neural Networks cost function $J(\theta)$ are (usually) **not a convex** function
- We can have local minima
- **When will this be a problem?**
- A problem if local optima has a **high cost**
- **Good thing:** This is probably not the situation!
- Many practitioners believe this is a problem



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Neural Networks cost function $J(\theta)$ are (usually) **not a convex** function
- We can have local minima
- **When will this be a problem?**
- A problem if local optima has a **high cost**
- **Good thing:** This is probably not the situation!
- Many practitioners believe this is a problem
- **Diagnosis:** Plot the gradient norm, $\mathbf{g}^T \mathbf{g}$, where $\mathbf{g} = \hat{\nabla} J(\theta)$



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Neural Networks cost function $J(\theta)$ are (usually) **not a convex** function
- We can have local minima
- **When will this be a problem?**
- A problem if local optima has a **high cost**
- **Good thing:** This is probably not the situation!
- Many practitioners believe this is a problem
- **Diagnosis:** Plot the gradient norm, $\mathbf{g}^T \mathbf{g}$, where $\mathbf{g} = \hat{\nabla} J(\theta)$



Problems: Plateaus and Saddle Points

- Another problem is **saddle points**

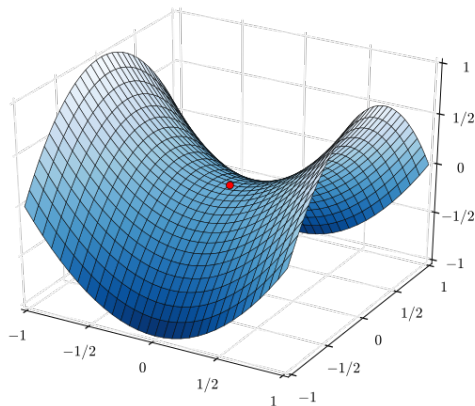


Figure: Saddle point for $z = x^2 - y^2$ (Wikipedia)

- What is the gradient in a saddle point?



Problems: Plateaus and Saddle Points

- Another problem is **saddle points**

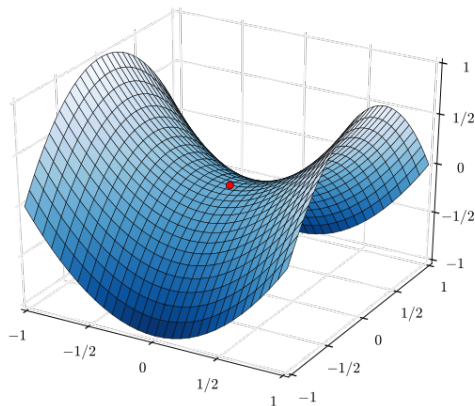


Figure: Saddle point for $z = x^2 - y^2$ (Wikipedia)

- What is the gradient in a saddle point?



Problems: Plateaus and Saddle Points

- In a saddle point the Hessian is **indefinite**, both positive and negative eigen values

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - **Plateaus and Saddle Points**
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Problems: Plateaus and Saddle Points

- In a saddle point the Hessian is **indefinite**, both positive and negative eigen values
- A local minimum: Hessian is **positive definite**, only positive eigen values



Problems: Plateaus and Saddle Points

- In a saddle point the Hessian is **indefinite**, both positive and negative eigen values
- A local minimum: Hessian is **positive definite**, only positive eigen values
- In random functions of **high dimension**: **most points are saddle points**
- **Intuition**: In random functions the sign of the eigen values of the Hessian is random

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - **Plateaus and Saddle Points**
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



Problems: Plateaus and Saddle Points

- In a saddle point the Hessian is **indefinite**, both positive and negative eigen values
- A local minimum: Hessian is **positive definite**, only positive eigen values
- In random functions of **high dimension**: **most points are saddle points**
- **Intuition**: In random functions the sign of the eigen values of the Hessian is random
- In random functions: eigen values become more positive in regions of lower cost. I.e. local minima more probable in low cost regions. Saddle points more common in high-cost regions

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - **Plateaus and Saddle Points**
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Problems: Plateaus and Saddle Points

- In a saddle point the Hessian is **indefinite**, both positive and negative eigen values
- A local minimum: Hessian is **positive definite**, only positive eigen values
- In random functions of **high dimension**: **most points are saddle points**
- **Intuition**: In random functions the sign of the eigen values of the Hessian is random
- In random functions: eigen values become more positive in regions of lower cost. I.e. local minima more probable in low cost regions. Saddle points more common in high-cost regions
- The problem of saddle points:
 - Probably a reason why second order methods (using the Hessian) has not succeed
 - Empirically, (stochastic) gradient descent seem to escape saddle points



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Problems: Cliffs

- Another problem is "cliffs" or large changes in gradients




Figure: Cliff (Goodfellow et al., 2017, Fig. 8.3)

- Can undo many iterative steps
- Common in Recurrent neural networks
- Mitigation: Gradient clipping



Problems: Exploding and vanishing gradients

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- In deep neural networks, gradients can **vanish or explode**
- As an example: We want to compute the gradient for a situation where the weights are multiplied t times. Then using eigendecomposition

$$W^t = V \text{diag}(\lambda)^t V^T$$

- The gradient is scaled wrt $\text{diag}(\lambda)^t$



Problems: Exploding and vanishing gradients

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- In deep neural networks, gradients can **vanish or explode**
- As an example: We want to compute the gradient for a situation where the weights are multiplied t times. Then using eigendecomposition

$$W^t = V \text{diag}(\lambda)^t V^T$$

- The gradient is scaled wrt $\text{diag}(\lambda)^t$
- Cliffs is an example of gradient explosion
- **Mitigation for gradient explosion:** Gradient clipping



Problems: Exploding and vanishing gradients

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- In deep neural networks, gradients can **vanish or explode**
- As an example: We want to compute the gradient for a situation where the weights are multiplied t times. Then using eigendecomposition

$$W^t = V \text{diag}(\lambda)^t V^T$$

- The gradient is scaled wrt $\text{diag}(\lambda)^t$
- Cliffs is an example of gradient explosion
- **Mitigation for gradient explosion:** Gradient clipping
- Common problem in **Recurrent neural networks**



- We need to have **starting values** for gradient descent

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - **Parameter initialization**
- Neural Networks in Practice



UPPSALA UNIVERSITET

Initial values

- We need to have **starting values** for gradient descent
- Initialization can be seen as a hyperparameter

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - **Parameter initialization**
- Neural Networks in Practice



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- We need to have **starting values** for gradient descent
- Initialization can be seen as a hyperparameter
- Bad initial values might
 - Bad convergence (local optimum)
 - Numerical problems



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - **Parameter initialization**
- Neural Networks in Practice

- We need to have **starting values** for gradient descent
- Initialization can be seen as a hyperparameter
- Bad initial values might
 - Bad convergence (local optimum)
 - Numerical problems
- We want to **break symmetry** between layers
 - Otherwise the same units will be updated in the same way (deterministically)



- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - **Parameter initialization**
- Neural Networks in Practice

- We need to have **starting values** for gradient descent
- Initialization can be seen as a hyperparameter
- Bad initial values might
 - Bad convergence (local optimum)
 - Numerical problems
- We want to **break symmetry** between layers
 - Otherwise the same units will be updated in the same way (deterministically)
- Good practice
 - Initialize values randomly close to zero (uniform or normal)



UPPSALA
UNIVERSITET

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

Section 6

Neural Networks in Practice



UPPSALA UNIVERSITET

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

TensorFlow

- Framework for large-scale machine learning and Neural Networks
- Developed by Google
- Can be used both from R and Python
- Used in both research and production
- What Tensorflow does:
 - Computing gradients (autodiff) for Neural Networks
 - Enable use of graphical processing units (GPU) and Tensor processing Units (TPU)
 - Enable training using common optimizers (such as Adam, RMSprop)
- Tensorflow Probability is a probabilistic programming framework using TF





UPPSALA UNIVERSITET

(Py)Torch

- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Similar to TensorFlow
- Developed by Meta AI
- Can be used both from R and Python
- Used in both research and production
- **pyro** is a probabilistic programming framework using torch





- Practicalities
- Introduction
- Feed-Forward Neural Networks
 - Feed-Forward Neural Networks
 - Hyper-parameters
- Regularization
- Optimization of Neural Networks
 - Local minima
 - Plateaus and Saddle Points
 - Cliffs, exploding and vanishing gradients
 - Parameter initialization
- Neural Networks in Practice

- Syntax for 'building' Neural Networks
- Available both in R and Python
- TensorFlow or Torch as backend

