UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Machine learning, big data and artificial intelligence – Block 6

Måns Magnusson
Department of Statistics, Uppsala University

HT 2020

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

This week's lectures

- Word embedding basics
- Recurrent Neural Networks
- Attention and Transformers
- BERT

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

How do we represent words?

- One-hot encoding
  - A vector of length V (vocabulary size)

$$\text{Uppsala} = [0, ..., 1, ..., 0] = 1_i$$

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

How do we represent words?

- One-hot encoding
  - A vector of length V (vocabulary size)

$$\text{Uppsala} = [0, ..., 1, ..., 0] = 1_i$$

- Word embeddings
  - A vector of length D (embedding dimension)

$$\text{Uppsala} = [-0.1231, ..., 1.9001, ..., 0.012]$$

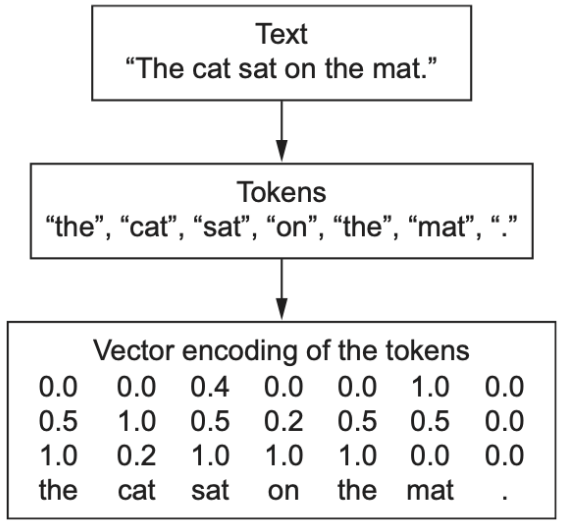

Figure: Representing words as word emnbeddings (Chollet and Allair, 2018, Fig. 6.1)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Word embeddings vs. One-Hot

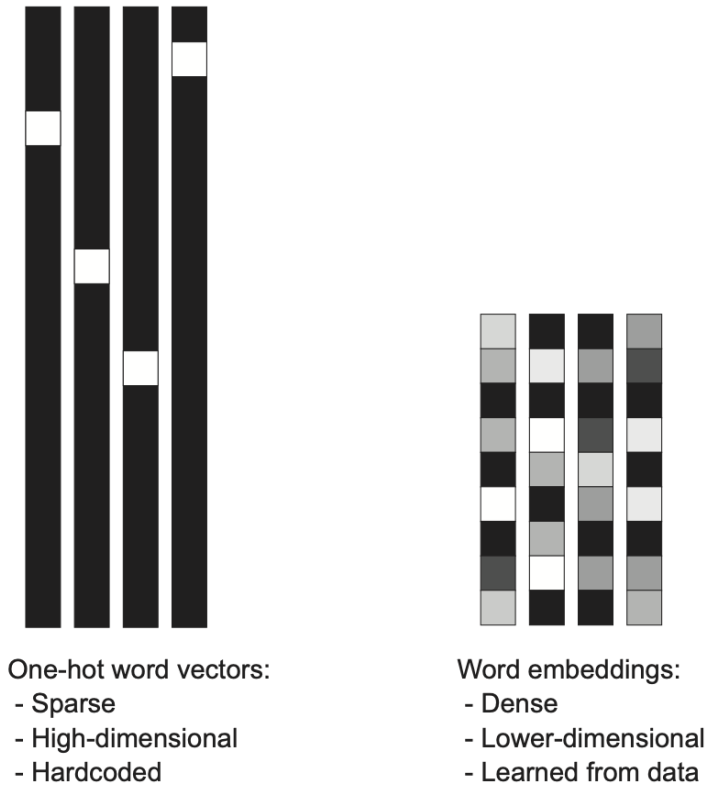

Figure: One-Hot vs. Word embeddings (Chollet and Allair, 2018, Fig. 6.2)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Word embeddings

- A word type represent meaning in a low-dimensional semanitic space Word embeddings
- The distributional hypothesis:
  - Harris (1954) and Firth (1957):
    "A word is characterized by the company it keeps"
  - Semantics (broadly defined) is captured by context
- Lots of different embeddings:
  `word2vec`, GloVe, Probabilistic Embeddings

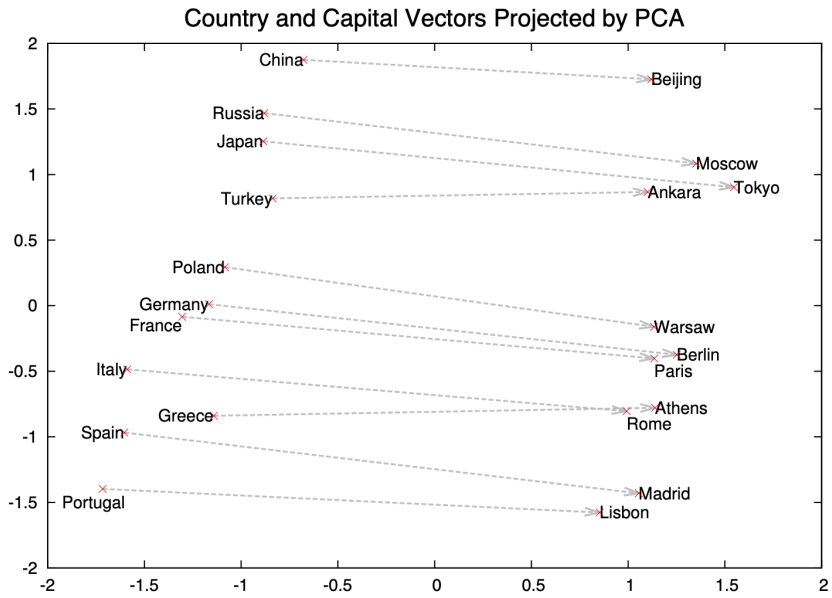

Figure: Word embedding properties (Mikolov et al, 2013)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Context Matters!

Figure: Context matters (Alammar, 2020)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Recurrent Neural Networks

- Recurrent Neural Networks, Recurrent Nets, RNN, ...
- Modeling of temporal data structures, such as
  - Time series data
  - Sequences of words (language models)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Recurrent Neural Networks

- Recurrent Neural Networks, Recurrent Nets, RNN, ...
- Modeling of temporal data structures, such as
  - Time series data
  - Sequences of words (language models)
- Examples of applications:
  - Text classification
  - Sequence / word classification
  - Time series predictions

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Recurrent Neural Networks

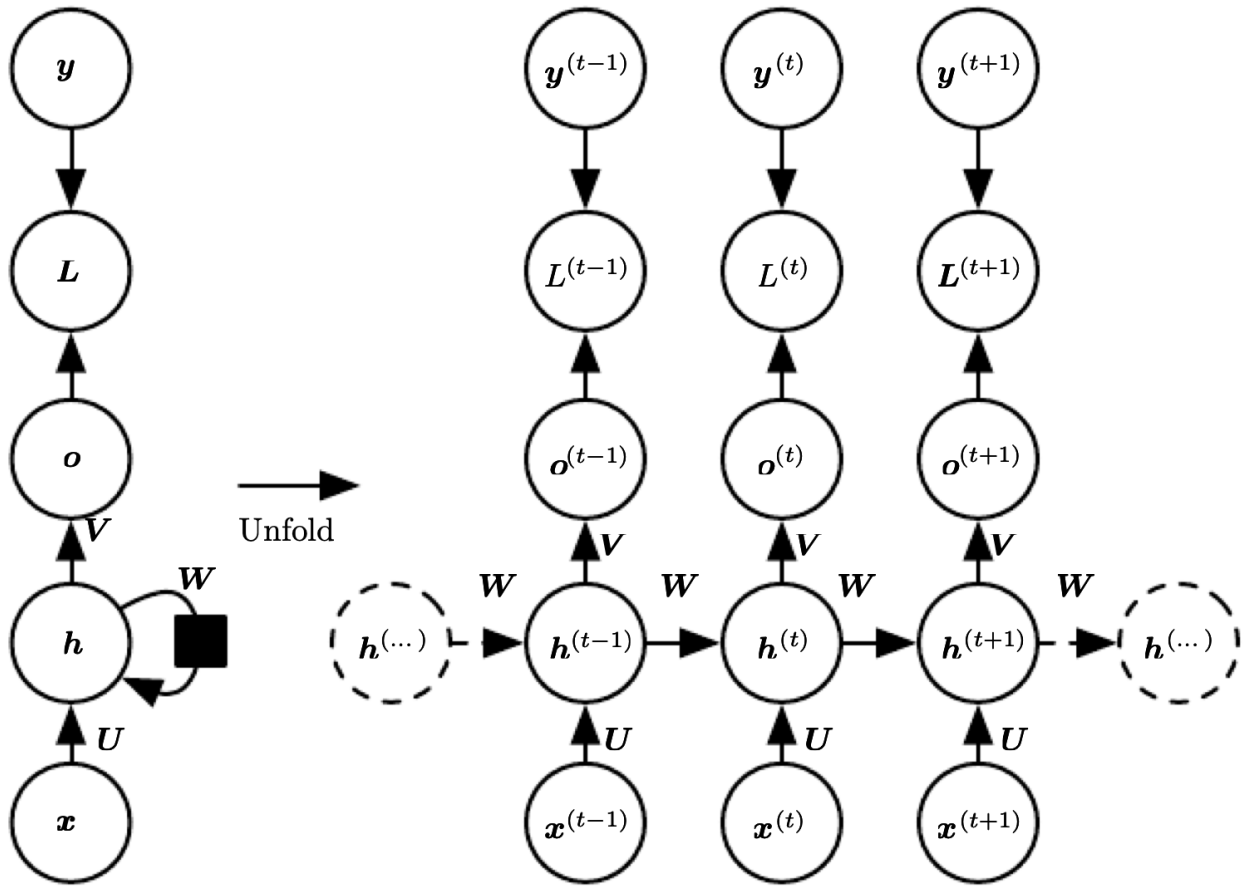

Figure: Recurrent Neural Network (Goodfellow et al, 2017, Fig. 10.3)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Recurrent Neural Networks

$$a_t = b_t + Wh_{t-1} + Ux_t$$
$$h_t = \sigma a_t$$
$$o_t = c + Vh_t$$
$$\hat{y}_t = \text{softmax}(o_t)$$

Think of h_t as the "state" at timepoint t

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Recurrent network with one output

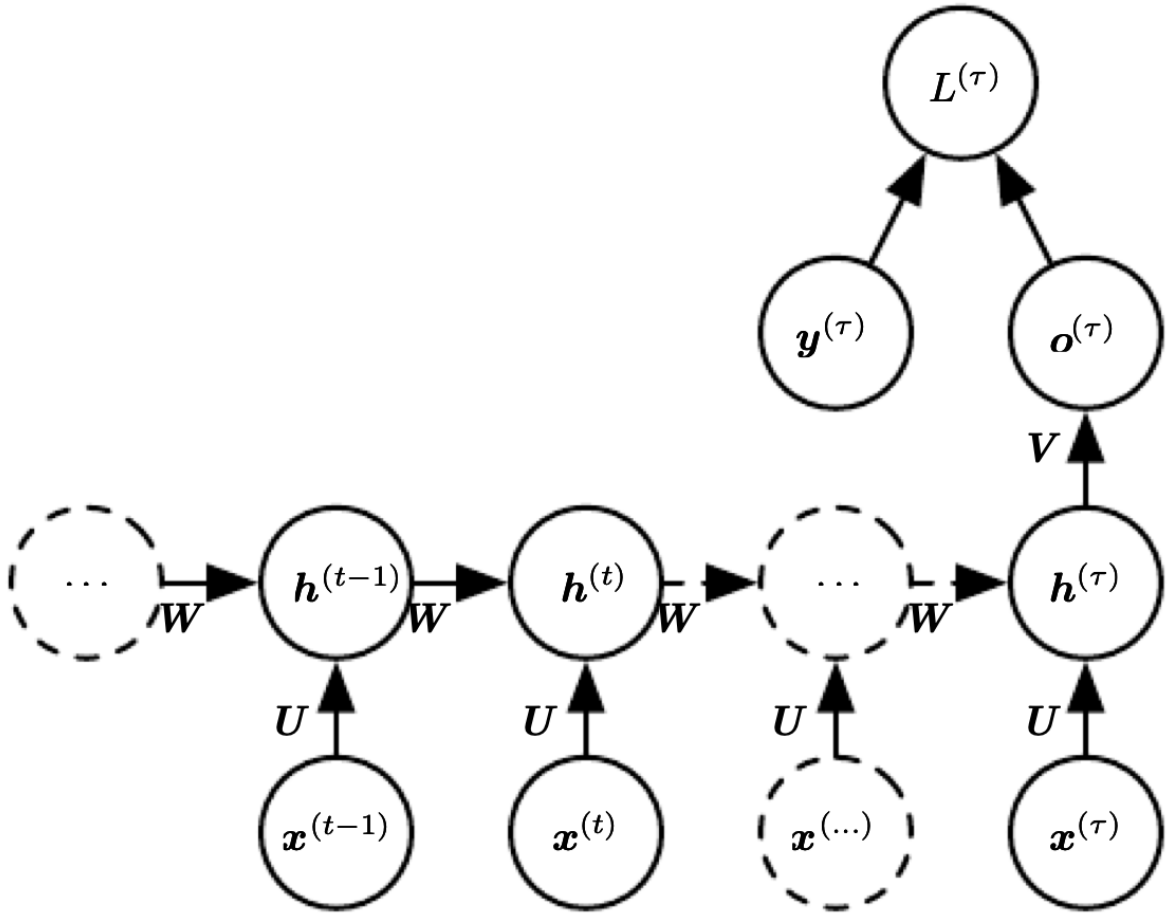

Figure: Recurrent Neural Network with one output (Goodfellow et al, 2017, Fig. 10.5)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Sequence to Sequence: Encoder-Decoder

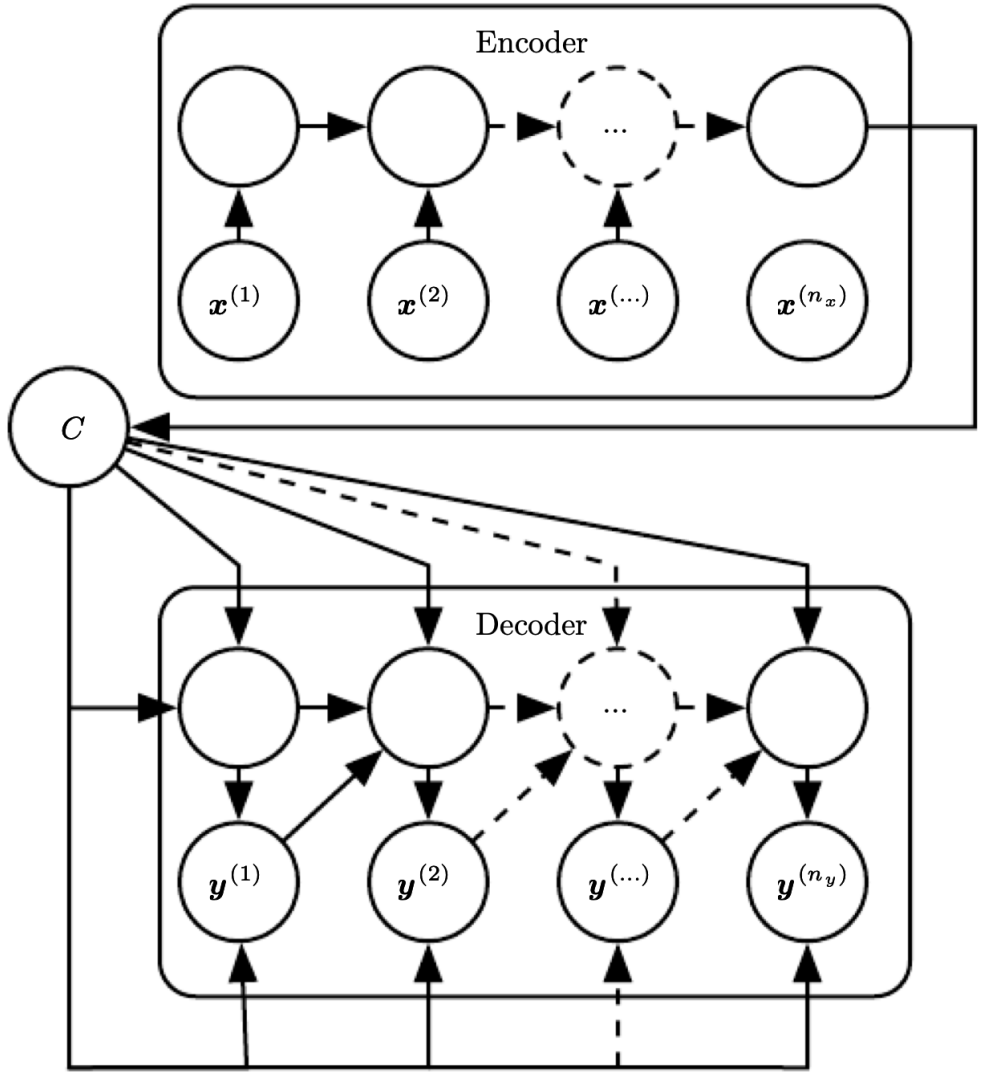

Figure: Encoder-Decoder Recurrent Networks (Goodfellow et al, 2017, Fig. 10.12)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Problems with RNN

- Predicting sequences of different lengths
- Exploding and vanishing gradients
- Long-term dependencies

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Bi-Directional RNN

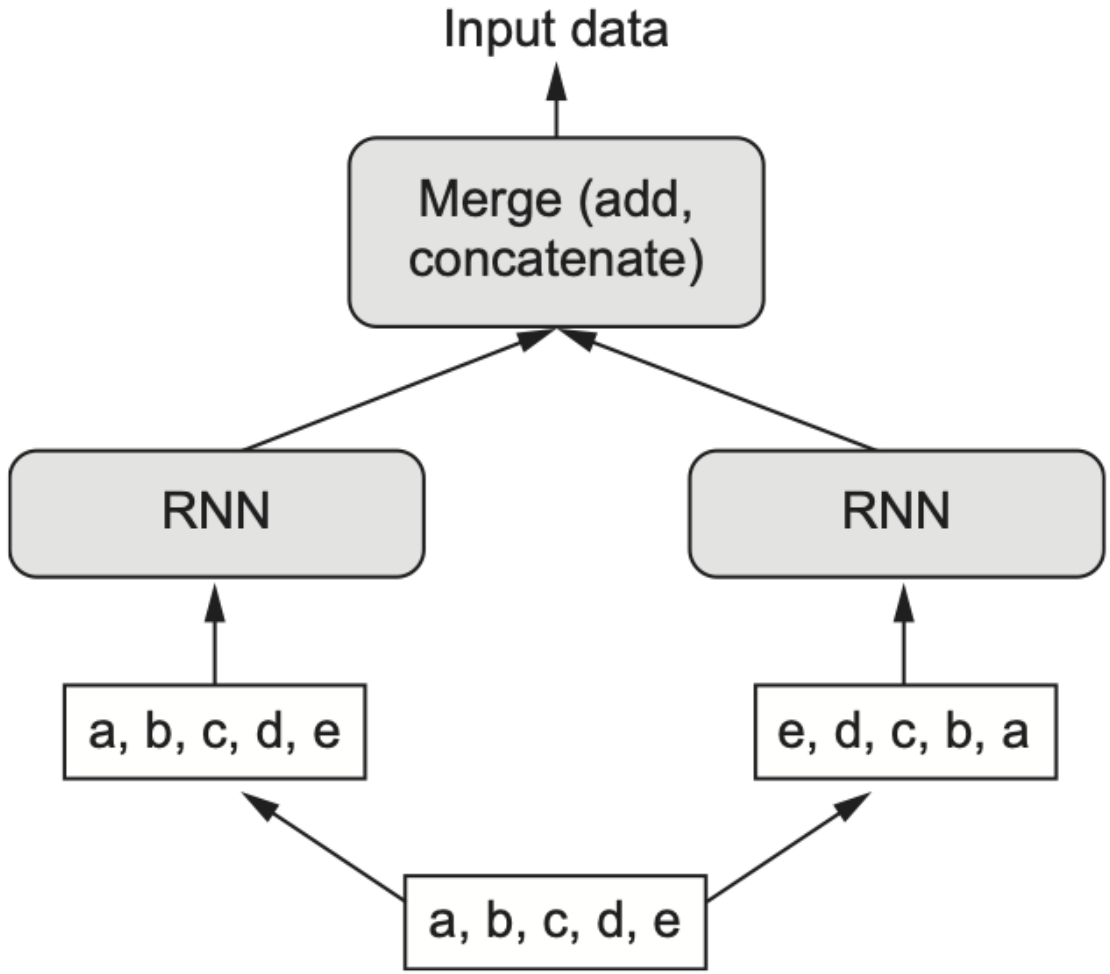

Figure: Bi-Directional RNN (Chollet and Allaire, 2018, Fig. 6.21)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Long Short-Term Memory (LSTM)

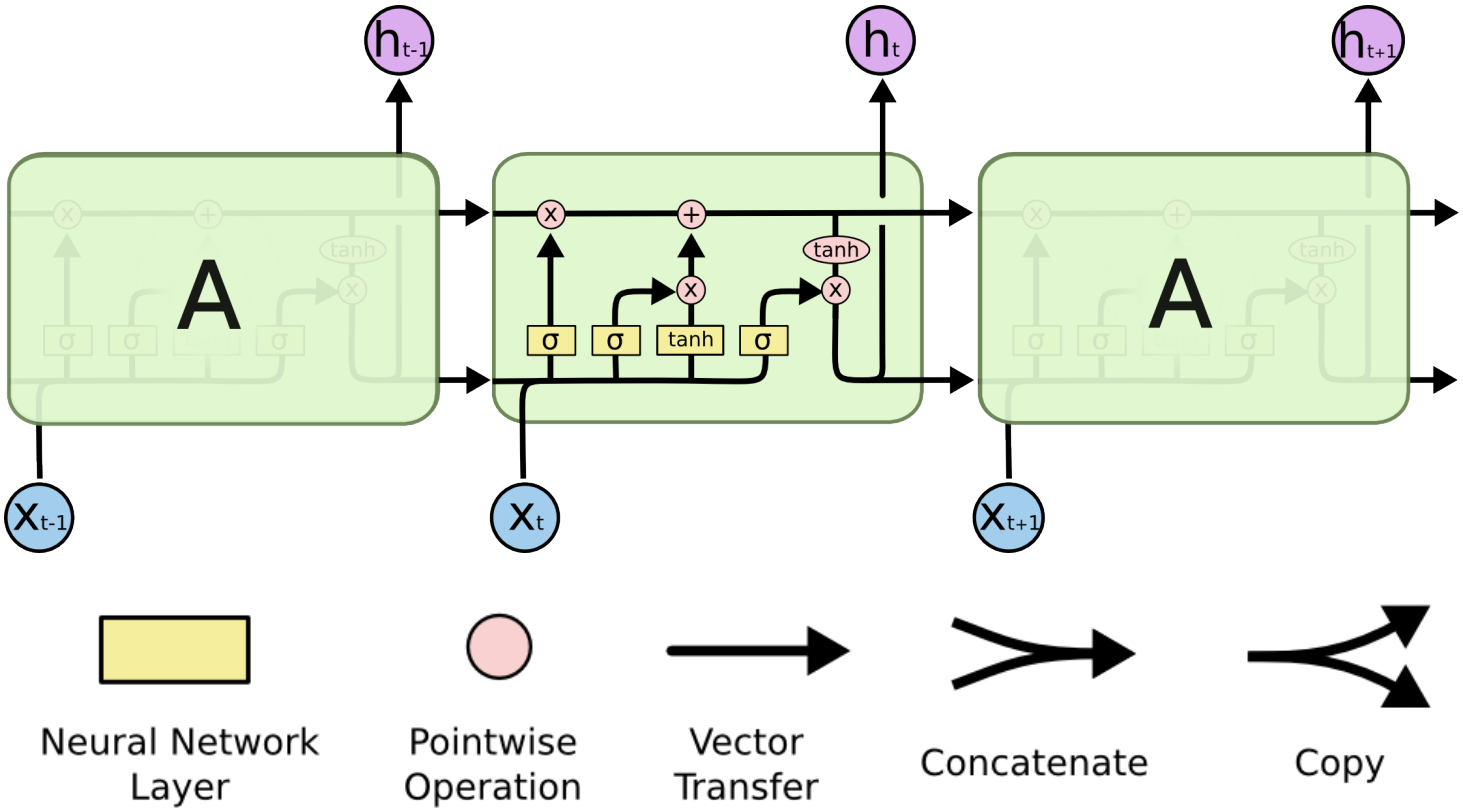

Figure: The LSTM (Olah, 2015)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

LSTM cell state

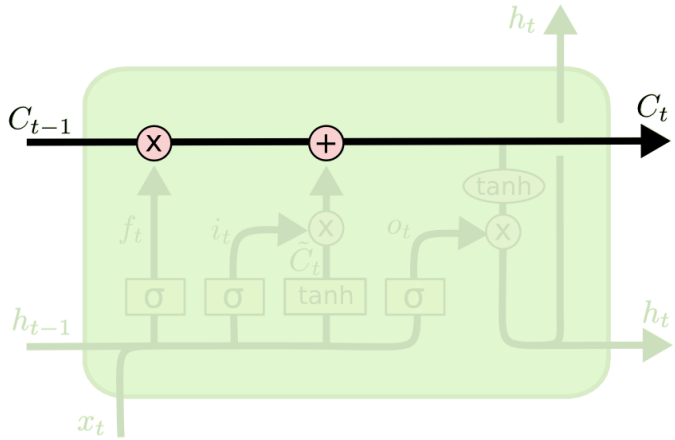

Figure: LSTM cell state, i.e. "carrybelt" (Olah, 2015)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

LSTM forget gate

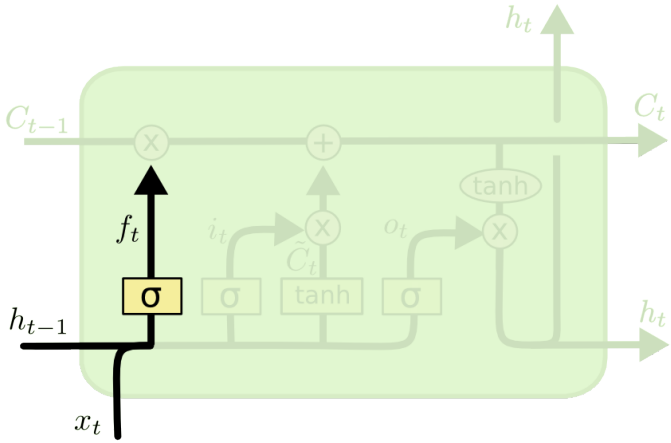

$$f_t = \sigma\left(W_f \cdot [h_{t-1}, x_t] \; + \; b_f\right)$$

Figure: LSTM forget gate (Olah, 2015)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

LSTM input gate

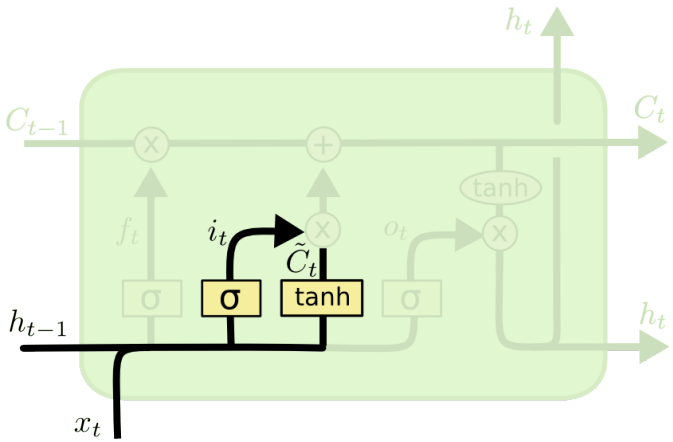

$$i_t = \sigma\left(W_i \cdot [h_{t-1}, x_t] \; + \; b_i\right)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] \; + \; b_C)$$

Figure: LSTM input gate (Olah, 2015)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

LSTM cell state update

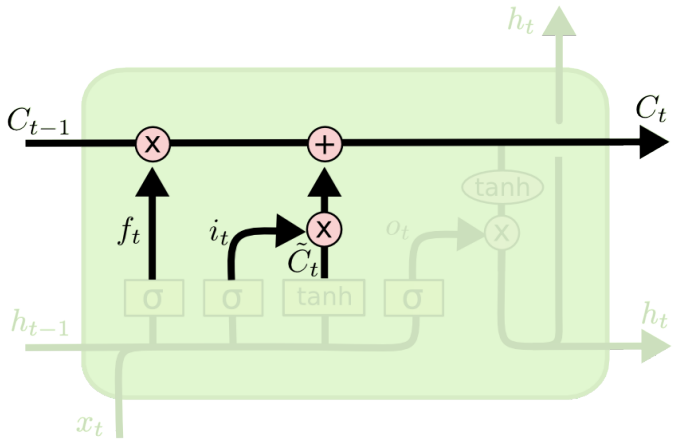

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Figure: Update cell state (Olah, 2015)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

LSTM output gate

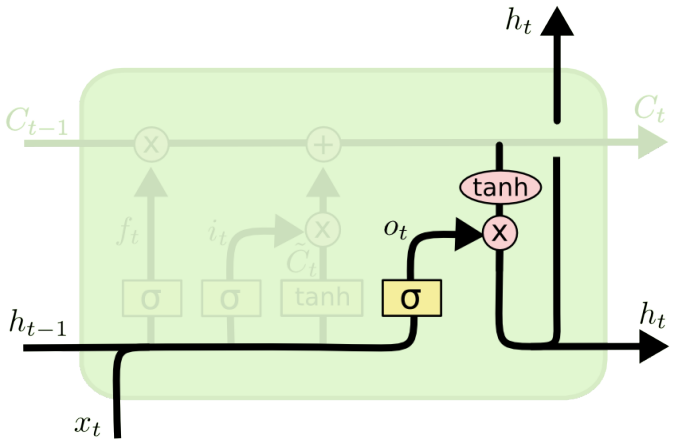

$$o_t = \sigma\left(W_o\left[h_{t-1}, x_t\right] + b_o\right)$$

$$h_t = o_t * \tanh\left(C_t\right)$$

Figure: LSTM output gate (Olah, 2015)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Problems

- Still a recurrent structure, (vanishing and exploding gradients)
- Long-term dependencies still difficult
- Hard to do transfer learning

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Transformer

- Introduced in 2017 in Vaswani et al. (2017)
- Behind the recent progress in NLP: BERT, GPT-2, GPT-3, etc.

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Transformer

- Introduced in 2017 in Vaswani et al. (2017)
- Behind the recent progress in NLP: BERT, GPT-2, GPT-3, etc.
- Becoming de-facto standard in industry and academia

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Transformer

- Introduced in 2017 in Vaswani et al. (2017)
- Behind the recent progress in NLP: BERT, GPT-2, GPT-3, etc.
- Becoming de-facto standard in industry and academia
- Brings transfer learning to NLP

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Transformer

- Introduced in 2017 in Vaswani et al. (2017)
- Behind the recent progress in NLP: BERT, GPT-2, GPT-3, etc.
- Becoming de-facto standard in industry and academia
- Brings transfer learning to NLP
- Two large benefits:
  - Enables more parallelism
  - Better handling of long-range dependencies

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

A Sequence-to-Sequence Model

Figure: Attention (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Stacked Encoder-Decoder Structure

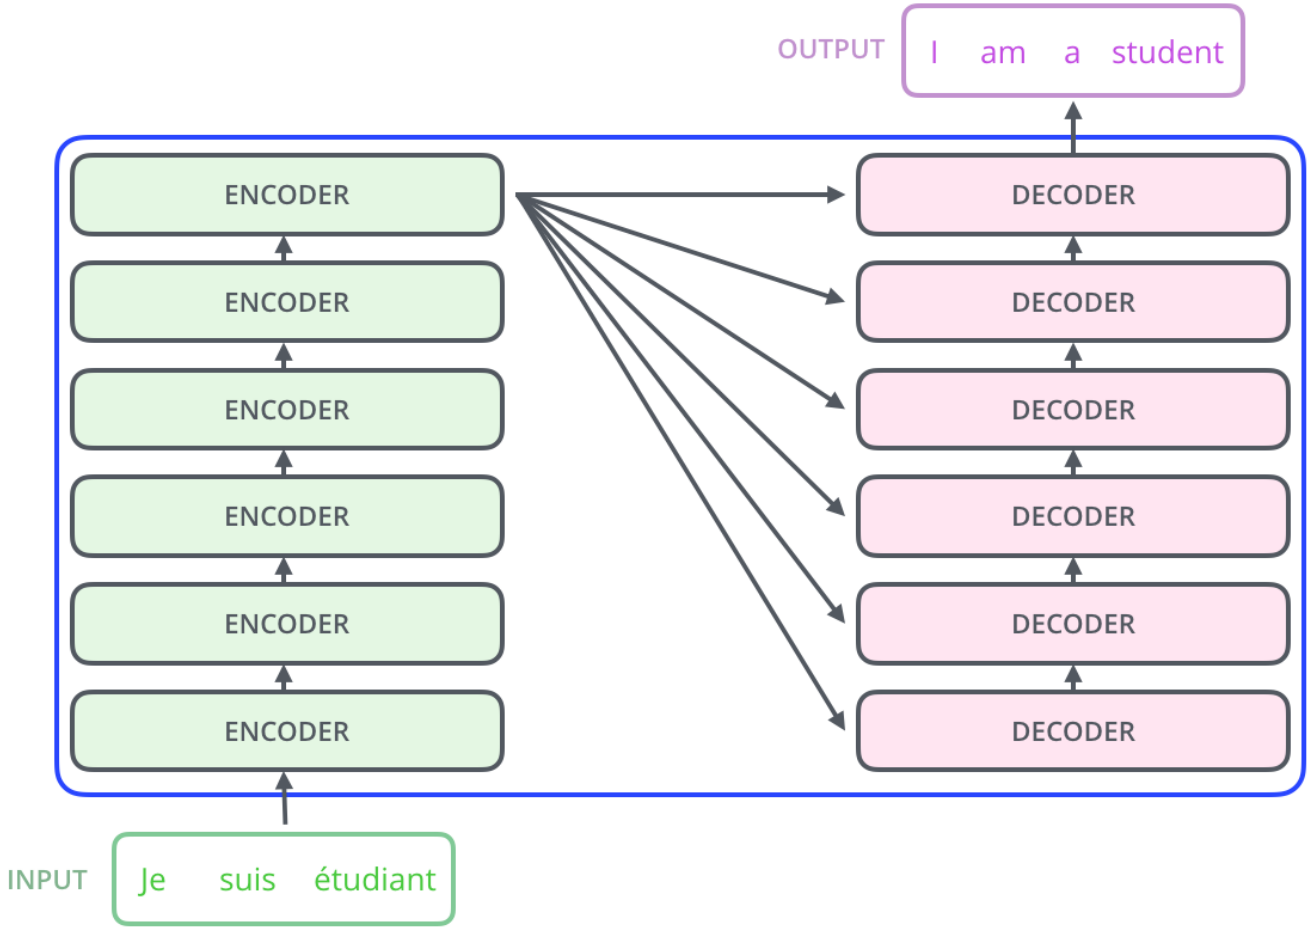

Figure: Attention (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Transformer

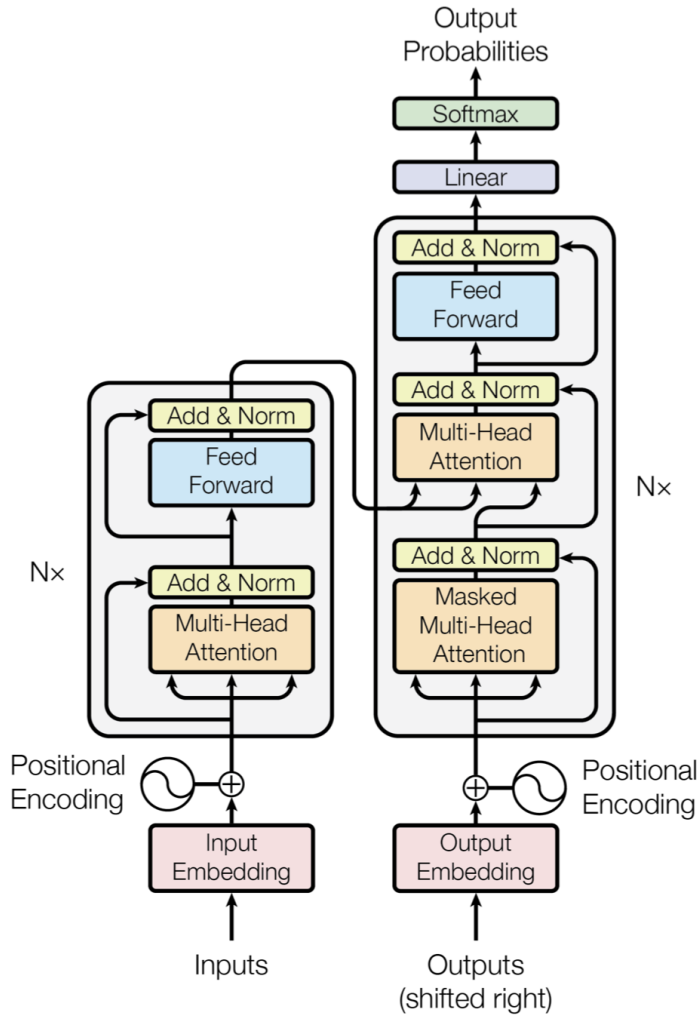

Figure: The Transformer Architecture (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

The encoder vs. the decoder

- Encoder:
  - Input: words (embeddings)
  - Output: contextualized embeddings
- Decoder:
  - Input: contextualized embeddings and previous words (embeddings)
  - Output: words (embeddings)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

The Encoder Layer

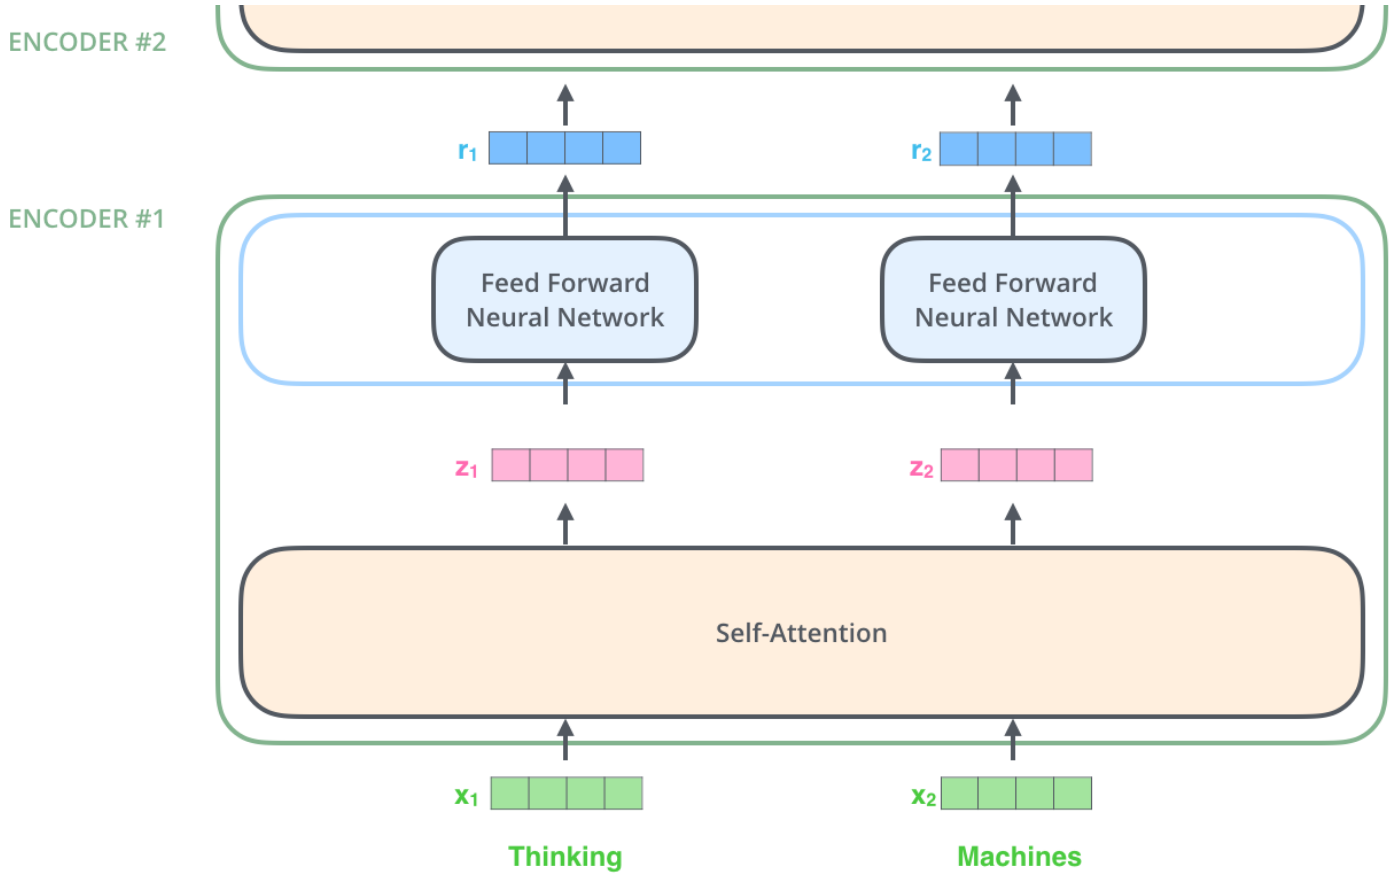

Figure: The Encoder Layer (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Scaled Dot-Product Attention

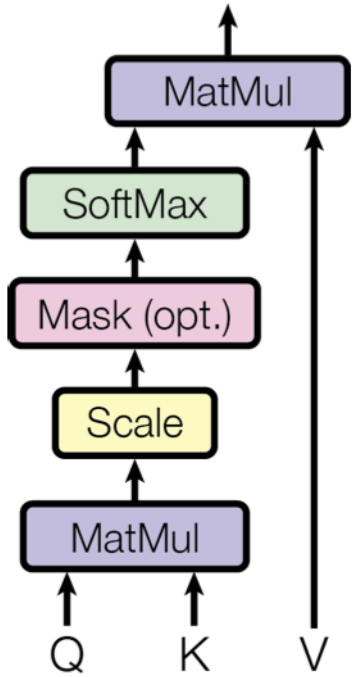

Figure: Scaled Dot-Product Attention (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

The encoder vs. the decoder

- (Q)uery: Word i query other words
- (K)ey: The other words return their key to i
- (V)alue: The value of the other words to i

Computing Q, V and K

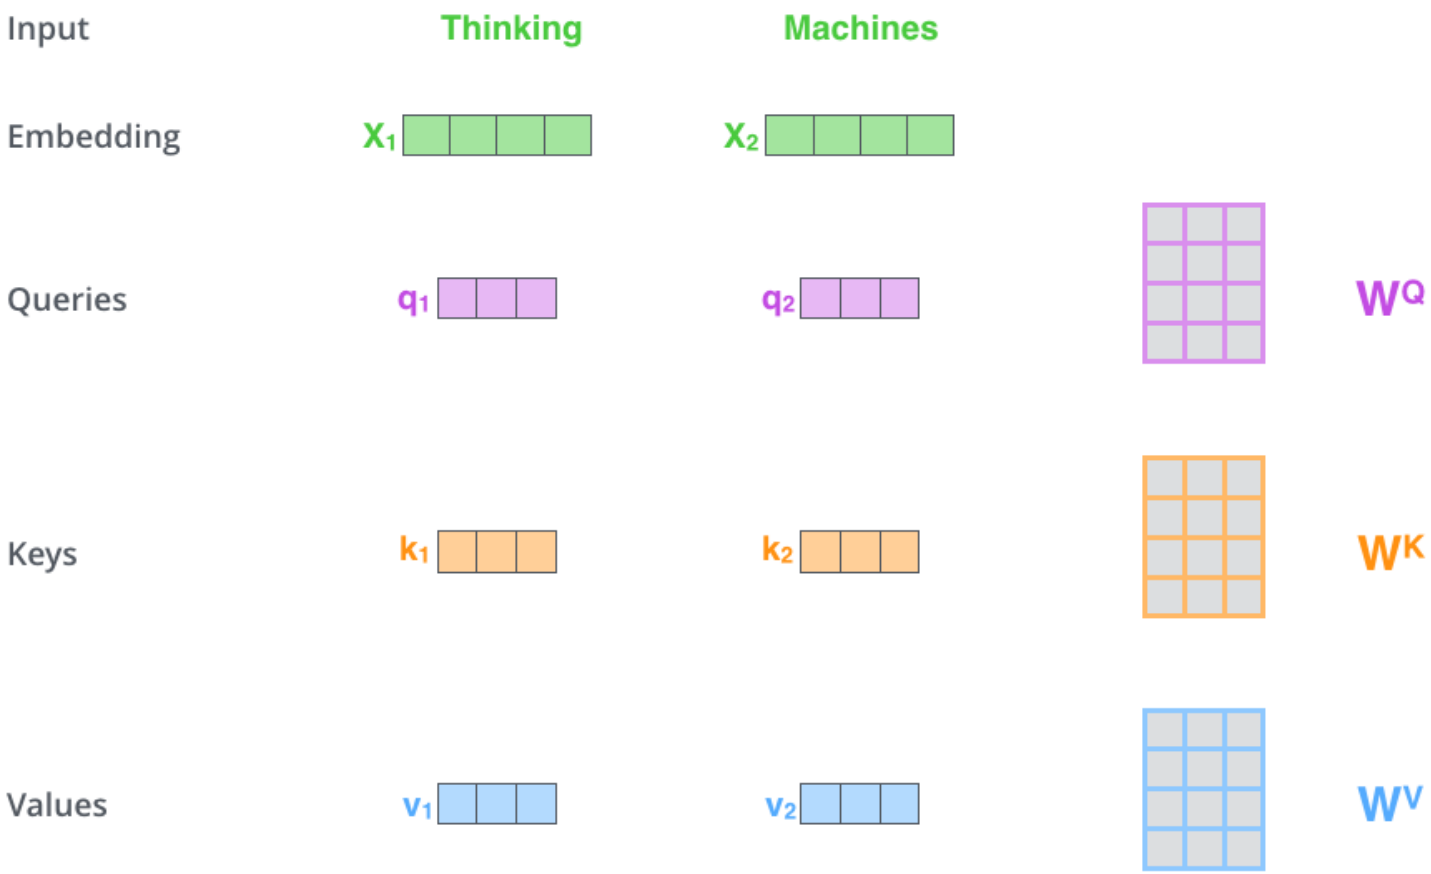

Figure: Attention heads (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Computing Self-Attention

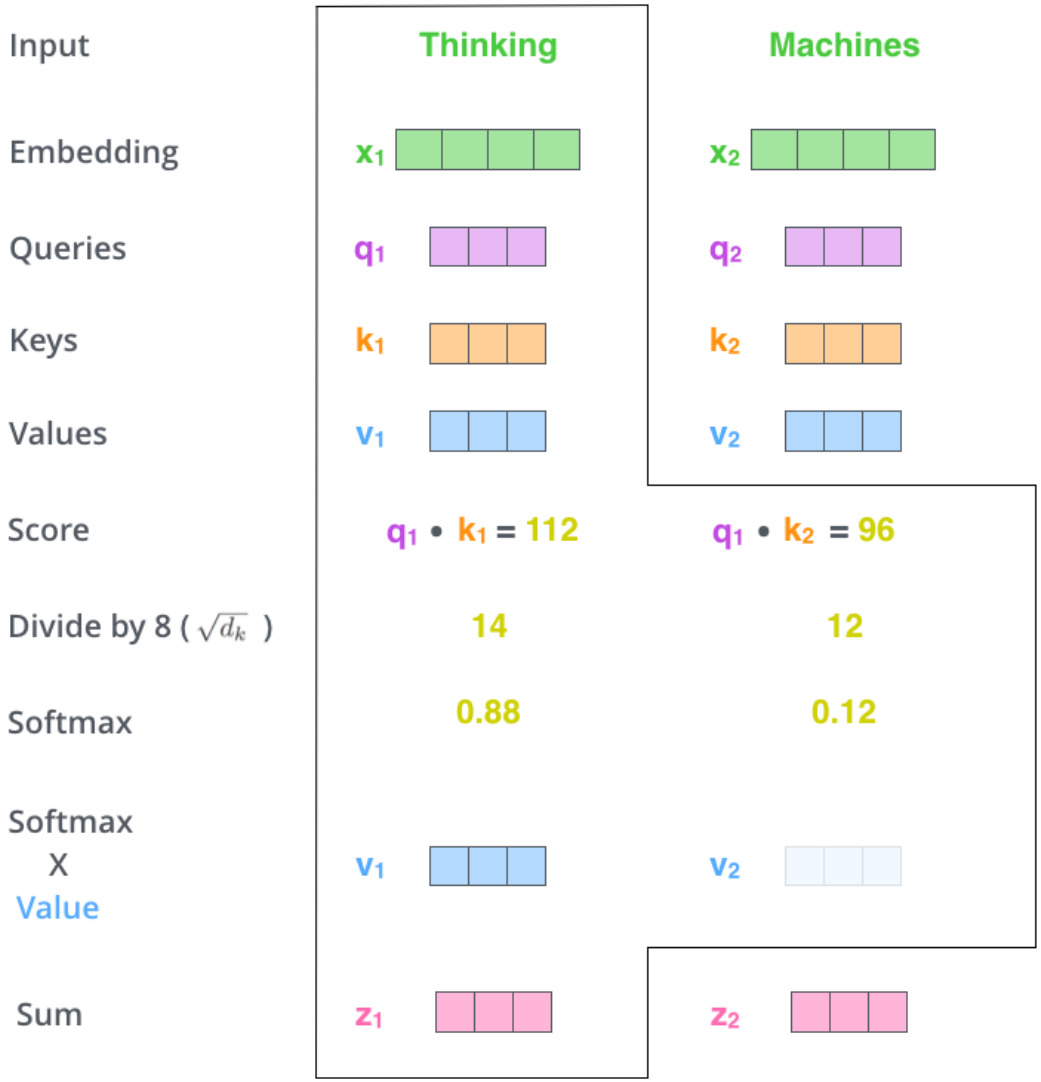

Figure: Attention (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Multi-Head Attention

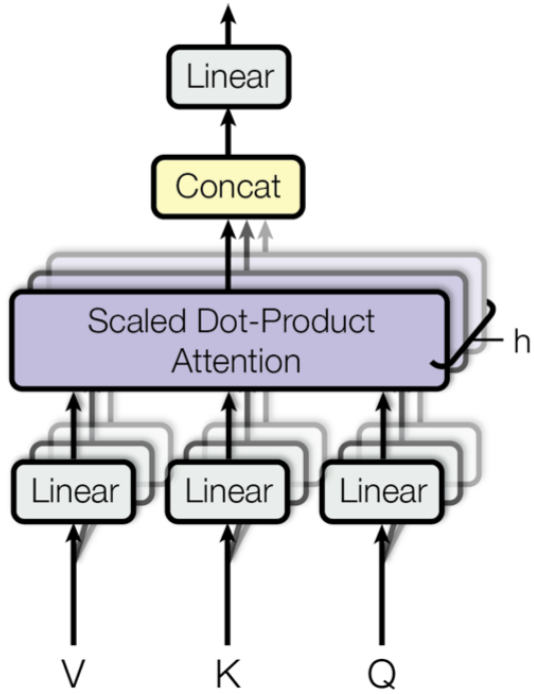

Figure: Scaled Dot-Product Attention (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Attentions Heads

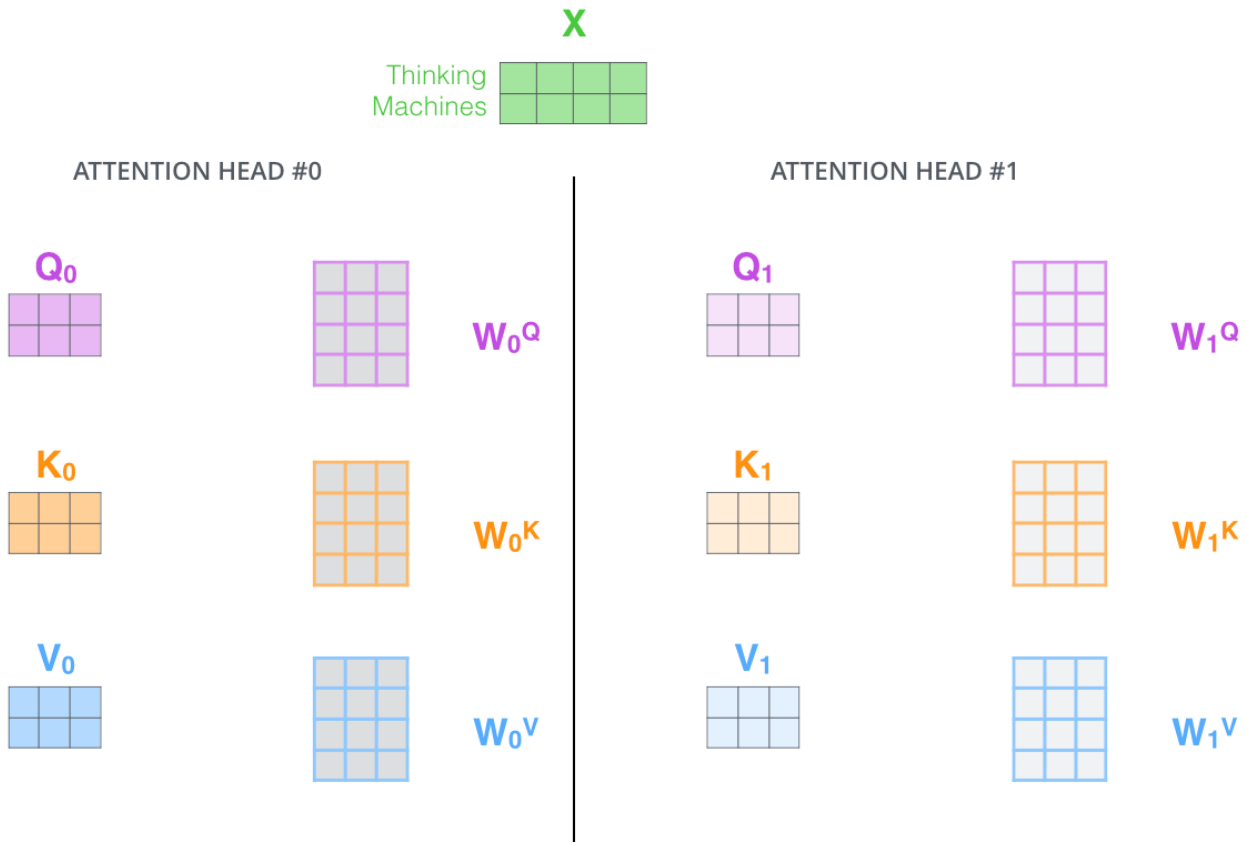

Figure: Attention heads (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Multi-head attention

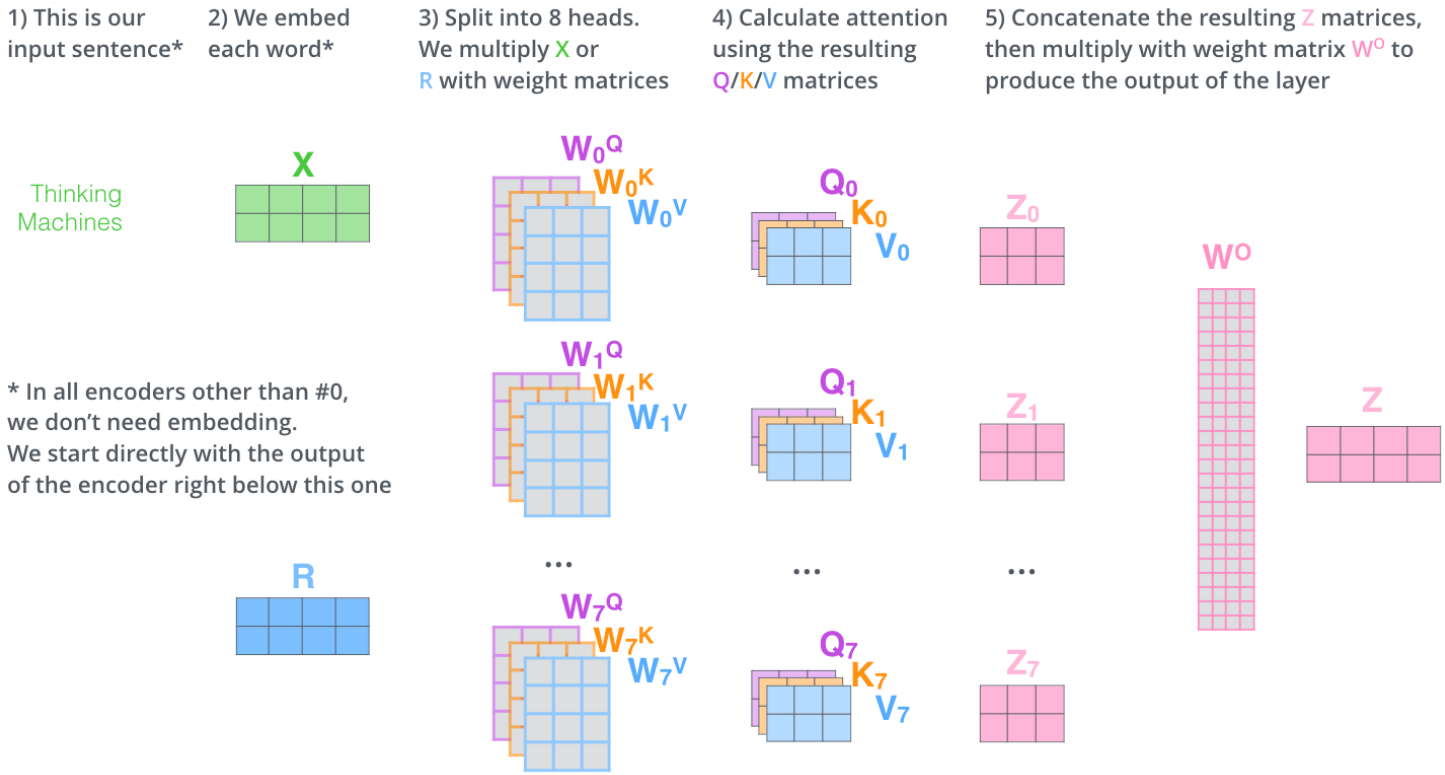

Figure: Attention heads (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Multi-Head Attention example

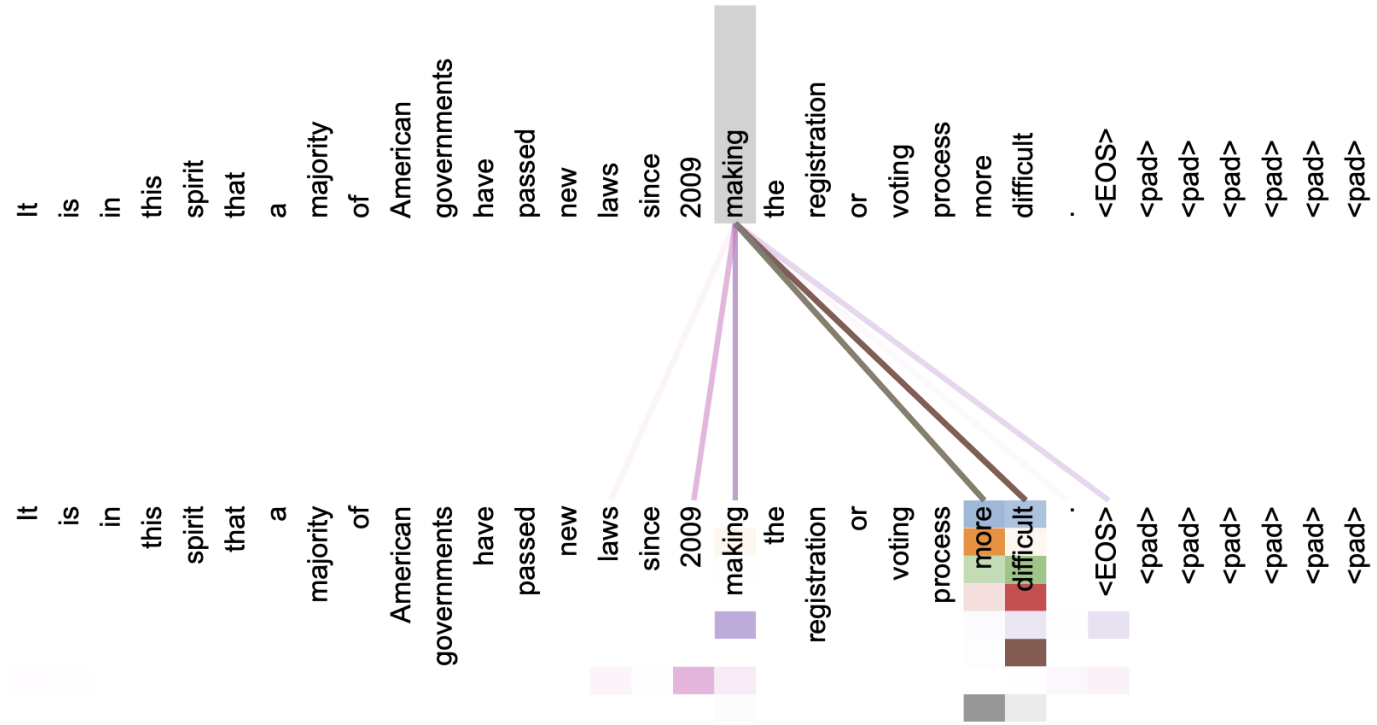

Figure 3: An example of the attention mechanism following long-distance dependencies in the encoder self-attention in layer 5 of 6. Many of the attention heads attend to a distant dependency of the verb 'making', completing the phrase 'making...more difficult'. Attentions here shown only for the word 'making'. Different colors represent different heads. Best viewed in color.

Figure: Attention (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Positional Encoding

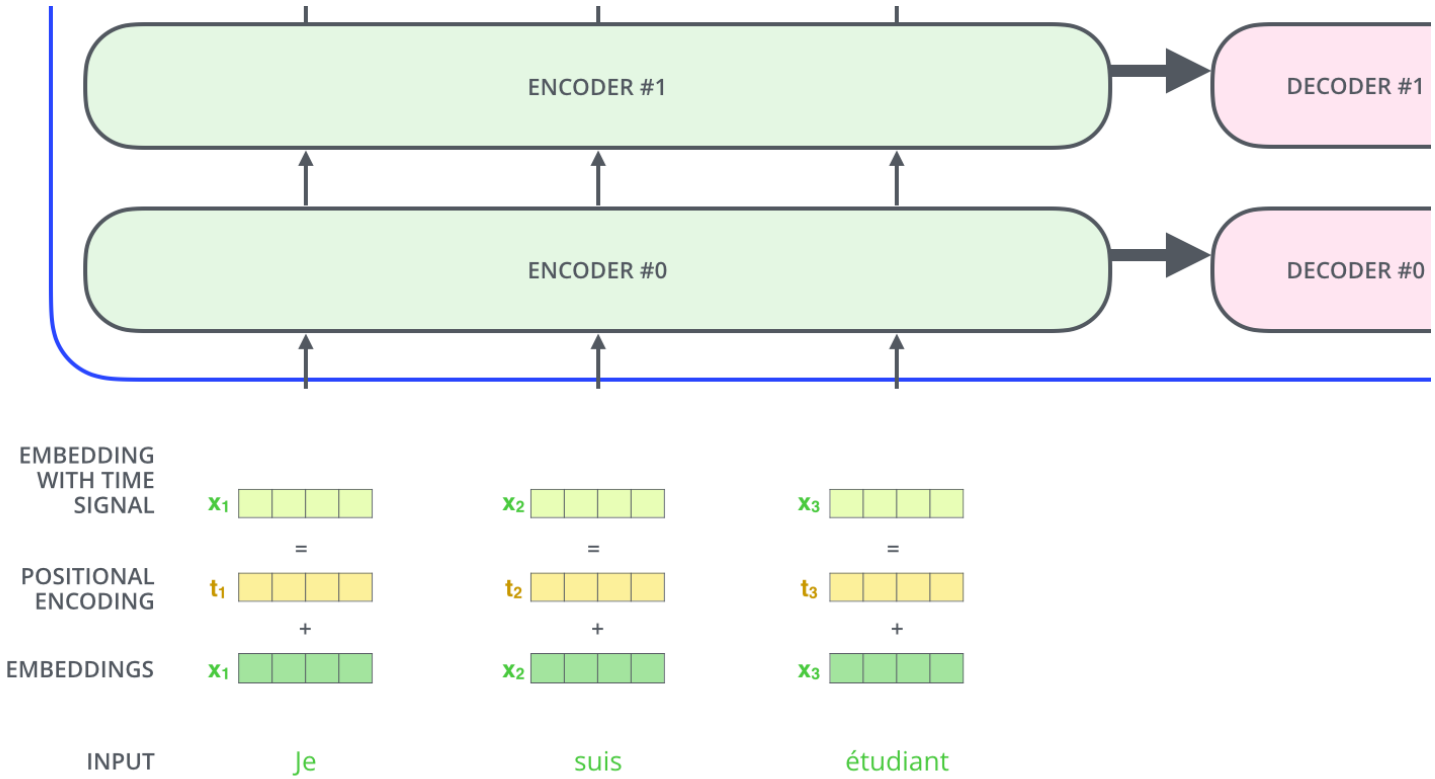

Figure: Attention heads (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Positional Encoding

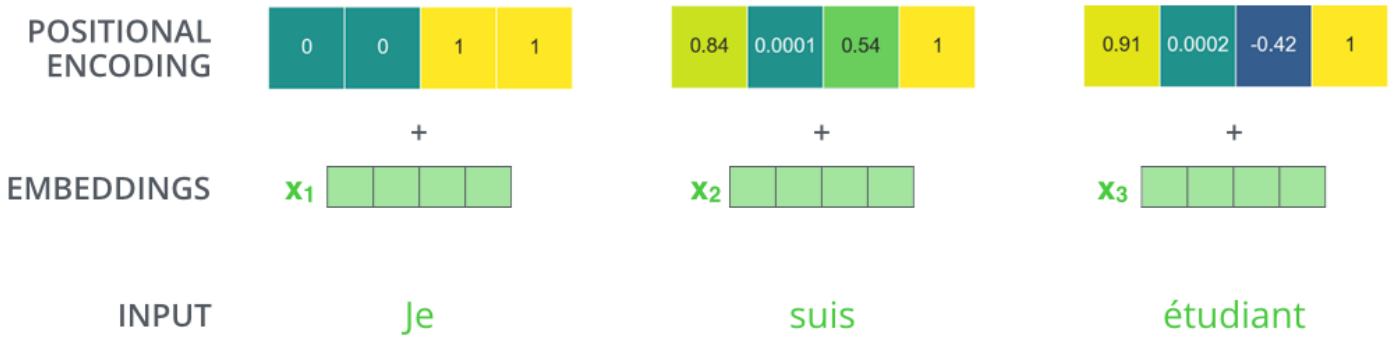

Figure: Adding positional encodings to embeddings (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Add and Normalize

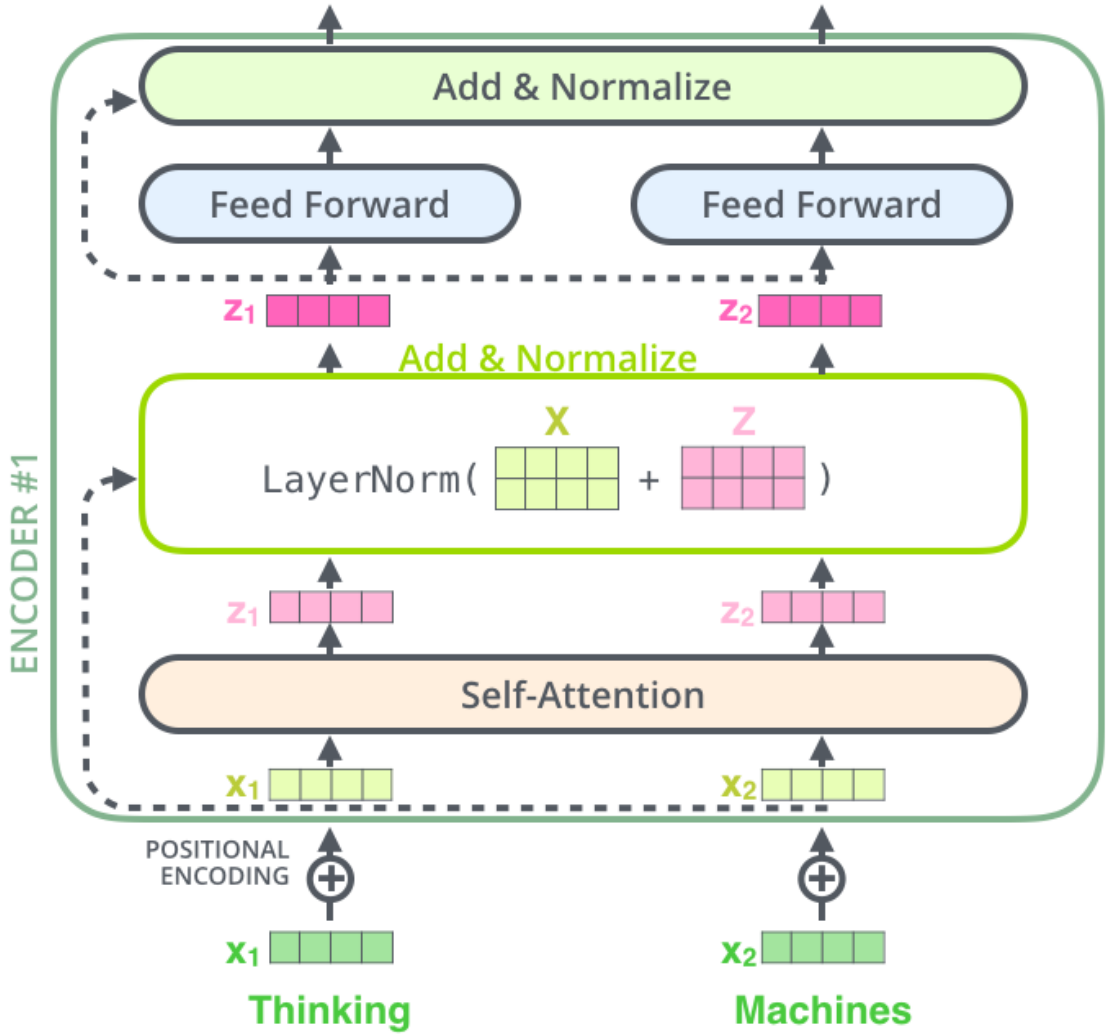

Figure: Add and Normalize (Alammar, 2018)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Transformer

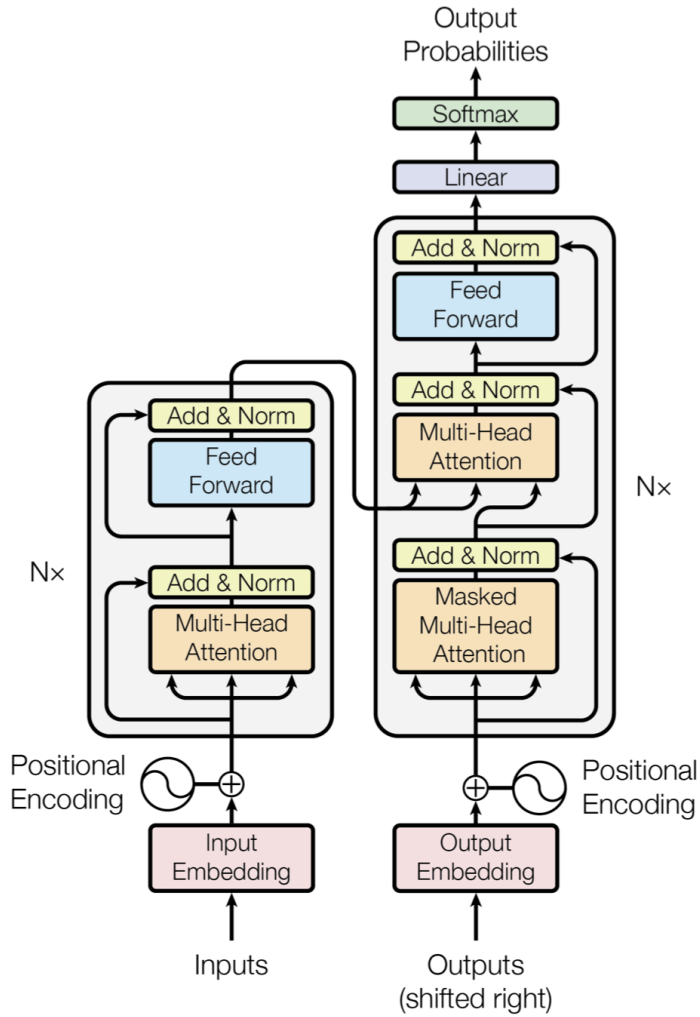

Figure: The Transformer Architecture (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT

- Bidirectional Encoder Representations from Transformers
- Introduced in 2018/2019 in Devlin et al. (2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT

- Bidirectional Encoder Representations from Transformers
- Introduced in 2018/2019 in Devlin et al. (2017)
- State-of-the-Art in many text prediction tasks, such as
  - Question-Answering
  - Named-Entity Recognition
  - Text Classification

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT

- Bidirectional Encoder Representations from Transformers
- Introduced in 2018/2019 in Devlin et al. (2017)
- State-of-the-Art in many text prediction tasks, such as
  - Question-Answering
  - Named-Entity Recognition
  - Text Classification
- Becoming the de facto new standard in industry and NLP

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT

- Bidirectional Encoder Representations from Transformers
- Introduced in 2018/2019 in Devlin et al. (2017)
- State-of-the-Art in many text prediction tasks, such as
  - Question-Answering
  - Named-Entity Recognition
  - Text Classification
- Becoming the de facto new standard in industry and NLP
- Pre-trained on a large number of books

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT

- Bidirectional Encoder Representations from Transformers
- Introduced in 2018/2019 in Devlin et al. (2017)
- State-of-the-Art in many text prediction tasks, such as
  - Question-Answering
  - Named-Entity Recognition
  - Text Classification
- Becoming the de facto new standard in industry and NLP
- Pre-trained on a large number of books
- Available both in English and Swedish (The National Library)

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT and transfer learning

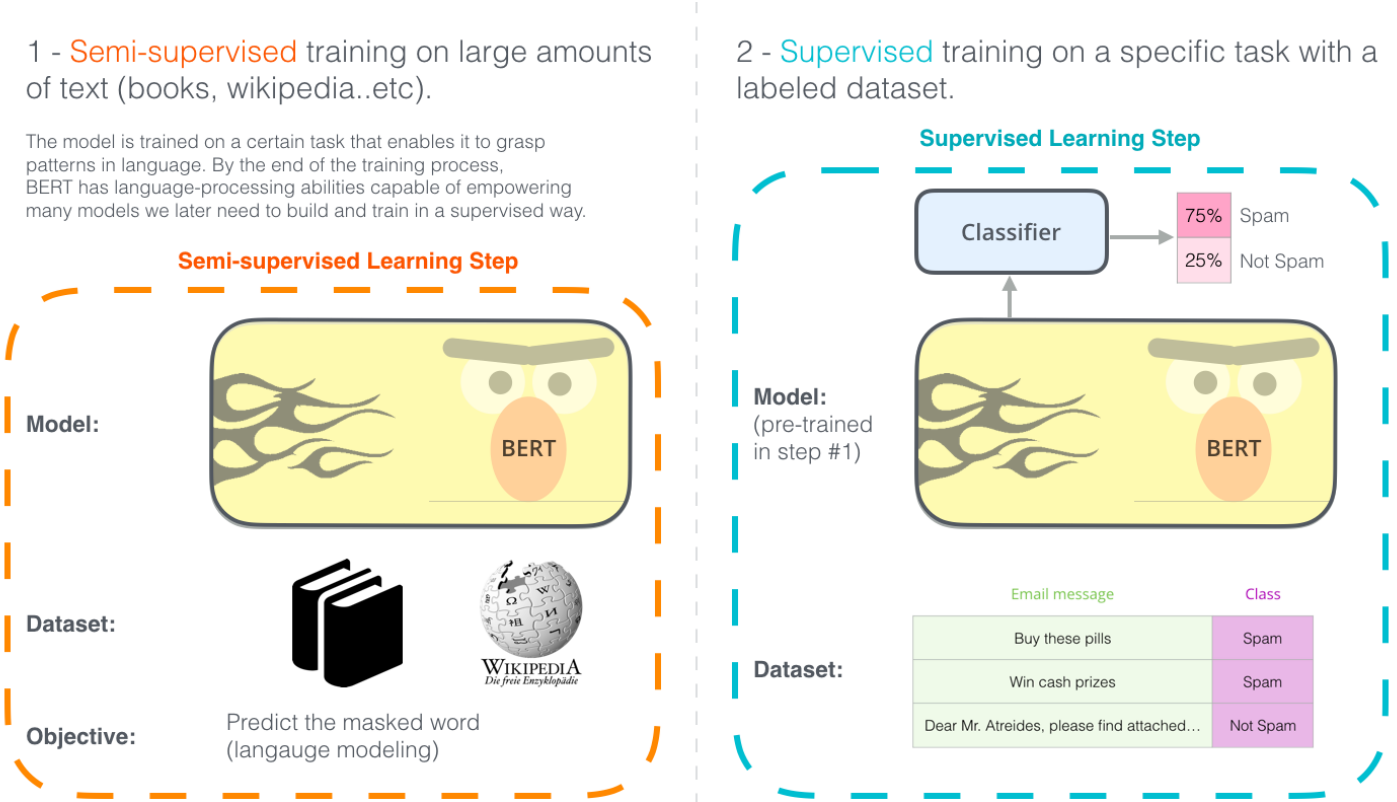

Figure: Using BERT for Transfer Learning (Alammar, 2019)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

The BERT model

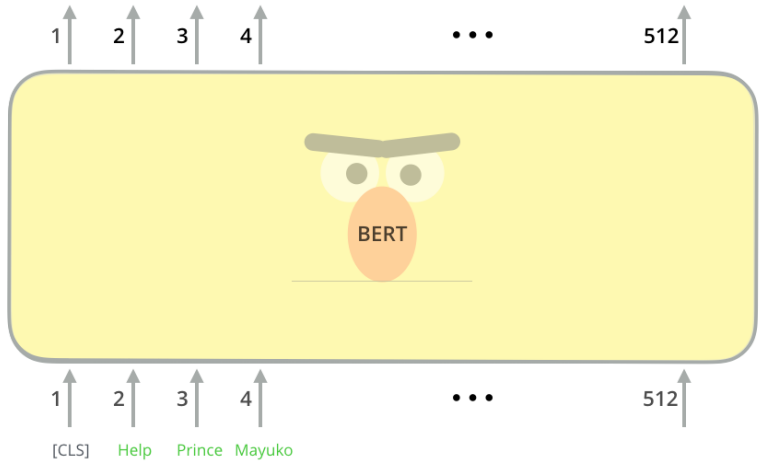

Figure: The BERT model (Alammar, 2019)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT Architecture

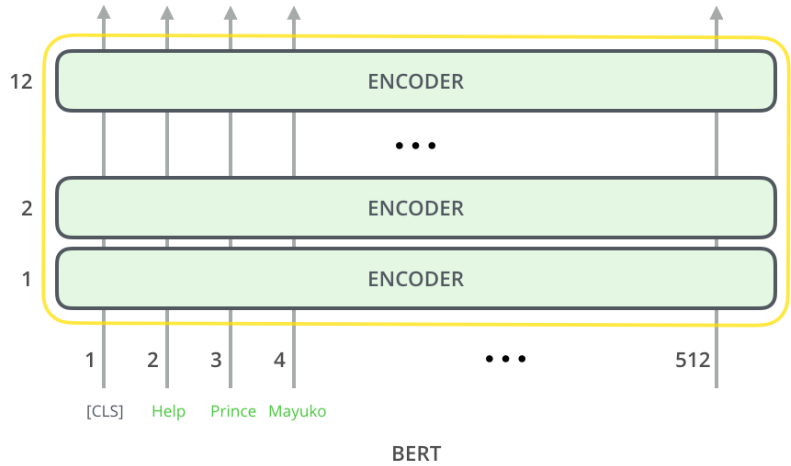

Figure: Opening up BERT (Vaswani et al., 2017)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Task 1: Masked Language Model

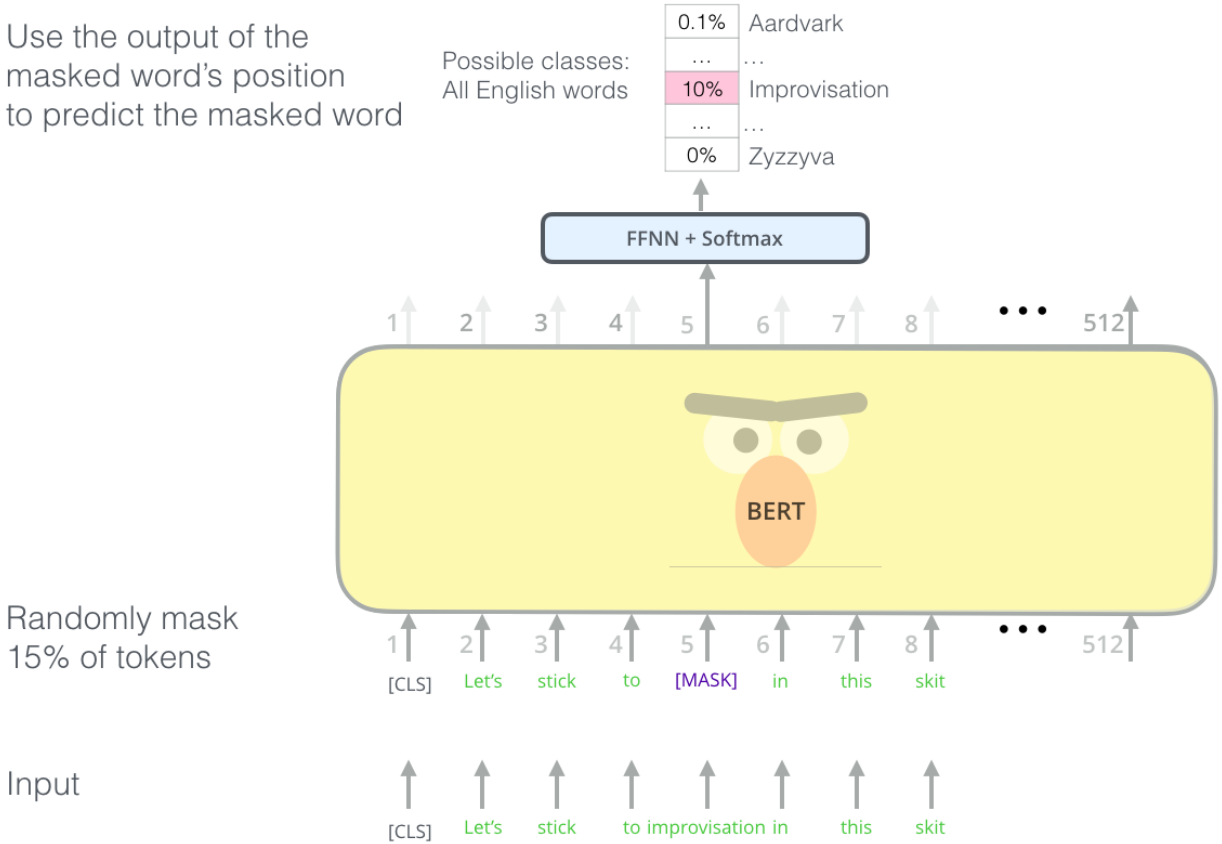

Figure: Masked Language Modeling (Alammar, 2019)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Next Sentence Prediction

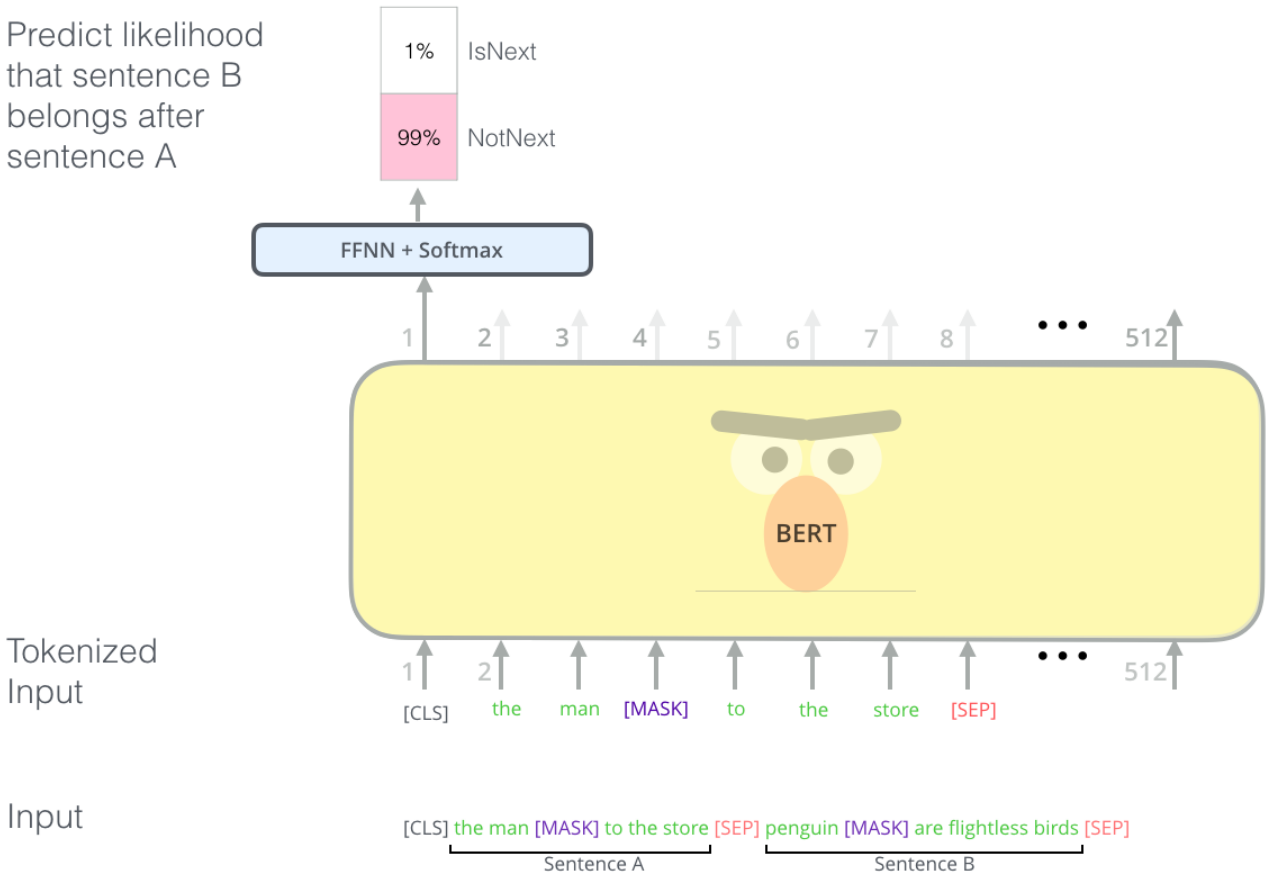

Figure: Next Sentence Prediction (Alammar, 2019)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Using BERT for Classification

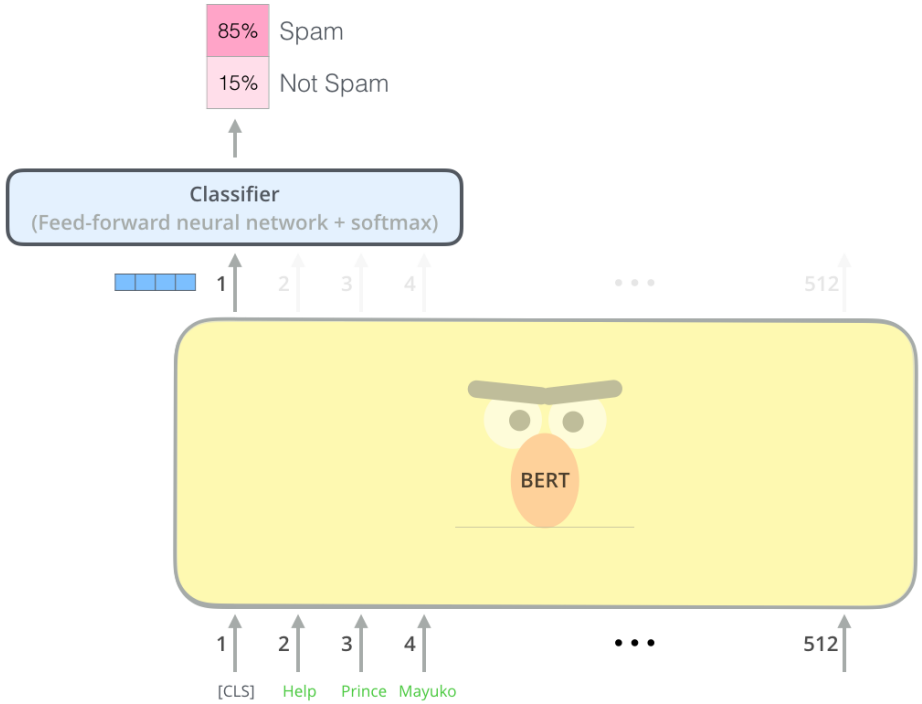

Figure: Using BERT for classification (Alammar, 2019)

UPPSALA UNIVERSITET

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

BERT and Contextualized embeddings

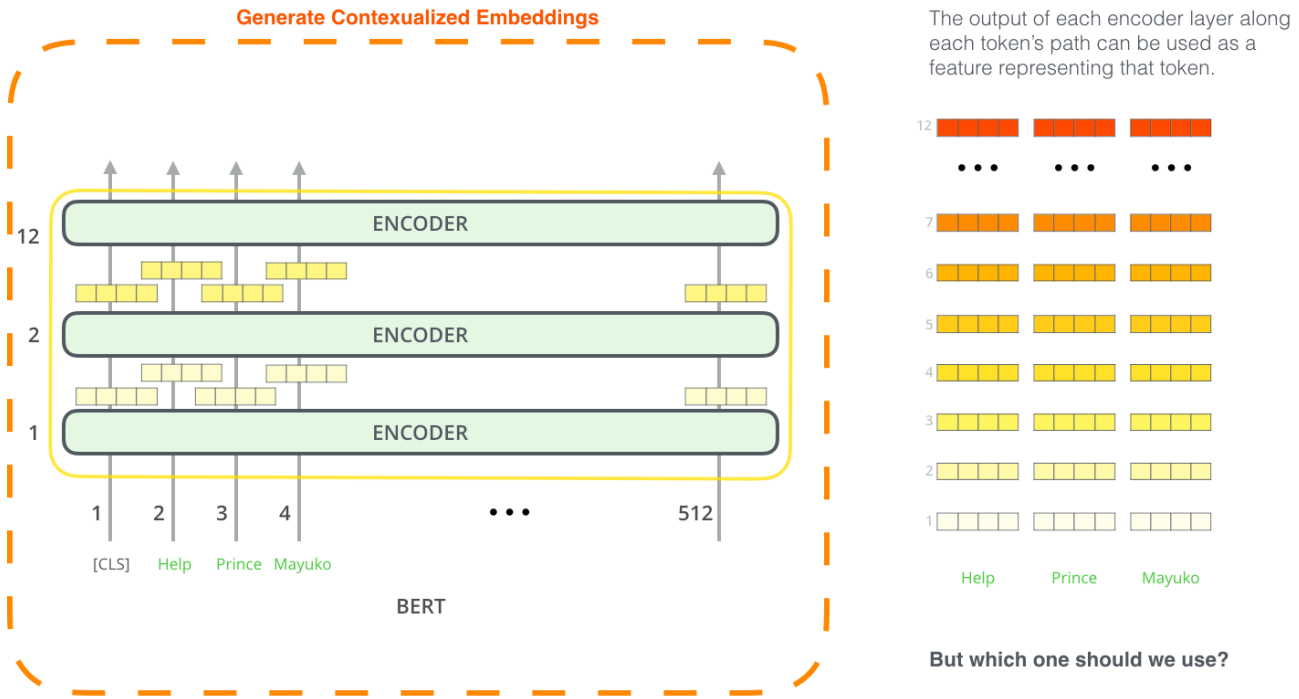

Figure: Contextualized Embeddings (Alammar, 2019)

- Word embeddings
- Recurrent Neural Networks
  - LSTM
- Attention and Transformers
  - Attention
  - Multi-Head Attention
  - Positional encoding
  - Add and Normalize
- BERT
  - Training BERT
  - Using BERT

Using Contextualized Embeddings

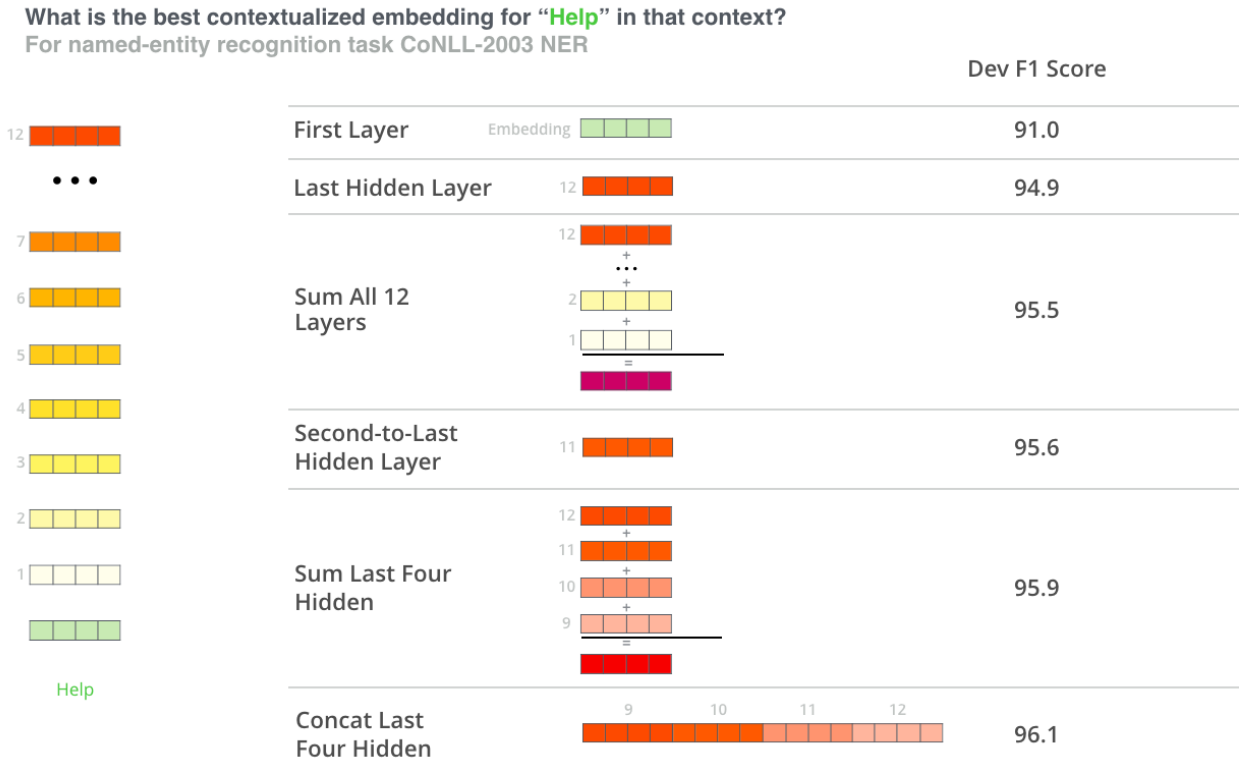

Figure: Using Contextualized Embeddings (Alammar, 2019)