



UPPSALA  
UNIVERSITET

# Machine learning – Block 7c

- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

Måns Magnusson  
Department of Statistics, Uppsala University

Autumn 2025



# Diffusion Models

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models





UPPSALA  
UNIVERSITET

# Diffusion Models

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models
- State-of-the-art for image generation



# Diffusion Models

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models
- State-of-the-art for image generation
- Reading: Bishop & Bishop, Chapter 20.



# Diffusion Models

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models
- State-of-the-art for image generation
- Reading: Bishop & Bishop, Chapter 20.
- Closely related to variational autoencoders



# Generative Models Recap

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- (Deep) latent variable models:
  - Variational Autoencoders (VAEs)
  - Diffusion Models
- All transform a simple (latent) distribution into complex data



# Generative Models Recap

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- (Deep) latent variable models:
  - Variational Autoencoders (VAEs)
  - Diffusion Models
- All transform a simple (latent) distribution into complex data
- Main difference: how this transformation is learned



# Why Diffusion Models?

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Stable training
- Excellent sample quality
- Main drawback: slow sampling

Demo: DALL·E 3: <https://openai.com/dall-e-3>



- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

## High-Level Idea

---

- Two processes:
  1. **Forward process:** gradually add noise to data
  2. **Reverse process:** learn to remove noise

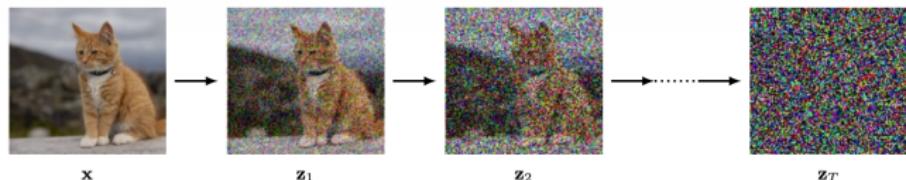


Figure: Figure 20.1 from Bishop & Bishop (2024).



- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

## High-Level Idea

---

- Two processes:
  1. **Forward process:** gradually add noise to data
  2. **Reverse process:** learn to remove noise

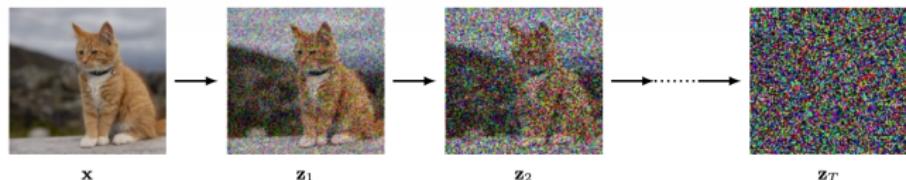


Figure: Figure 20.1 from Bishop & Bishop (2024).

- Generation starts from pure noise
- Data is generated by iteratively denoising



# Diffusion Models and VAEs

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models can be viewed as hierarchical VAEs
- Key differences:
  - Encoder is **fixed** (noise process)



# Diffusion Models and VAEs

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models can be viewed as hierarchical VAEs
- Key differences:
  - Encoder is **fixed** (noise process)
  - Decoder is learned (denoising network)



# Diffusion Models and VAEs

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models can be viewed as hierarchical VAEs
- Key differences:
  - Encoder is **fixed** (noise process)
  - Decoder is learned (denoising network)
  - Many latent variables instead of a few



- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

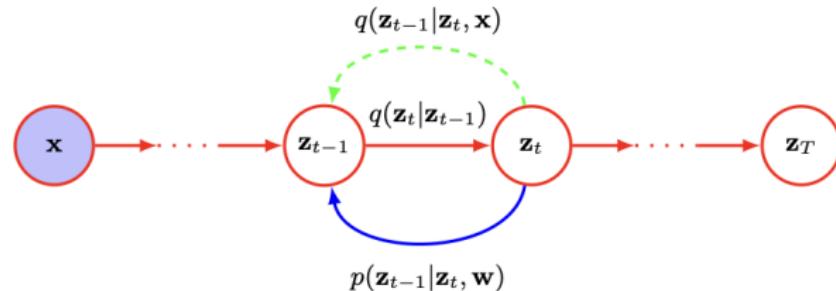


Figure: Figure 20.2 from Bishop & Bishop (2024).

- Start with data point  $x$
- Add small Gaussian noise repeatedly
- After many steps: result is Gaussian noise



## Forward Process: Single Step

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

For  $t = 1, \dots, T$ :

$$z_t = \sqrt{1 - \beta_t} z_{t-1} + \sqrt{\beta_t} \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, I)$$

- $\beta_t$  controls noise level
- Noise increases with  $t$
- Defines a Markov chain



# Probabilistic Encoder

---

- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

- The forward process defines:

$$q(z_t \mid z_{t-1}) = \mathcal{N} \left( \sqrt{1 - \beta_t} z_{t-1}, \beta_t I \right)$$

- This plays the role of an **encoder**
- Unlike VAEs: encoder is fixed



# Closed-Form Noising

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

Key result:

$$z_t = \sqrt{\alpha_t}x + \sqrt{1 - \alpha_t}\varepsilon, \quad \varepsilon \sim \mathcal{N}(0, I)$$

- $\alpha_t = \prod_{\tau=1}^t (1 - \beta_\tau)$
- Allows direct sampling of  $z_t$  from  $x$
- Very important for training



# End of Forward Process

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- As  $t \rightarrow T$ :

$$z_T \sim \mathcal{N}(0, I)$$

- All information about  $x$  is lost
- This distribution is easy to sample from



- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

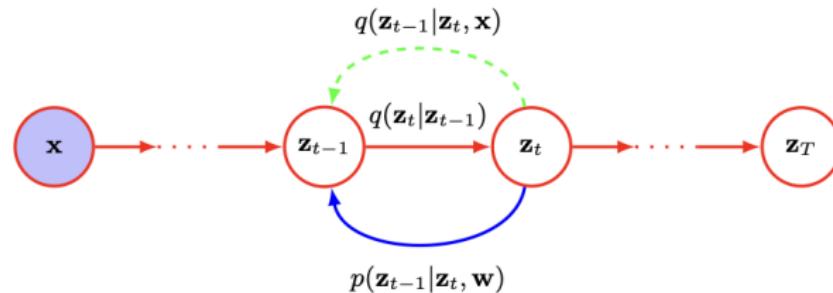


Figure: Figure 20.2 from Bishop & Bishop (2024).

- Goal: reverse the noise process

$$p(z_{t-1} | z_t)$$

- Exact reverse distribution is intractable
- We learn an approximation using a neural network



# Probabilistic Decoder

---

- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

We model:

$$p(z_{t-1} \mid z_t, w) = \mathcal{N}(\mu(z_t, t), \beta_t I)$$

- Mean predicted by neural network
- Variance often fixed
- Decoder shared across all steps



# Sampling from the Model

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

1. Sample  $z_T \sim \mathcal{N}(0, I)$
2. For  $t = T, \dots, 1$ :
  - Sample  $z_{t-1} \sim p(z_{t-1} | z_t)$
3. Final sample corresponds to data point  $x$



- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

# Sampling from the Model

## Algorithm 20.2: Sampling from a denoising diffusion probabilistic model

**Input:** Trained denoising network  $\mathbf{g}(\mathbf{z}, \mathbf{w}, t)$

Noise schedule  $\{\beta_1, \dots, \beta_T\}$

**Output:** Sample vector  $\mathbf{x}$  in data space

$\mathbf{z}_T \sim \mathcal{N}(\mathbf{z}|\mathbf{0}, \mathbf{I})$  // Sample from final latent space

for  $t \in T, \dots, 2$  do

$\alpha_t \leftarrow \prod_{\tau=1}^t (1 - \beta_\tau)$  // Calculate alpha

// Evaluate network output

$\mu(\mathbf{z}_t, \mathbf{w}, t) \leftarrow \frac{1}{\sqrt{1-\beta_t}} \left\{ \mathbf{z}_t - \frac{\beta_t}{\sqrt{1-\alpha_t}} \mathbf{g}(\mathbf{z}_t, \mathbf{w}, t) \right\}$

$\epsilon \sim \mathcal{N}(\epsilon|\mathbf{0}, \mathbf{I})$  // Sample a noise vector

$\mathbf{z}_{t-1} \leftarrow \mu(\mathbf{z}_t, \mathbf{w}, t) + \sqrt{\beta_t} \epsilon$  // Add scaled noise

end for

$\mathbf{x} = \frac{1}{\sqrt{1-\beta_1}} \left\{ \mathbf{z}_1 - \frac{\beta_1}{\sqrt{1-\alpha_1}} \mathbf{g}(\mathbf{z}_1, \mathbf{w}, t) \right\}$  // Final denoising step

return  $\mathbf{x}$

Figure: Algorithm 20.2 from Bishop & Bishop (2024).



# Training Objective

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Exact likelihood is intractable
- Maximize an Evidence Lower Bound (ELBO)
- Very similar to VAE training



# ELBO Interpretation

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Reconstruction term
- Consistency terms between forward and reverse processes
- Encoder distribution is fixed
- Only decoder parameters are learned



# Predicting Noise Instead of Data

---

- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

Key idea:

- Instead of predicting  $z_{t-1}$
- Predict the noise  $\varepsilon$  added to  $x$

$$z_t = \sqrt{\alpha_t}x + \sqrt{1 - \alpha_t}\varepsilon$$



# Final Training Loss

---

- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

The loss simplifies to:

$$\mathbb{E}_{x,t,\varepsilon} [\|\varepsilon - g(z_t, t)\|^2]$$

- Simple squared error loss
- Very stable optimization
- Central result of diffusion models



- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

# Training a Denoising diffusion model

## Algorithm 20.1: Training a denoising diffusion probabilistic model

**Input:** Training data  $\mathcal{D} = \{\mathbf{x}_n\}$

Noise schedule  $\{\beta_1, \dots, \beta_T\}$

**Output:** Network parameters  $\mathbf{w}$

**for**  $t \in \{1, \dots, T\}$  **do**

$\alpha_t \leftarrow \prod_{\tau=1}^t (1 - \beta_\tau)$  // Calculate alphas from betas

**end for**

**repeat**

$\mathbf{x} \sim \mathcal{D}$  // Sample a data point

$t \sim \{1, \dots, T\}$  // Sample a point along the Markov chain

$\boldsymbol{\epsilon} \sim \mathcal{N}(\boldsymbol{\epsilon} | \mathbf{0}, \mathbf{I})$  // Sample a noise vector

$\mathbf{z}_t \leftarrow \sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}$  // Evaluate noisy latent variable

$\mathcal{L}(\mathbf{w}) \leftarrow \|\mathbf{g}(\mathbf{z}_t, \mathbf{w}, t) - \boldsymbol{\epsilon}\|^2$  // Compute loss term

Take optimizer step

**until** converged

**return**  $\mathbf{w}$

Figure: Algorithm 20.1 from Bishop & Bishop (2024).



# Diffusion Models vs VAEs

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

	VAE	Diffusion
Encoder	Learned	Fixed
Latent dim.	Low	High
Decoder	One step	Many steps
Training	ELBO	ELBO
Sampling	Fast	Slow



# Conditional Diffusion Models

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- So far: unconditional generation
  - Sample images from  $p(x)$
- Often we want **conditional generation**
  - Generate images given text
  - Generate images given class labels
- Goal: model  $p(x | c)$ , where  $c$  is conditioning information



# Conditioning on Text

---

- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

- Let  $c$  be a text description
- Text is encoded using a language model

$$c \rightarrow \text{embedding } e(c)$$

- Diffusion model is conditioned on this embedding
- Reverse process becomes:

$$p(z_{t-1} \mid z_t, c)$$



- Diffusion Models
  - Core Idea
  - Forward Process (Encoder)
  - Reverse Process (Decoder)
  - Training
  - Summary
  - Conditional Diffusion Models

## How Conditioning Is Used

- Conditioning information is provided to the neural network
- Noise prediction network becomes:
$$g(z_t, t, c)$$
- Intuition:
  - Noise removal is guided by the text
  - Different texts lead to different denoising trajectories
- Used in systems such as:
  - Text-to-image generation
  - Image editing and inpainting





# How Conditioning Is Used

- Used in systems such as:
  - Text-to-image generation
  - Image editing and inpainting

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

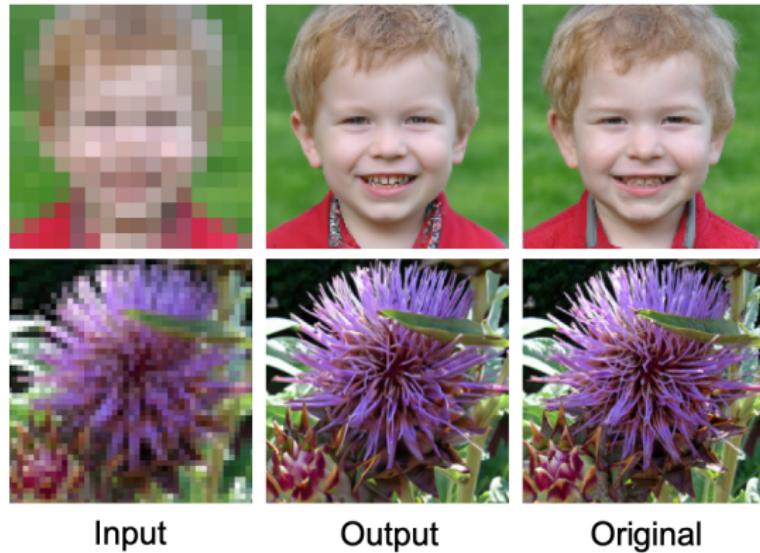


Figure: Figure 20.8 from Bishop & Bishop (2024).



# Key Takeaways

---

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models are deep probabilistic models
- Closely related to VAEs
- Built around denoising Gaussian noise
- Foundation for modern image generation systems conditional on textual input