



UPPSALA
UNIVERSITET

Machine learning – Block 7c

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

Måns Magnusson
Department of Statistics, Uppsala University

Autumn 2025



Diffusion Models

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models





UPPSALA
UNIVERSITET

Diffusion Models

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models
- State-of-the-art for image generation



Diffusion Models

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models
- State-of-the-art for image generation
- Reading: Bishop & Bishop, Chapter 20.



Diffusion Models

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- A new class of deep probabilistic generative models
- State-of-the-art for image generation
- Reading: Bishop & Bishop, Chapter 20.
- Closely related to variational autoencoders



Generative Models Recap

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- (Deep) latent variable models:
 - Variational Autoencoders (VAEs)
 - Diffusion Models
- All transform a simple (latent) distribution into complex data



Generative Models Recap

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- (Deep) latent variable models:
 - Variational Autoencoders (VAEs)
 - Diffusion Models
- All transform a simple (latent) distribution into complex data
- Main difference: how this transformation is learned



Why Diffusion Models?

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Stable training
- Excellent sample quality
- Main drawback: slow sampling

Demo: DALL·E 3: <https://openai.com/dall-e-3>



- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

High-Level Idea

- Two processes:
 1. **Forward process:** gradually add noise to data
 2. **Reverse process:** learn to remove noise

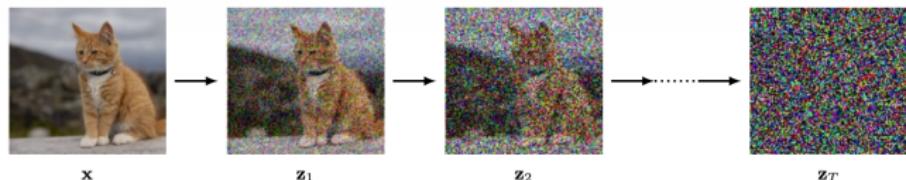


Figure: Figure 20.1 from Bishop & Bishop (2024).



- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

High-Level Idea

- Two processes:
 1. **Forward process:** gradually add noise to data
 2. **Reverse process:** learn to remove noise

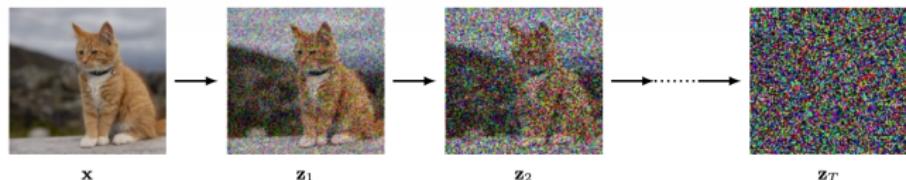


Figure: Figure 20.1 from Bishop & Bishop (2024).

- Generation starts from pure noise
- Data is generated by iteratively denoising



Diffusion Models and VAEs

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models can be viewed as hierarchical VAEs
- Key differences:
 - Encoder is **fixed** (noise process)



Diffusion Models and VAEs

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models can be viewed as hierarchical VAEs
- Key differences:
 - Encoder is **fixed** (noise process)
 - Decoder is learned (denoising network)



Diffusion Models and VAEs

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models can be viewed as hierarchical VAEs
- Key differences:
 - Encoder is **fixed** (noise process)
 - Decoder is learned (denoising network)
 - Many latent variables instead of a few



- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

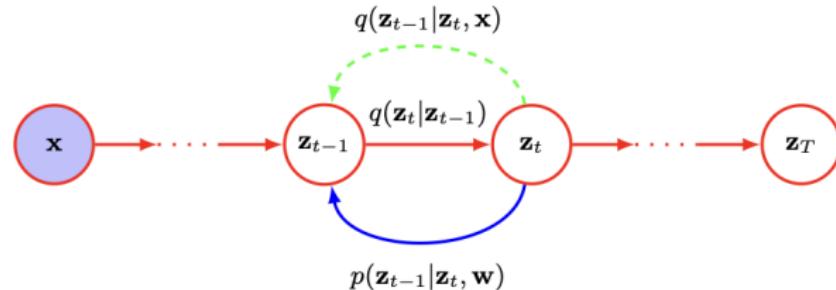


Figure: Figure 20.2 from Bishop & Bishop (2024).

- Start with data point x
- Add small Gaussian noise repeatedly
- After many steps: result is Gaussian noise



Forward Process: Single Step

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

For $t = 1, \dots, T$:

$$z_t = \sqrt{1 - \beta_t} z_{t-1} + \sqrt{\beta_t} \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, I)$$

- β_t controls noise level
- Noise increases with t
- Defines a Markov chain



Probabilistic Encoder

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

- The forward process defines:

$$q(z_t \mid z_{t-1}) = \mathcal{N} \left(\sqrt{1 - \beta_t} z_{t-1}, \beta_t I \right)$$

- This plays the role of an **encoder**
- Unlike VAEs: encoder is fixed



Closed-Form Noising

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

Key result:

$$z_t = \sqrt{\alpha_t}x + \sqrt{1 - \alpha_t}\varepsilon, \quad \varepsilon \sim \mathcal{N}(0, I)$$

- $\alpha_t = \prod_{\tau=1}^t (1 - \beta_\tau)$
- Allows direct sampling of z_t from x
- Very important for training



End of Forward Process

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

- As $t \rightarrow T$:

$$z_T \sim \mathcal{N}(0, I)$$

- All information about x is lost
- This distribution is easy to sample from



- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

Learning the Reverse Process

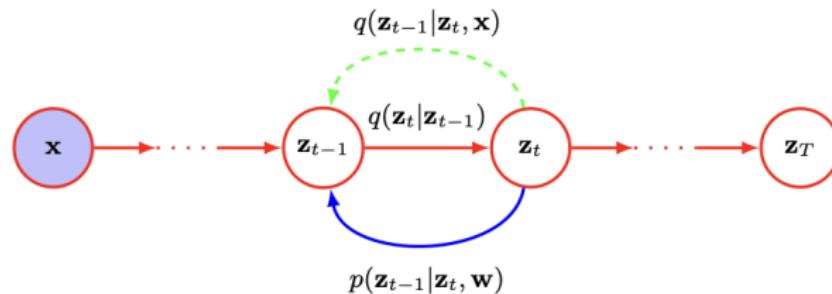


Figure: Figure 20.2 from Bishop & Bishop (2024).

- Goal: reverse the noise process

$$p(z_{t-1} | z_t)$$

- Exact reverse distribution is intractable
- We learn an approximation using a neural network



Probabilistic Decoder

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

We model:

$$p(z_{t-1} \mid z_t, w) = \mathcal{N}(\mu(z_t, t), \beta_t I)$$

- Mean predicted by neural network
- Variance often fixed
- Decoder shared across all steps



Sampling from the Model

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

1. Sample $z_T \sim \mathcal{N}(0, I)$
2. For $t = T, \dots, 1$:
 - Sample $z_{t-1} \sim p(z_{t-1} | z_t)$
3. Final sample corresponds to data point x



- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

Sampling from the Model

Algorithm 20.2: Sampling from a denoising diffusion probabilistic model

Input: Trained denoising network $\mathbf{g}(\mathbf{z}, \mathbf{w}, t)$

Noise schedule $\{\beta_1, \dots, \beta_T\}$

Output: Sample vector \mathbf{x} in data space

$\mathbf{z}_T \sim \mathcal{N}(\mathbf{z}|\mathbf{0}, \mathbf{I})$ // Sample from final latent space

for $t \in T, \dots, 2$ do

$\alpha_t \leftarrow \prod_{\tau=1}^t (1 - \beta_\tau)$ // Calculate alpha

// Evaluate network output

$\mu(\mathbf{z}_t, \mathbf{w}, t) \leftarrow \frac{1}{\sqrt{1-\beta_t}} \left\{ \mathbf{z}_t - \frac{\beta_t}{\sqrt{1-\alpha_t}} \mathbf{g}(\mathbf{z}_t, \mathbf{w}, t) \right\}$

$\epsilon \sim \mathcal{N}(\epsilon|\mathbf{0}, \mathbf{I})$ // Sample a noise vector

$\mathbf{z}_{t-1} \leftarrow \mu(\mathbf{z}_t, \mathbf{w}, t) + \sqrt{\beta_t} \epsilon$ // Add scaled noise

end for

$\mathbf{x} = \frac{1}{\sqrt{1-\beta_1}} \left\{ \mathbf{z}_1 - \frac{\beta_1}{\sqrt{1-\alpha_1}} \mathbf{g}(\mathbf{z}_1, \mathbf{w}, t) \right\}$ // Final denoising step

return \mathbf{x}

Figure: Algorithm 20.2 from Bishop & Bishop (2024).



Training Objective

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Exact likelihood is intractable
- Maximize an Evidence Lower Bound (ELBO)
- Very similar to VAE training



ELBO Interpretation

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Reconstruction term
- Consistency terms between forward and reverse processes
- Encoder distribution is fixed
- Only decoder parameters are learned



Predicting Noise Instead of Data

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

Key idea:

- Instead of predicting z_{t-1}
- Predict the noise ε added to x : $g(z_t, t)$



Final Training Loss

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

The loss simplifies to:

$$\mathbb{E}_{x,t,\varepsilon} [\|\varepsilon - g(z_t, t)\|^2]$$

- Simple squared error loss
- Very stable optimization
- Central result of diffusion models



- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

Training a Denoising diffusion model

Algorithm 20.1: Training a denoising diffusion probabilistic model

Input: Training data $\mathcal{D} = \{\mathbf{x}_n\}$

Noise schedule $\{\beta_1, \dots, \beta_T\}$

Output: Network parameters \mathbf{w}

for $t \in \{1, \dots, T\}$ **do**

$\alpha_t \leftarrow \prod_{\tau=1}^t (1 - \beta_\tau)$ // Calculate alphas from betas

end for

repeat

$\mathbf{x} \sim \mathcal{D}$ // Sample a data point

$t \sim \{1, \dots, T\}$ // Sample a point along the Markov chain

$\boldsymbol{\epsilon} \sim \mathcal{N}(\boldsymbol{\epsilon} | \mathbf{0}, \mathbf{I})$ // Sample a noise vector

$\mathbf{z}_t \leftarrow \sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}$ // Evaluate noisy latent variable

$\mathcal{L}(\mathbf{w}) \leftarrow \|\mathbf{g}(\mathbf{z}_t, \mathbf{w}, t) - \boldsymbol{\epsilon}\|^2$ // Compute loss term

Take optimizer step

until converged

return \mathbf{w}

Figure: Algorithm 20.1 from Bishop & Bishop (2024).



Diffusion Models vs VAEs

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

	VAE	Diffusion
Encoder	Learned	Fixed
Latent dim.	Low	High
Decoder	One step	Many steps
Training	ELBO	ELBO
Sampling	Fast	Slow



Conditional Diffusion Models

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- So far: unconditional generation
 - Sample images from $p(x)$
- Often we want **conditional generation**
 - Generate images given text
 - Generate images given class labels
- Goal: model $p(x | c)$, where c is conditioning information



Conditioning on Text

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Let c be a text description
- Text is encoded using a language model

$$c \rightarrow \text{embedding } e(c)$$

- Diffusion model is conditioned on this embedding
- Reverse process becomes:

$$p(z_{t-1} \mid z_t, c)$$



How Conditioning Is Used

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

- Conditioning information is provided to the neural network
- Noise prediction network becomes:
$$g(z_t, t, c)$$
- Intuition:
 - Noise removal is guided by the text
 - Different texts lead to different denoising trajectories
- Used in systems such as:
 - Text-to-image generation
 - Image editing and inpainting



How Conditioning Is Used

- Used in systems such as:
 - Text-to-image generation
 - Image editing and inpainting

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

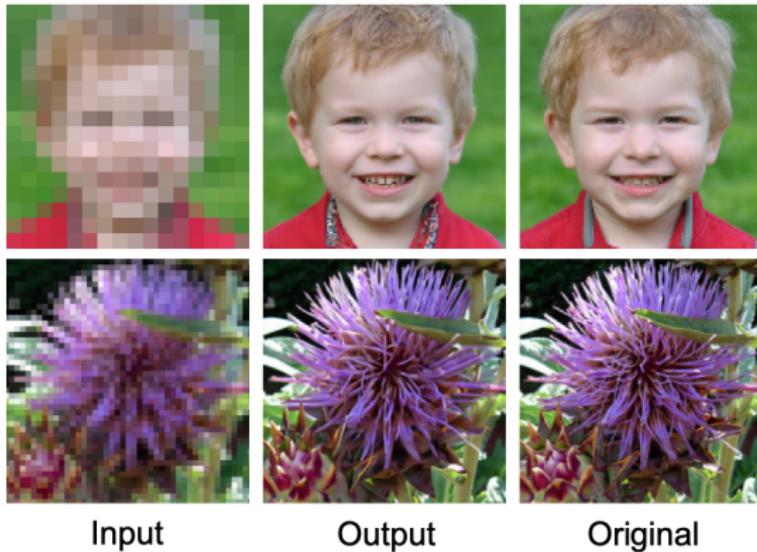


Figure: Figure 20.8 from Bishop & Bishop (2024).



Diffusion Ideas Beyond Images

- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

- Diffusion models popularized a key idea:
 - Learn complex objects by **iterative denoising**



- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

Diffusion Ideas Beyond Images

- Diffusion models popularized a key idea:
 - Learn complex objects by **iterative denoising**
- **AlphaFold 2** applies a related principle to protein structures:
 - Protein geometry is refined step-by-step
 - Each step reduces uncertainty and inconsistency
- Important distinction:
 - AlphaFold 2 is **not** a Denoising Diffusion Probabilistic Model
 - But it uses a diffusion-*like* denoising process



- Diffusion Models
 - Core Idea
 - Forward Process (Encoder)
 - Reverse Process (Decoder)
 - Training
 - Summary
 - Conditional Diffusion Models

Diffusion Ideas Beyond Images

- Diffusion models popularized a key idea:
 - Learn complex objects by **iterative denoising**
- **AlphaFold 2** applies a related principle to protein structures:
 - Protein geometry is refined step-by-step
 - Each step reduces uncertainty and inconsistency
- Important distinction:
 - AlphaFold 2 is **not** a Denoising Diffusion Probabilistic Model
 - But it uses a diffusion-*like* denoising process
- New generative biology models now use **explicit diffusion models**
 - Protein backbone generation (Watson et al, 2023)
 - Molecular design (Hoogeboom et al, 2022)



Key Takeaways

- Diffusion Models

- Core Idea
- Forward Process (Encoder)
- Reverse Process (Decoder)
- Training
- Summary
- Conditional Diffusion Models

- Diffusion models are deep probabilistic models
- Closely related to VAEs
- Built around denoising Gaussian noise
- Foundation for modern image generation systems conditional on textual input