



UPPSALA  
UNIVERSITET

# Machine learning – Block 8

Måns Magnusson  
Department of Statistics, Uppsala University

Autumn 2024

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes



UPPSALA  
UNIVERSITET

# This week's lectures

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Introduction to Reinforcement learning



UPPSALA  
UNIVERSITET

# Introduction to Reinforcement Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Another type of Machine Learning:
  - Supervised Learning
  - Unsupervised Learning
  - Reinforcement learning



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# Introduction to Reinforcement Learning

---

- Another type of Machine Learning:
  - Supervised Learning
  - Unsupervised Learning
  - Reinforcement learning
- Computational approach of learning from interaction



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# Introduction to Reinforcement Learning

---

- Another type of Machine Learning:
  - Supervised Learning
  - Unsupervised Learning
  - Reinforcement learning
- Computational approach of learning from interaction
- Closest to human and animal learning: trial, error, and planning.



UPPSALA  
UNIVERSITET

# Introduction to Reinforcement Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Another type of Machine Learning:
  - Supervised Learning
  - Unsupervised Learning
  - Reinforcement learning
- Computational approach of learning from interaction
- Closest to human and animal learning: trial, error, and planning.
- The learner is *not* told which actions to take



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# Introduction to Reinforcement Learning

---

- Another type of Machine Learning:
  - Supervised Learning
  - Unsupervised Learning
  - Reinforcement learning
- Computational approach of learning from interaction
- Closest to human and animal learning: trial, error, and planning.
- The learner is *not* told which actions to take
- Connections to:
  - Game Theory
  - Control Theory
  - Multi-agent systems
  - Swarm intelligence
  - Information theory
  - Statistics



UPPSALA  
UNIVERSITET

# Introduction to Reinforcement Learning

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- **Goal:** maximize return over a sequence of actions





UPPSALA  
UNIVERSITET

# Introduction to Reinforcement Learning

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- **Goal**: maximize return over a sequence of actions
  - Three characteristics:
    1. **Closed-loop**: early actions affects later actions



UPPSALA  
UNIVERSITET

# Introduction to Reinforcement Learning

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- **Goal**: maximize return over a sequence of actions
  - Three characteristics:
    1. **Closed-loop**: early actions affects later actions
    2. **No instructions**



UPPSALA  
UNIVERSITET

# Introduction to Reinforcement Learning

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- **Goal**: maximize return over a sequence of actions
  - Three characteristics:
    1. **Closed-loop**: early actions affects later actions
    2. **No instructions**
    3. **Reward signals** over a long period of time



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

## Recent Achievements



Figure: Mnih et al (2013) "Playing Atari with Deep Reinforcement Learning"



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# Recent Achievements



Figure: Lee Sedol vs. Alpha Go in 2016



UPPSALA  
UNIVERSITET

## Recent Achievements, also

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes
- Industry automation: RL is used to reduce the energy cost of datacenter cooling



UPPSALA  
UNIVERSITET

## Recent Achievements, also

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes
- Industry automation: RL is used to reduce the energy cost of datacenter cooling
- Automated trading



UPPSALA  
UNIVERSITET

## Recent Achievements, also

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes
- Industry automation: RL is used to reduce the energy cost of datacenter cooling
- Automated trading
- Elevator scheduling





UPPSALA  
UNIVERSITET

## Recent Achievements, also

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- Industry automation: RL is used to reduce the energy cost of datacenter cooling
  - Automated trading
  - Elevator scheduling
  - A/B testing and personalized recommendations



UPPSALA  
UNIVERSITET

# The different parts in RL

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

## 1. The **Agent**: The learning agent.



UPPSALA  
UNIVERSITET

# The different parts in RL

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. The **Agent**: The learning agent.
2. The **Environment**: Where the agent performs actions.



UPPSALA  
UNIVERSITET

# The different parts in RL

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. The **Agent**: The learning agent.
2. The **Environment**: Where the agent performs actions.
3. **Actions**: Made by the agent and affects the environment.



UPPSALA  
UNIVERSITET

# The different parts in RL

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. The **Agent**: The learning agent.
2. The **Environment**: Where the agent performs actions.
3. **Actions**: Made by the agent and affects the environment.
4. **Reward**: The evaluation of an action. A scalar value.  
"Pleasure and pain" of actions.



# The different parts in RL

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. The **Agent**: The learning agent.
2. The **Environment**: Where the agent performs actions.
3. **Actions**: Made by the agent and affects the environment.
4. **Reward**: The evaluation of an action. A scalar value.  
"Pleasure and pain" of actions.
5. **Return**: The aggregated reward over a long period.



UPPSALA  
UNIVERSITET

# The different parts in RL

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

## 1. Agents:

- 1.1 Have a **goal** (maximize return)
- 1.2 **Sense** aspect of their environment
- 1.3 Choose **actions**
- 1.4 Possibility to **improve performance over time**



# The different parts in RL

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

## 1. Agents:

- 1.1 Have a **goal** (maximize return)
- 1.2 **Sense** aspect of their environment
- 1.3 Choose **actions**
- 1.4 Possibility to **improve performance over time**

## 2. Usually an **uncertainty** about the environment

## 3. Represent uncertainty of environment: **Probability**





UPPSALA  
UNIVERSITET

# Sub-elements of agents

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. **Policy**: How the agent choose actions. Determines behaviour.
2. **Model**: The agent's model of the environment. Used for planning



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

## Sub-elements of agents

---

1. **Policy**: How the agent choose actions. Determines behaviour.
2. **Model**: The agent's model of the environment. Used for planning
3. **Value function**: The long-term value (the expected long-term return following a policy)



# Sub-elements of agents

---

1. **Policy**: How the agent choose actions. Determines behaviour.
  2. **Model**: The agent's model of the environment. Used for **planning**
  3. **Value function**: The long-term value (the expected long-term return following a policy)
- Outside the agent: **Reward signal**: The instant value of an action



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. **Policy**: How the agent choose actions. Determines behaviour.
  2. **Model**: The agent's model of the environment. Used for planning
  3. **Value function**: The long-term value (the expected long-term return following a policy)
- Outside the agent: **Reward signal**: The instant value of an action
  - Problem: **Balance** the trade-off between long-term and short-term rewards



UPPSALA  
UNIVERSITET

# Supervised, Unsupervised and Reinforcement Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

## 1. Static vs. Dynamic



UPPSALA  
UNIVERSITET

# Supervised, Unsupervised and Reinforcement Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic
2. No Gold Standard



UPPSALA  
UNIVERSITET

# Supervised, Unsupervised and Reinforcement Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic
2. No Gold Standard
3. Multiple-Decision Process: Return vs. reward



UPPSALA  
UNIVERSITET

# Supervised, Unsupervised and Reinforcement Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic
2. No Gold Standard
3. Multiple-Decision Process: Return vs. reward
4. Need for exploration





UPPSALA  
UNIVERSITET

# Supervised, Unsupervised and Reinforcement Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic
2. No Gold Standard
3. Multiple-Decision Process: Return vs. reward
4. Need for exploration
5. Evaluates actions - not only instruct actions



UPPSALA  
UNIVERSITET

# Exploration vs Exploitation

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal:** Maximize the return (the total reward), i.e.



UPPSALA  
UNIVERSITET

# Exploration vs Exploitation

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal**: Maximize the return (the total reward), i.e.
- **Exploit** the best actions



UPPSALA  
UNIVERSITET

# Exploration vs Exploitation

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal**: Maximize the return (the total reward), i.e.
- **Exploit** the best actions
- **Explore** to know the best actions



UPPSALA  
UNIVERSITET

# Evolution vs Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Set a policy without learning: **Evolutionary** Methods
- Good when the agent cannot sense the environment



UPPSALA  
UNIVERSITET

# Evolution vs Learning

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Set a policy without learning: **Evolutionary** Methods
- Good when the agent cannot sense the environment
- **Example:** Bacteria don't learn, they evolve



# Setting the goal for the Agent

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- Setting the goal: **defining the reward** signal (reward function)
  - **Example:** If you want the agent to do something quick, give -1 per action.



# Setting the goal for the Agent

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Setting the goal: **defining the reward** signal (reward function)
- **Example:** If you want the agent to do something quick, give -1 per action.
- We should give rewards for correct **behaviour**
- Do **not** use reward to guide **how** to reach the goal
- **Be careful what you wish for...**





UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

## Section 2

### Bandits



UPPSALA  
UNIVERSITET

# The $k$ -armed bandit problem

---

- **Goal:** Maximize the total or average reward after  $N$  actions

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# The $k$ -armed bandit problem

---

- **Goal:** Maximize the total or average reward after  $N$  actions
- **The actions:** Choose between  $k$  arms, i.e.  $A_t \in \{1, \dots, k\}$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# The $k$ -armed bandit problem

---

- **Goal:** Maximize the total or average reward after  $N$  actions
- **The actions:** Choose between  $k$  arms, i.e.  $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(R_t|a),$$

where  $\mathbb{E}(R_t|A_t = a) = q^*(a)$ .



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# The $k$ -armed bandit problem

---

- **Goal:** Maximize the total or average reward after  $N$  actions
- **The actions:** Choose between  $k$  arms, i.e.  $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(R_t|a),$$

where  $\mathbb{E}(R_t|A_t = a) = q^*(a)$ .

- $q^*(a)$  is **unknown**.



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# The $k$ -armed bandit problem

---

- **Goal:** Maximize the total or average reward after  $N$  actions
- **The actions:** Choose between  $k$  arms, i.e.  $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(R_t|a),$$

where  $\mathbb{E}(R_t|A_t = a) = q^*(a)$ .

- $q^*(a)$  is **unknown**.
- The estimated (expected) value if action  $a$  at step  $t$ :  $Q_t(a)$ .



# The $k$ -armed bandit problem

- **Goal:** Maximize the total or average reward after  $N$  actions
- **The actions:** Choose between  $k$  arms, i.e.  $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(R_t|a),$$

where  $\mathbb{E}(R_t|A_t = a) = q^*(a)$ .

- $q^*(a)$  is **unknown**.
- The estimated (expected) value if action  $a$  at step  $t$ :  $Q_t(a)$ .
- This is a **tabular** method/problem:  
We can represent the actions in a table.
- Tabular methods works in small problems  
e.g. A/B testing and dynamic web pages.



# Exploration vs. Exploitation

---

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- Two types of actions:
  1. Exploitation: Choose the action with highest expected reward (short term)
  2. Exploration: Choose action to improve  $Q_t(a)$ , but reduces the reward (long term)
- The **conflict** between exploration and exploitation





UPPSALA  
UNIVERSITET

## $\epsilon$ -greedy

---

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- $\epsilon$ -greedy:  $P(\text{exploration}) = \epsilon$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- $\epsilon$ -greedy:  $P(\text{exploration}) = \epsilon$ 
  - Exploitation:

$$A_t = \arg \max_a Q_t(a)$$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- $\epsilon$ -greedy:  $P(\text{exploration}) = \epsilon$

- Exploitation:

$$A_t = \arg \max_a Q_t(a)$$

- Exploration:

$$A_t \sim U(1, \dots, k)$$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- $\epsilon$ -greedy:  $P(\text{exploration}) = \epsilon$

- Exploitation:

$$A_t = \arg \max_a Q_t(a)$$

- Exploration:

$$A_t \sim U(1, \dots, k)$$

- $Q_1(a) = 0$  (or used to encourage initial exploration)



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- $\epsilon$ -greedy:  $P(\text{exploration}) = \epsilon$

- Exploitation:

$$A_t = \arg \max_a Q_t(a)$$

- Exploration:

$$A_t \sim U(1, \dots, k)$$

- $Q_1(a) = 0$  (or used to encourage initial exploration)
- For any  $\epsilon > 0$ ,  $Q_t(a) \rightarrow q^*(a)$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

## $\epsilon$ -greedy

---

- We estimate  $q^*(a)$  using  $Q_t(a)$  as

$$Q_T(a) = \frac{1}{N(a)} \sum_t^{T-1} R_{t, A_t=a} ,$$

where  $N(a)$  is the total number of times action  $a$  has been taken.



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

## $\epsilon$ -greedy

---

- We estimate  $q^*(a)$  using  $Q_t(a)$  as

$$Q_T(a) = \frac{1}{N(a)} \sum_t^{T-1} R_{t,A_t=a} ,$$

where  $N(a)$  is the total number of times action  $a$  has been taken.

- When should we explore?
  - Large  $V(R_t)$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

## $\epsilon$ -greedy

---

- We estimate  $q^*(a)$  using  $Q_t(a)$  as

$$Q_T(a) = \frac{1}{N(a)} \sum_t^{T-1} R_{t,A_t=a},$$

where  $N(a)$  is the total number of times action  $a$  has been taken.

- When should we explore?
  - Large  $V(R_t)$
  - Large  $\mathcal{A}$





- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

## $\epsilon$ -greedy

---

- We estimate  $q^*(a)$  using  $Q_t(a)$  as

$$Q_T(a) = \frac{1}{N(a)} \sum_t^{T-1} R_{t, A_t=a},$$


where  $N(a)$  is the total number of times action  $a$  has been taken.

- When should we explore?
  - Large  $V(R_t)$
  - Large  $\mathcal{A}$
  - Non-stationarity



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# Bandit example




**Figure:** The 10-armed bandit environment (Sutton and Barto, 2017, Fig. 2.1)



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# Bandit example




**Figure:** The  $\epsilon$ -greedy algorithm result in the 10-armed bandit (Sutton and Barto, 2017, Fig. 2.2)



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

## Bandit example: Optimistic initialization



**Figure:** The  $\epsilon$ -greedy algorithm and optimistic initialization (Sutton and Barto, 2017, Fig. 2.3)



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- Compute  $Q_t(a)$  on the fly:

$$Q_T(a) = Q_{T-1} + \frac{1}{N_t(a)} (R_{t,A_t=a} - Q_{T-1}(a))$$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- Compute  $Q_t(a)$  on the fly:

$$Q_T(a) = Q_{T-1} + \frac{1}{N_t(a)}(R_{t,A_t=a} - Q_{T-1}(a))$$

- Handling non-stationarity:

$$Q_T(a) = Q_{T-1} + \alpha(t)(R_{t,A_t=a} - Q_{T-1}(a))$$



UPPSALA  
UNIVERSITET

# Efficient computation and non-stationarity

---

$$Q_T(a) = Q_{T-1} + \alpha(t)(R_{t,A_t=a} - Q_{T-1}(a))$$

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- **Examples:**

- $\alpha(t) = 1$ :  $Q_T(a) = R_{t,A_t=a}$



$$Q_T(a) = Q_{T-1} + \alpha(t)(R_{t,A_t=a} - Q_{T-1}(a))$$

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Examples:
  - $\alpha(t) = 1$ :  $Q_T(a) = R_{t,A_t=a}$
  - $\alpha(t) = 0$ :  $Q_T(a) = Q_1(a)$





UPPSALA  
UNIVERSITET

# Efficient computation and non-stationarity

$$Q_T(a) = Q_{T-1} + \alpha(t)(R_{t,A_t=a} - Q_{T-1}(a))$$

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- **Examples:**

- $\alpha(t) = 1$ :  $Q_T(a) = R_{t,A_t=a}$
- $\alpha(t) = 0$ :  $Q_T(a) = Q_1(a)$
- $\alpha(t) = \frac{1}{N_t(a)}$ : Average reward



$$Q_T(a) = Q_{T-1} + \alpha(t)(R_{t,A_t=a} - Q_{T-1}(a))$$

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- Examples:
  - $\alpha(t) = 1$ :  $Q_T(a) = R_{t,A_t=a}$
  - $\alpha(t) = 0$ :  $Q_T(a) = Q_1(a)$
  - $\alpha(t) = \frac{1}{N_t(a)}$ : Average reward
- $Q_T(a) \rightarrow q^*(a)$ , if:



$$Q_T(a) = Q_{T-1} + \alpha(t)(R_{t,A_t=a} - Q_{T-1}(a))$$

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Examples:

- $\alpha(t) = 1$ :  $Q_T(a) = R_{t,A_t=a}$
- $\alpha(t) = 0$ :  $Q_T(a) = Q_1(a)$
- $\alpha(t) = \frac{1}{N_t(a)}$ : Average reward
- $Q_T(a) \rightarrow q^*(a)$ , if:
  1.  $\sum_t \alpha_t = \infty$
  2.  $\sum_t \alpha_t^2 < \infty$
- Where have we seen these criterias before?



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# The $\epsilon$ -greedy algorithm

## A simple bandit algorithm

Initialize, for  $a = 1$  to  $k$ :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Repeat forever:

$$A \leftarrow \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \epsilon \quad (\text{breaking ties randomly}) \\ \text{a random action} & \text{with probability } \epsilon \end{cases}$$

$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$

$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

Figure: The  $\epsilon$ -greedy algorithm



UPPSALA  
UNIVERSITET

# The Upper-Confidence-Bound method

---

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

- Explore based on our uncertainty of  $Q_t(a)$



- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# The Upper-Confidence-Bound method

- Explore based on our uncertainty of  $Q_t(a)$
- The Upper-Confidence-Bound (UCB) method

$$A_t = \arg \max_a \left( Q_t + c \sqrt{\frac{\log t}{N_t(a)}} \right)$$

An analogy:

$$A_t = \arg \max_a \left( Q_t + c \sqrt{\frac{\hat{\sigma}^2(a)}{N_t(a)}} \right)$$



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- **Bandits**
- Markov Decision Processes

# The UCB algorithm

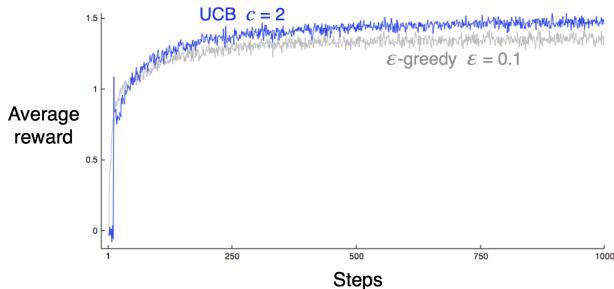


Figure: The UCB algorithm



# The Bayesian Bandit: Thompson sampling

- A Bayesian Bandit
  1. Setup a likelihood for  $R$ ,  $p(R|a, \theta)$
  2. Setup a prior for  $\theta$ ,  $p(\theta)$
  3. Compute posterior for  $\theta$ ,  $p(\theta|R, a)$
  4. Choose action  $A_t$  proportional to

$$\int I[\mathbb{E}(R|\theta, a^*) = \arg \max_{a'} \mathbb{E}(R|\theta, a')] p(\theta|R, a) d\theta$$

where  $I$  is the indicator function.

- Repeat step 3-4





# The Bayesian Bandit: Thompson sampling

- A Bayesian Bandit

1. Setup a likelihood for  $R$ ,  $p(R|a, \theta)$
2. Setup a prior for  $\theta$ ,  $p(\theta)$
3. Compute posterior for  $\theta$ ,  $p(\theta|R, a)$
4. Choose action  $A_t$  proportional to

$$\int I[\mathbb{E}(R|\theta, a^*) = \arg \max_{a'} \mathbb{E}(R|\theta, a')] p(\theta|R, a) d\theta$$

where  $I$  is the indicator function.

- Repeat step 3-4
- Monte Carlo approximation of step 4:
  1. Draw one sample from the posterior  $\tilde{\theta}$

$$\tilde{\theta} \sim p(\theta|R, a)$$

2. Conditional on  $\tilde{\theta}$ , choose action  $A_t$  proportional to

$$\arg \max_{a'} \mathbb{E}(R|\tilde{\theta}, a'),$$

i.e. sample  $A_t$  based on the probability of being the best action.



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

## Section 3

# Markov Decision Processes



UPPSALA  
UNIVERSITET

# The Markov Decision process

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Bandits does not have a *state*.



UPPSALA  
UNIVERSITET

# The Markov Decision process

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Bandits does not have a *state*.
- An action might **change** the environment.
- An action might be different in different **states**



# The Markov Decision process

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Bandits does not have a *state*.
- An action might **change** the environment.
- An action might be different in different **states**
- **Example:** In chess, we want to make a move based on the current position of all pieces



# The Markov Decision process

---


- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Bandits does not have a *state*.
- An action might **change** the environment.
- An action might be different in different **states**
- **Example:** In chess, we want to make a move based on the current position of all pieces
- To capture this we use a **Markov Decision process**
- One of the most important concepts in Reinforcement Learning



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# The Markov Decision process



**Figure:** The (finite) Markov Decision Process (Sutton and Barto, 2017, Fig 3.1)

- States  $S_t \in \mathcal{S}$ : Basis for action
- Actions  $A_t \in \mathcal{A}$
- Rewards  $R_t \in \mathbb{R}$



UPPSALA  
UNIVERSITET

# The Markov Decision process

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- Boundry between Agent and Environment:
    - The **total control** of the action





# The Markov Decision process

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- Boundry between Agent and Environment:
    - The **total control** of the action
    - Reward is **external** to agent: Pain and pleasure
    - The agent **should not be able to change the reward function**



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Boundry between Agent and Environment:
  - The **total control** of the action
  - Reward is **external** to agent: Pain and pleasure
  - The agent **should not be able to change the reward function**
- The policy ( $\pi(A_t|S_t = s)$ ):
  - We make an action given the current state  $S_t$
- **The goal**: (Again) maximize return  $G_t = R_{t+1} + \dots + R_T$



## Return and discount

---

- Two type of interactions
  - **Episodic**:  $T < \infty$ , has terminal state
  - **Continuing**:  $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



## Return and discount

---

- Two type of interactions
  - **Episodic**:  $T < \infty$ , has terminal state
  - **Continuing**:  $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discount rate  $\gamma$ :
  - $0 \leq \gamma \leq 1$



# Return and discount

- Two type of interactions
  - **Episodic**:  $T < \infty$ , has terminal state
  - **Continuing**:  $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discount rate  $\gamma$ :
  - $0 \leq \gamma \leq 1$
  - $\gamma = 1$ : **No discount**
  - $\gamma = 0$ : **Full discount**: Only next reward counts
  - $\gamma < 1$  and  $R_t$  is bounded:  $G_t < \infty$



## Return and discount

- Two type of interactions
  - **Episodic**:  $T < \infty$ , has terminal state
  - **Continuing**:  $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discount rate  $\gamma$ :
  - $0 \leq \gamma \leq 1$
  - $\gamma = 1$ : **No discount**
  - $\gamma = 0$ : **Full discount**: Only next reward counts
  - $\gamma < 1$  and  $R_t$  is bounded:  $G_t < \infty$
- For episodic problem we assume  $R_{T+i} = 0$  for all  $i \in \mathbb{N}^+$



UPPSALA  
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# The Markov Decision Process

---

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- Eq. (1) **fully specify** a MDP



# The Markov Decision Process

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- Eq. (1) **fully specify** a MDP
- Markov property:

$$P(S_{t+1} = s', R_{t+1} = r | S_1 = s, A_1 = a, \dots, S_t = s, A_t = a) = \\ P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a)$$

- The MDP is a good **approximation or model**:  
All models are wrong, but some are useful.





- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# The Markov Decision Process Marginals

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- From Eq. (1) we can get marginals of interest:
  - State-action rewards:

$$r(s, a) = \mathbb{E}(R_{t+1} | S_t = s, A_t = a)$$



# The Markov Decision Process Marginals

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- From Eq. (1) we can get marginals of interest:
  - State-action rewards:

$$r(s, a) = \mathbb{E}(R_{t+1} | S_t = s, A_t = a)$$

- State-transition probability:

$$p(s' | s, a) = P(S_{t+1} = s' | S_t = s, A_t = a)$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- The value function  $v_\pi(s)$ :  
the long-term value of  $s$  given a policy  $\pi(a|s)$ , i.e.

$$v_\pi(s) = \mathbb{E}_\pi(G_t | S_t = s) = \mathbb{E}_\pi \left( \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right)$$



# The Value function

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- The value function  $v_\pi(s)$ :  
the long-term value of  $s$  given a policy  $\pi(a|s)$ , i.e.

$$v_\pi(s) = \mathbb{E}_\pi(G_t | S_t = s) = \mathbb{E}_\pi \left( \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right)$$

- Informally: How "good" is a state for the agent with the policy  $\pi$ .



UPPSALA  
UNIVERSITET

# The Value function

---

- Estimating  $v_{\pi}(s)$  is one of the most important problem in RL

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# The Value function

---

- Estimating  $v_\pi(s)$  is one of the most important problem in RL
- Value functions are **recursive**:

$$\begin{aligned}v_\pi(s) &= \mathbb{E}_\pi(G_t | S_t = s) \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a)(r + \gamma v_\pi(s'))\end{aligned}$$

- This is the **Bellman equation** for  $v_\pi(s)$ :  
The relationship between the values of the state and its successor states.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

# The Value function

- Estimating  $v_\pi(s)$  is one of the most important problem in RL
- Value functions are **recursive**:



$$\begin{aligned}v_\pi(s) &= \mathbb{E}_\pi(G_t | S_t = s) \\&= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a)(r + \gamma v_\pi(s'))\end{aligned}$$

- This is the **Bellman equation** for  $v_\pi(s)$ :  
The relationship between the values of the state and its successor states.
- Bellman equation is the basis for computing  $v_\pi(s)$  (not part of this course)



# The value function

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes



Actions

3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

Figure: The gridworld equiprobable policy value function





# The Optimal Policy

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- A policy  $\pi$  is **better** than  $\pi'$  if  $v_\pi(s) \geq v_{\pi'}(s)$  for all  $s \in \mathcal{S}$ .



# The Optimal Policy

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- A policy  $\pi$  is **better** than  $\pi'$  if  $v_\pi(s) \geq v_{\pi'}(s)$  for all  $s \in \mathcal{S}$ .
- A policy that is better or equal to all other policies is the **optimal policy**  $\pi_\star$ .



# The Optimal Policy

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- A policy  $\pi$  is **better** than  $\pi'$  if  $v_\pi(s) \geq v_{\pi'}(s)$  for all  $s \in \mathcal{S}$ .
- A policy that is better or equal to all other policies is the **optimal policy**  $\pi_\star$ .
- The optimal value function:  $v_{\pi_\star}(s) = v_\star(s)$



# The Optimal Policy

---

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- A policy  $\pi$  is **better** than  $\pi'$  if  $v_\pi(s) \geq v_{\pi'}(s)$  for all  $s \in \mathcal{S}$ .
- A policy that is better or equal to all other policies is the **optimal policy**  $\pi_\star$ .
- The optimal value function:  $v_{\pi_\star}(s) = v_\star(s)$
- The optimal policy  $\pi_\star$  is greedy wrt  $v_{\pi_\star}(s)$ :  
The best long term strategy



UPPSALA  
UNIVERSITET

# The Optimal Policy

---

- Introduction to Reinforcement Learning
  - Bandits
  - Markov Decision Processes
- Computing optimal value function might be **impossible**: we need to estimate/re-estimate/approximate it
  - **Example**: Chess, we cannot compute the optimal long-term moves, we need to approximate/estimate (based on computational budget)




- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Computing optimal value function might be **impossible**: we need to estimate/re-estimate/approximate it
- **Example**: Chess, we cannot compute the optimal long-term moves, we need to approximate/estimate (based on computational budget)
- We might also estimate  $v_*(s)$  better for commonly encountered states



# The value function


- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes



Gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

$v_*$



$\pi_*$

Figure: The gridworld optimal value function and policy