



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Machine learning – Block 8

Måns Magnusson
Department of Statistics, Uppsala University

Autumn 2025



UPPSALA
UNIVERSITET

This week's lectures

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
-
- Introduction to reinforcement learning



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Introduction to Reinforcement Learning

- Another type of machine learning:
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement learning



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Introduction to Reinforcement Learning

- Another type of machine learning:
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement learning
- Computational approach of learning from interaction



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Introduction to Reinforcement Learning

- Another type of machine learning:
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement learning
- Computational approach of learning from interaction
- Closest to human and animal learning: trial, error, and planning.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Introduction to Reinforcement Learning

- Another type of machine learning:
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement learning
- Computational approach of learning from interaction
- Closest to human and animal learning: trial, error, and planning.
- The learner is *not* told which actions to take



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Introduction to Reinforcement Learning

- Another type of machine learning:
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement learning
- Computational approach of **learning from interaction**
- Closest to human and animal learning: **trial, error, and planning.**
- The learner is *not* told which actions to take
- Connections to:
 - Game Theory
 - Control Theory
 - Multi-agent systems
 - Swarm intelligence
 - Information theory
 - Statistics



UPPSALA
UNIVERSITET

Introduction to Reinforcement Learning

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
- **Goal:** maximize return over a sequence of actions





UPPSALA
UNIVERSITET

Introduction to Reinforcement Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal:** maximize return over a sequence of actions
- Three characteristics:
 1. **Closed-loop:** early actions affect later actions



Introduction to Reinforcement Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal:** maximize return over a sequence of actions
- Three characteristics:
 1. **Closed-loop:** early actions affect later actions
 2. **No instructions**



Introduction to Reinforcement Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal:** maximize return over a sequence of actions
- Three characteristics:
 1. **Closed-loop:** early actions affect later actions
 2. **No instructions**
 3. **Reward signals** over a long period of time



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Recent Achievements



Figure: Mnih et al (2013) "Playing Atari with Deep Reinforcement Learning"



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Recent Achievements



Figure: Lee Sedol vs. Alpha Go in 2016



UPPSALA
UNIVERSITET

Recent Achievements, also

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes
- Industry automation: RL is used to reduce the energy cost of datacenter cooling



UPPSALA
UNIVERSITET

Recent Achievements, also

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Industry automation: RL is used to reduce the energy cost of datacenter cooling
- Automated trading



UPPSALA
UNIVERSITET

Recent Achievements, also

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Industry automation: RL is used to reduce the energy cost of datacenter cooling
- Automated trading
- Elevator scheduling



Recent Achievements, also

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Industry automation: RL is used to reduce the energy cost of datacenter cooling
- Automated trading
- Elevator scheduling
- A/B testing and personalized recommendations



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The different parts in RL

1. The Agent: The learning agent.



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The different parts in RL

1. The **Agent**: The learning agent.
2. The **Environment**: Where the agent performs actions.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The different parts in RL

1. The **Agent**: The learning agent.
2. The **Environment**: Where the agent performs actions.
3. **Actions**: Made by the agent and affect the environment.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The different parts in RL

1. **The Agent:** The learning agent.
2. **The Environment:** Where the agent performs actions.
3. **Actions:** Made by the agent and affect the environment.
4. **Reward:** The evaluation of an action. A scalar value.
"Pleasure and pain" of actions.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The different parts in RL

1. **The Agent:** The learning agent.
2. **The Environment:** Where the agent performs actions.
3. **Actions:** Made by the agent and affect the environment.
4. **Reward:** The evaluation of an action. A scalar value.
"Pleasure and pain" of actions.
5. **Return:** The aggregated reward over a long period.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The different parts in RL

1. Agents:

- 1.1 Have a **goal** (maximize return)
- 1.2 **Sense** aspect of their environment
- 1.3 Choose **actions**
- 1.4 Possibility to **improve performance over time**



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The different parts in RL

1. Agents:
 - 1.1 Have a **goal** (maximize return)
 - 1.2 **Sense** aspect of their environment
 - 1.3 Choose **actions**
 - 1.4 Possibility to **improve performance over time**
2. Usually an **uncertainty** about the environment
3. Represent uncertainty of environment: **Probability**



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Sub-elements of agents

1. **Policy**: How the agent chooses actions. Determines behaviour.
2. **Model**: The agent's model of the environment. Used for planning



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Sub-elements of agents

1. **Policy**: How the agent chooses actions. Determines behaviour.
2. **Model**: The agent's model of the environment. Used for planning
3. **Value function**: The long-term value (the expected long-term return following a policy)



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Sub-elements of agents

1. **Policy**: How the agent chooses actions. Determines behaviour.
 2. **Model**: The agent's model of the environment. Used for planning
 3. **Value function**: The long-term value (the expected long-term return following a policy)
-
- Outside the agent: **Reward signal**: The instant value of an action



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Sub-elements of agents

1. **Policy**: How the agent chooses actions. Determines behaviour.
 2. **Model**: The agent's model of the environment. Used for planning
 3. **Value function**: The long-term value (the expected long-term return following a policy)
-
- Outside the agent: **Reward signal**: The instant value of an action
 - Problem: **Balance** the trade-off between long-term and short-term rewards



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Key Distinctions in Reinforcement Learning

Reward vs. Return

- *Reward R_{t+1} :*
 - Immediate, scalar feedback from the environment.
 - Measures short-term desirability of an action.
- *Return G_t :*
 - Cumulative (possibly discounted) future reward.
 - Objective the agent seeks to maximize.

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Policy vs. Value Function

- *Policy $\pi(a | s)$:*
 - Defines how the agent chooses actions.
 - A probability distribution over actions given a state.
- *Value function $v_{\pi}(s)$:*
 - Evaluates how good a state is under a policy.
 - Expected return starting from state s and following π .



Supervised, Unsupervised and Reinforcement Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic



Supervised, Unsupervised and Reinforcement Learning

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
1. Static vs. Dynamic
 2. No Gold Standard





Supervised, Unsupervised and Reinforcement Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic
2. No Gold Standard
3. Multiple-Decision Process: Return vs. reward



Supervised, Unsupervised and Reinforcement Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic
2. No Gold Standard
3. Multiple-Decision Process: Return vs. reward
4. Need for exploration



Supervised, Unsupervised and Reinforcement Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

1. Static vs. Dynamic
2. No Gold Standard
3. Multiple-Decision Process: Return vs. reward
4. Need for exploration
5. Evaluates actions rather than being explicitly instructed which actions to take



Exploration vs Exploitation

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
-
- **Goal:** Maximize the return (the total reward), i.e.



Exploration vs Exploitation

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal:** Maximize the return (the total reward), i.e.
- **Exploit** the best actions



Exploration vs Exploitation

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- **Goal:** Maximize the return (the total reward), i.e.
- **Exploit** the best actions
- **Explore** to know the best actions



Evolution vs Learning

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Set a policy without learning: **Evolutionary** Methods
- Good when the agent cannot sense the environment



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Evolution vs Learning

- Set a policy without learning: Evolutionary Methods
- Good when the agent cannot sense the environment
- Example: Bacteria don't learn, they evolve



Setting the goal for the Agent

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Setting the goal: defining the reward signal (reward function)
- Example: If you want the agent to do something quick, give -1 per action.



Setting the goal for the Agent

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Setting the goal: **defining the reward** signal (reward function)
- **Example:** If you want the agent to do something quick, give -1 per action.
- We should give rewards for correct **behaviour**
- Do **not** use reward to guide **how** to reach the goal
- Be careful what you wish for...



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Bandits vs. Markov Decision Processes

Two levels of sequential decision-making

Bandits

- **No state:** each action is independent of past actions.
- **No dynamics:** actions do not change the environment.
- Feedback is **myopic**: reward depends only on the chosen action.
- **Objective:** identify and exploit the best action.



Bandits vs. Markov Decision Processes

Two levels of sequential decision-making

Bandits

- **No state:** each action is independent of past actions.
- **No dynamics:** actions do not change the environment.
- Feedback is **myopic**: reward depends only on the chosen action.
- **Objective:** identify and exploit the best action.

Markov Decision Processes (MDPs)

- Explicit state describing the environment.
- Actions influence future states and rewards.
- Feedback is **delayed**: consequences unfold over time.
- **Objective:** plan a sequence of actions to maximize long-term return.



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Section 2

Bandits



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The k -armed bandit problem

- **Goal:** Maximize the total or average reward after N actions



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The k -armed bandit problem

- **Goal:** Maximize the total or average reward after N actions
- **The actions:** Choose between k arms, i.e. $A_t \in \{1, \dots, k\}$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The k -armed bandit problem

- **Goal:** Maximize the total or average reward after N actions
- **The actions:** Choose between k arms, i.e. $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(\cdot \mid A_t = a), \quad \mathbb{E}[R_t \mid A_t = a] = q^*(a).$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The k -armed bandit problem

- **Goal:** Maximize the total or average reward after N actions
- **The actions:** Choose between k arms, i.e. $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(\cdot | A_t = a), \quad \mathbb{E}[R_t | A_t = a] = q^*(a).$$

- $q^*(a)$ is **unknown**.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The k -armed bandit problem

- **Goal:** Maximize the total or average reward after N actions
- **The actions:** Choose between k arms, i.e. $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(\cdot \mid A_t = a), \quad \mathbb{E}[R_t \mid A_t = a] = q^*(a).$$

- $q^*(a)$ is **unknown**.
- Estimated expected reward of action a at time t .



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The k -armed bandit problem

- **Goal:** Maximize the total or average reward after N actions
- **The actions:** Choose between k arms, i.e. $A_t \in \{1, \dots, k\}$
- The reward signal:

$$R_t \sim p(\cdot | A_t = a), \quad \mathbb{E}[R_t | A_t = a] = q^*(a).$$

- $q^*(a)$ is **unknown**.
- Estimated expected reward of action a at time t .
- This is a **tabular** method/problem:
We can represent the actions in a table.
- Tabular methods works in small problems
e.g. A/B testing and dynamic web pages.



Exploration vs. Exploitation

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes
- Two types of actions:
 1. Exploitation: Choose the action with highest expected reward (short term)
 2. Exploration: Choose action to improve $Q_t(a)$, but reduces the reward (long term)
- The **conflict** between exploration and exploitation



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- ϵ -greedy: $P(\text{exploration}) = \epsilon$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- ϵ -greedy: $P(\text{exploration}) = \epsilon$

- Exploitation:

$$A_t = \arg \max_a Q_t(a)$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- ϵ -greedy: $P(\text{exploration}) = \epsilon$

- Exploitation:

$$A_t = \arg \max_a Q_t(a)$$

- Exploration:

$$A_t \sim U(1, \dots k)$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- ϵ -greedy: $P(\text{exploration}) = \epsilon$

- Exploitation:

$$A_t = \arg \max_a Q_t(a)$$

- Exploration:

$$A_t \sim U(1, \dots k)$$

- $Q_1(a) = 0$ (or used to encourage initial exploration)



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- ϵ -greedy: $P(\text{exploration}) = \epsilon$

- Exploitation:

$$A_t = \arg \max_a Q_t(a)$$

- Exploration:

$$A_t \sim U(1, \dots k)$$

- $Q_1(a) = 0$ (or used to encourage initial exploration)
- For any fixed $\epsilon > 0$, under stationarity and with sample-average updates,

$$Q_t(a) \xrightarrow{\text{a.s.}} q^*(a) \quad \text{as } t \rightarrow \infty.$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- We estimate $q^*(a)$ using $Q_t(a)$ as

$$Q_T(a) = \frac{1}{N_T(a)} \sum_{t=1}^T R_t \mathbf{1}\{A_t = a\}.$$

where $N_T(a)$ is the total number of times action a has been taken at time point T .



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- We estimate $q^*(a)$ using $Q_t(a)$ as

$$Q_T(a) = \frac{1}{N_T(a)} \sum_{t=1}^T R_t \mathbf{1}\{A_t = a\}.$$

where $N_T(a)$ is the total number of times action a has been taken at time point T .

- When should we explore?
 - Large $V(R_t)$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- We estimate $q^*(a)$ using $Q_t(a)$ as

$$Q_T(a) = \frac{1}{N_T(a)} \sum_{t=1}^T R_t \mathbf{1}\{A_t = a\}.$$

where $N_T(a)$ is the total number of times action a has been taken at time point T .

- When should we explore?
 - Large $V(R_t)$
 - Large \mathcal{A}



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

ϵ -greedy

- We estimate $q^*(a)$ using $Q_t(a)$ as

$$Q_T(a) = \frac{1}{N_T(a)} \sum_{t=1}^T R_t \mathbf{1}\{A_t = a\}.$$

where $N_T(a)$ is the total number of times action a has been taken at time point T .

- When should we explore?
 - Large $V(R_t)$
 - Large \mathcal{A}
 - Non-stationarity



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The ϵ -greedy algorithm

A simple bandit algorithm

Initialize, for $a = 1$ to k :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Repeat forever:

$$A \leftarrow \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \varepsilon \\ \text{a random action} & \text{with probability } \varepsilon \end{cases} \quad (\text{breaking ties randomly})$$

$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$

$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

Figure: The ϵ -greedy algorithm



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

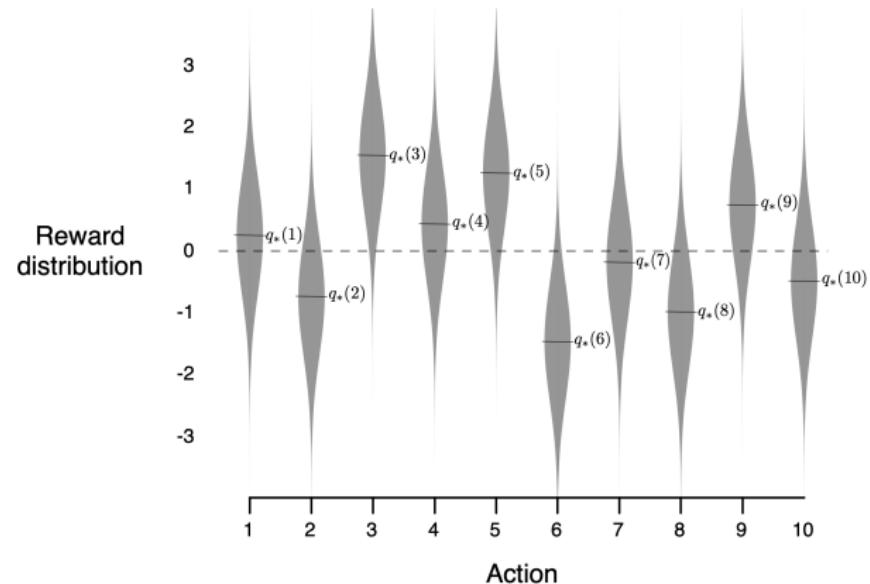


Figure: The 10-armed bandit environment (Sutton and Barto, 2017, Fig. 2.1)



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Bandit example

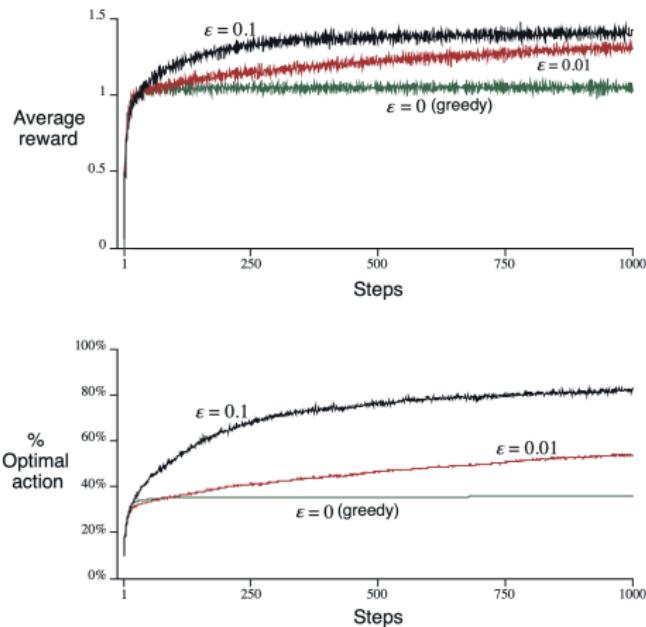


Figure: The ϵ -greedy algorithm result in the 10-armed bandit (Sutton and Barto, 2017, Fig. 2.2)



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Bandit example: Optimistic initialization

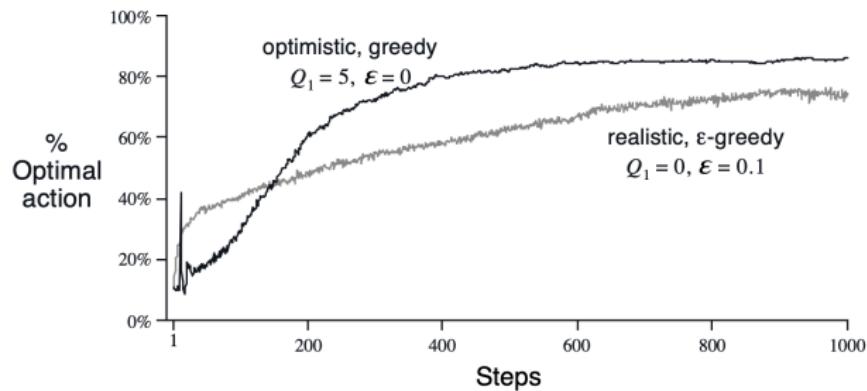


Figure: The ϵ -greedy algorithm and optimistic initialization (Sutton and Barto, 2017, Fig. 2.3)



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Compute $Q_t(a)$ on the fly:

$$Q_T(a) = Q_{T-1} + \frac{1}{N_t(a)}(R_{t,A_t=a} - Q_{T-1}(a))$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Compute $Q_t(a)$ on the fly:

$$Q_T(a) = Q_{T-1} + \frac{1}{N_t(a)}(R_{t,A_t=a} - Q_{T-1}(a))$$

- Handling non-stationarity:

$$Q_T(a) = Q_{T-1} + \alpha(t)(R_{t,A_t=a} - Q_{T-1}(a))$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

$$Q_{t+1}(a) = Q_t(a) + \alpha_t(R_{t+1} - Q_t(a)).$$

- Examples:
 - $\alpha(t) = 1$: $Q_T(a) = R_{t,A_t=a}$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

$$Q_{t+1}(a) = Q_t(a) + \alpha_t(R_{t+1} - Q_t(a)).$$

- Examples:
 - $\alpha(t) = 1$: $Q_T(a) = R_{t,A_t=a}$
 - $\alpha(t) = 0$: $Q_T(a) = Q_1(a)$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Efficient computation and non-stationarity

$$Q_{t+1}(a) = Q_t(a) + \alpha_t (R_{t+1} - Q_t(a)).$$

- Examples:
 - $\alpha(t) = 1$: $Q_T(a) = R_{t,A_t=a}$
 - $\alpha(t) = 0$: $Q_T(a) = Q_1(a)$
 - $\alpha(t) = \frac{1}{N_t(a)}$: Average reward



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Efficient computation and non-stationarity

$$Q_{t+1}(a) = Q_t(a) + \alpha_t (R_{t+1} - Q_t(a)).$$

- Examples:
 - $\alpha(t) = 1$: $Q_T(a) = R_{t,A_t=a}$
 - $\alpha(t) = 0$: $Q_T(a) = Q_1(a)$
 - $\alpha(t) = \frac{1}{N_t(a)}$: Average reward



Stochastic Approximation Conditions

$$Q_{t+1}(a) = Q_t(a) + \alpha_t(R_{t+1} - Q_t(a)).$$

Robbins–Monro conditions

For convergence of stochastic approximation algorithms, the learning rates must satisfy

$$\sum_{t=1}^{\infty} \alpha_t = \infty, \quad \sum_{t=1}^{\infty} \alpha_t^2 < \infty.$$

Interpretation

- $\sum_t \alpha_t = \infty$: ensures continued learning and prevents premature convergence.
- $\sum_t \alpha_t^2 < \infty$: controls the variance of the stochastic noise and prevents divergence.



Stochastic Approximation Conditions

$$Q_{t+1}(a) = Q_t(a) + \alpha_t(R_{t+1} - Q_t(a)).$$

$$\sum_{t=1}^{\infty} \alpha_t = \infty, \quad \sum_{t=1}^{\infty} \alpha_t^2 < \infty.$$

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
- These are *not heuristics*, but results from **stochastic approximation theory** (Robbins & Monro, 1951).
 - They guarantee convergence of value estimates under stationarity, sufficient exploration, and bounded rewards.
 - Many RL algorithms (e.g. bandits, TD learning, Q-learning) rely on these conditions.

Examples

- Sample-average updates: $\alpha_t = \frac{1}{t}$ (satisfies both conditions).
- Constant step-size: $\alpha_t = \alpha$ ($\sum_t \alpha_t^2 < \infty$; useful in non-stationary settings).



The Upper-Confidence-Bound method

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Explore based on our uncertainty of $Q_t(a)$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Upper-Confidence-Bound method

- Explore based on our uncertainty of $Q_t(a)$
- The Upper-Confidence-Bound (UCB) method

$$A_t = \arg \max_a \left(Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}} \right)$$

where $c > 0$ is a hyperparameter controlling exploration.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Upper-Confidence-Bound method

- Explore based on our uncertainty of $Q_t(a)$
- The Upper-Confidence-Bound (UCB) method

$$A_t = \arg \max_a \left(Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}} \right)$$

where $c > 0$ is a hyperparameter controlling exploration.
An analogy:

$$A_t = \arg \max_a \left(Q_t + c \sqrt{\frac{\hat{\sigma}^2(a)}{N_t(a)}} \right)$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The UCB algorithm

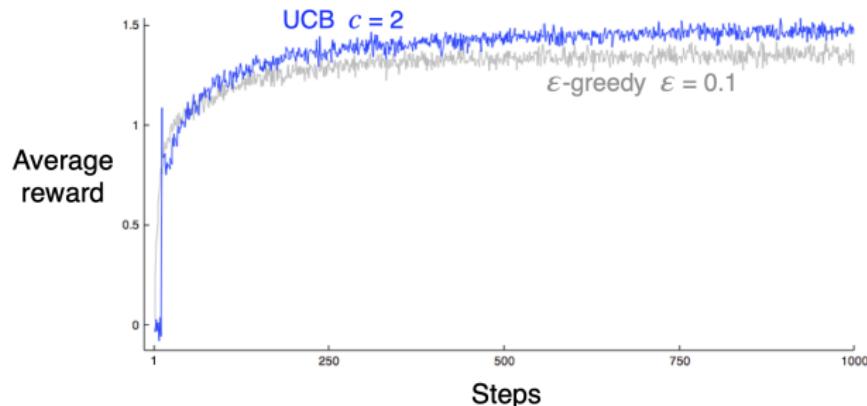


Figure: The UCB algorithm



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Bayesian Bandit: Thompson sampling

- A Bayesian Bandit

1. Setup a likelihood for R , $p(R|a, \theta)$
2. Setup a prior for θ , $p(\theta)$
3. Compute posterior for θ , $p(\theta|D, a)$
4. Choose action A_t by

$$\theta \sim p(\theta | D), \quad A_t = \arg \max_a \mathbb{E}[R | a, \theta],$$

where $D = \{(A_1, R_1), \dots, (A_{t-1}, R_{t-1})\}$.

5. Collect reward R_t .
- Repeat step 3-5. Step 4 can be approximated using Monte Carlo.



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- **Markov Decision Processes**

Section 3

Markov Decision Processes



UPPSALA
UNIVERSITET

The Markov Decision process

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
- Bandits don't have a *state*.





UPPSALA
UNIVERSITET

The Markov Decision process

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
-
- Bandits don't have a *state*.
 - An action might **change** the environment.
 - An action might be different in different **states**



The Markov Decision process

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Bandits don't have a *state*.
- An action might **change** the environment.
- An action might be different in different **states**
- **Example:** In chess, we want to make a move based on the current position of all pieces



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Markov Decision process

- Bandits don't have a *state*.
- An action might **change** the environment.
- An action might be different in different **states**
- **Example:** In chess, we want to make a move based on the current position of all pieces
- To capture this we use a **Markov Decision process**
- One of the most important concepts in reinforcement learning



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Markov Decision process

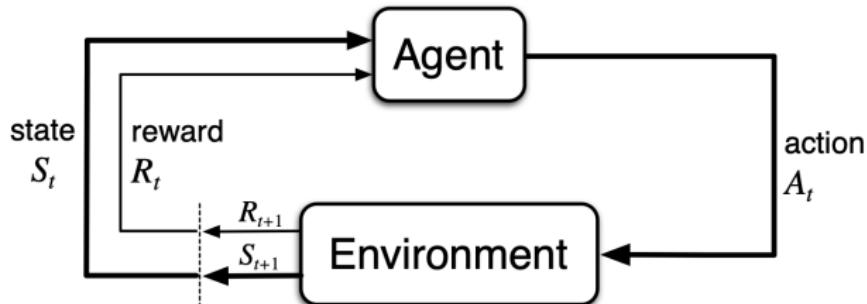


Figure: The (finite) Markov Decision Process (Sutton and Barto, 2017, Fig 3.1)

- States $S_t \in \mathcal{S}$: Basis for action
- Actions $A_t \in \mathcal{A}$
- Rewards $R_t \in \mathbb{R}$



UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Markov Decision process

- Boundary between Agent and Environment:
 - The **total control** of the action



The Markov Decision process

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

- Boundary between Agent and Environment:
 - The **total control** of the action
 - Reward is **external** to agent: Pain and pleasure
 - The agent **should not be able to change the reward function**



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Markov Decision process

- Boundary between Agent and Environment:
 - The **total control** of the action
 - Reward is **external** to agent: Pain and pleasure
 - The agent **should not be able to change the reward function**
- The policy ($\pi(A_t|S_t = s)$):
 - We make an action given the current state S_t
- The goal: (Again) maximize return $G_t = R_{t+1} + \dots + R_T$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Return and discount

- Two types of interactions
 - **Episodic:** $T < \infty$, has terminal state
 - **Continuing:** $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Return and discount

- Two types of interactions
 - **Episodic:** $T < \infty$, has terminal state
 - **Continuing:** $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discount rate γ :
 - $0 \leq \gamma \leq 1$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Return and discount

- Two types of interactions
 - **Episodic:** $T < \infty$, has terminal state
 - **Continuing:** $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discount rate γ :
 - $0 \leq \gamma \leq 1$
 - $\gamma = 1$: **No discount**
 - $\gamma = 0$: **Full discount**: Only next reward counts
 - $\gamma < 1$ and R_t is bounded: $G_t < \infty$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

Return and discount

- Two types of interactions
 - **Episodic:** $T < \infty$, has terminal state
 - **Continuing:** $T = \infty$
- Discounting:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discount rate γ :
 - $0 \leq \gamma \leq 1$
 - $\gamma = 1$: **No discount**
 - $\gamma = 0$: **Full discount**: Only next reward counts
 - $\gamma < 1$ and R_t is bounded: $G_t < \infty$
- For episodic problem we assume $R_{T+i} = 0$ for all $i \in \mathbb{N}^+$



- Introduction to Reinforcement Learning
- Bandits
- **Markov Decision Processes**

The Markov Decision Process

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- An MDP is fully specified by the tuple $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$.



- Introduction to Reinforcement Learning
- Bandits
- **Markov Decision Processes**

The Markov Decision Process

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- An MDP is fully specified by the tuple $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$.
- Markov property:

$$\begin{aligned} P(S_{t+1} = s', R_{t+1} = r | S_1 = s, A_1 = a, \dots, S_t = s, A_t = a) = \\ P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \end{aligned}$$

- The MDP is a good **approximation or model**:
All models are wrong, but some are useful.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Markov Decision Process Marginals

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- From Eq. (1) we can get marginals of interest:
 - State-action rewards:

$$r(s, a) = \mathbb{E}(R_{t+1} | S_t = s, A_t = a)$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Markov Decision Process Marginals

- The Markov Decision process (MDP):

$$P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a) \quad (1)$$

- From Eq. (1) we can get marginals of interest:

- State-action rewards:

$$r(s, a) = \mathbb{E}(R_{t+1} | S_t = s, A_t = a)$$

- State-transition probability:

$$p(s' | s, a) = P(S_{t+1} = s' | S_t = s, A_t = a)$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Value function

- The value function $v_\pi(s)$:
the long-term value of s given a policy $\pi(a|s)$, i.e.

$$v_\pi(s) = \mathbb{E}_\pi(G_t | S_t = s) = \mathbb{E}_\pi \left(\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right)$$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Value function

- The value function $v_\pi(s)$:
the long-term value of s given a policy $\pi(a|s)$, i.e.

$$v_\pi(s) = \mathbb{E}_\pi(G_t | S_t = s) = \mathbb{E}_\pi \left(\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right)$$

- Informally: How "good" is a state for the agent with the policy π .



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Value function

- Estimating $v_\pi(s)$ is one of the most important problem in RL



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Value function

- Estimating $v_\pi(s)$ is one of the most important problem in RL
- Value functions are **recursive**:

$$\begin{aligned}v_\pi(s) &= \mathbb{E}_\pi(G_t | S_t = s) \\&= \sum_a \pi(a|s) \sum_{s',r} p(s', r | s, a) (r + \gamma v_\pi(s'))\end{aligned}$$

- This is the **Bellman equation** for $v_\pi(s)$:
The relationship between the values of the state and its successor states.



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Value function

- Estimating $v_\pi(s)$ is one of the most important problem in RL
- Value functions are **recursive**:

$$\begin{aligned}v_\pi(s) &= \mathbb{E}_\pi(G_t | S_t = s) \\&= \sum_a \pi(a|s) \sum_{s',r} p(s', r | s, a) (r + \gamma v_\pi(s'))\end{aligned}$$

- This is the **Bellman equation** for $v_\pi(s)$:
The relationship between the values of the state and its successor states.
- Bellman equation is the basis for computing $v_\pi(s)$ (not part of this course)



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The value function

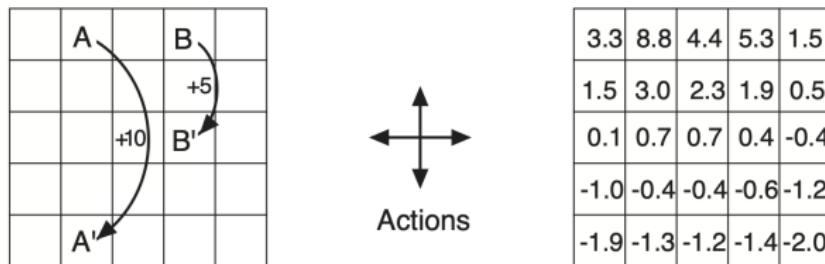


Figure: The gridworld equiprobable policy value function



UPPSALA
UNIVERSITET

The Optimal Policy

- Introduction to Reinforcement Learning
 - Bandits
 - Markov Decision Processes
-
- A policy π is **better** than π' if $v_\pi(s) \geq v_{\pi'}(s)$ for all $s \in \mathcal{S}$.





UPPSALA
UNIVERSITET

- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Optimal Policy

- A policy π is **better** than π' if $v_\pi(s) \geq v_{\pi'}(s)$ for all $s \in \mathcal{S}$.
- A policy that is better or equal to all other policies is the **optimal policy** π_* .



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Optimal Policy

- A policy π is **better** than π' if $v_\pi(s) \geq v_{\pi'}(s)$ for all $s \in \mathcal{S}$.
- A policy that is better or equal to all other policies is the **optimal policy** π_* .
- The optimal value function: $v_{\pi_*}(s) = v_*(s)$



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Optimal Policy

- A policy π is **better** than π' if $v_\pi(s) \geq v_{\pi'}(s)$ for all $s \in \mathcal{S}$.
- A policy that is better or equal to all other policies is the **optimal policy** π_* .
- The optimal value function: $v_{\pi_*}(s) = v_*(s)$
- The optimal policy π_* is greedy wrt $v_{\pi_*}(s)$:
The best long term strategy



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

The Optimal Policy

- Computing optimal value function might be **impossible**: we need to estimate/re-estimate/approximate it
- **Example:** Chess, we cannot compute the optimal long-term moves, we need to approximate/estimate (based on computational budget)



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

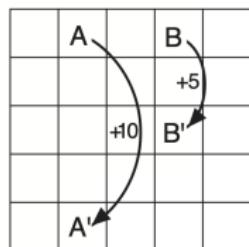
The Optimal Policy

- Computing optimal value function might be **impossible**: we need to estimate/re-estimate/approximate it
- **Example:** Chess, we cannot compute the optimal long-term moves, we need to approximate/estimate (based on computational budget)
- We might also estimate $v_*(s)$ better for commonly encountered states



- Introduction to Reinforcement Learning
- Bandits
- Markov Decision Processes

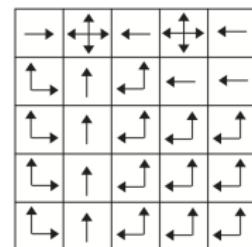
The value function



Gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

v_*



π_*

Figure: The gridworld optimal value function and policy